

EEL-4736/5737

Principles of Computer System Design

Lecture Slides 5

Textbook Chapter 2

Case Study: UNIX File System Layering and
Naming

Introduction

- File system:
 - Builds upon the file abstraction
 - Creates a system to organize and control access to files
 - Hierarchical organization, with support for links
- Objects: files and directories
 - Files hold data
 - Directories hold other objects
- Meta-data
 - Permissions (r/w/x), owner, times

Application Programming Interface

- System calls exposed by O/S
- `fd = OPEN(name, flag, mode)`
 - Flag: e.g. CREATE
 - Mode: r/w/x
 - Pointer/cursor=0 (beginning of file)
 - Return descriptor (or error code)
- `READ/WRITE(fd, buffer, n)`
- `SEEK(fd,offset,whence)`
 - Set cursor to offset relative to begin/end/current
- `CLOSE(fd)`

Application Programming Interface

- **MKDIR(name), RMDIR(name)**
 - Create/remove directory
- **CHDIR(name)**
 - Change working directory
- **CHROOT(name)**
 - Change default root directory
- **MOUNT(name, device)**
 - Associate file system on device with name
- **STAT(name)**
 - Return meta-data

Naming layers


| Layer | Purpose | |
|---------------------------------|--|------------------------|
| Symbolic link | Integrate multiple file systems | User-oriented names |
| Absolute path name layer | Provide a root for naming hierarchies | |
| Path name layer | Organize files into naming hierarchies | |
| File name layer | Human-oriented names | User/machine Interface |
| Inode number layer | Machine-oriented names | Machine-oriented names |
| File layer | Organize blocks into files | |
| Block layer | Identify disk blocks | |

The Block Layer

- Sector – minimum cell of information addressable in a disk
- Block – small number of disk sectors
 - E.g. 8 Kbytes
- Names are block numbers of a disk device
 - E.g. offset from block 0
 - Note: device may have its own layer mapping of block to physical location
 - E.g. hard disks: bad sectors

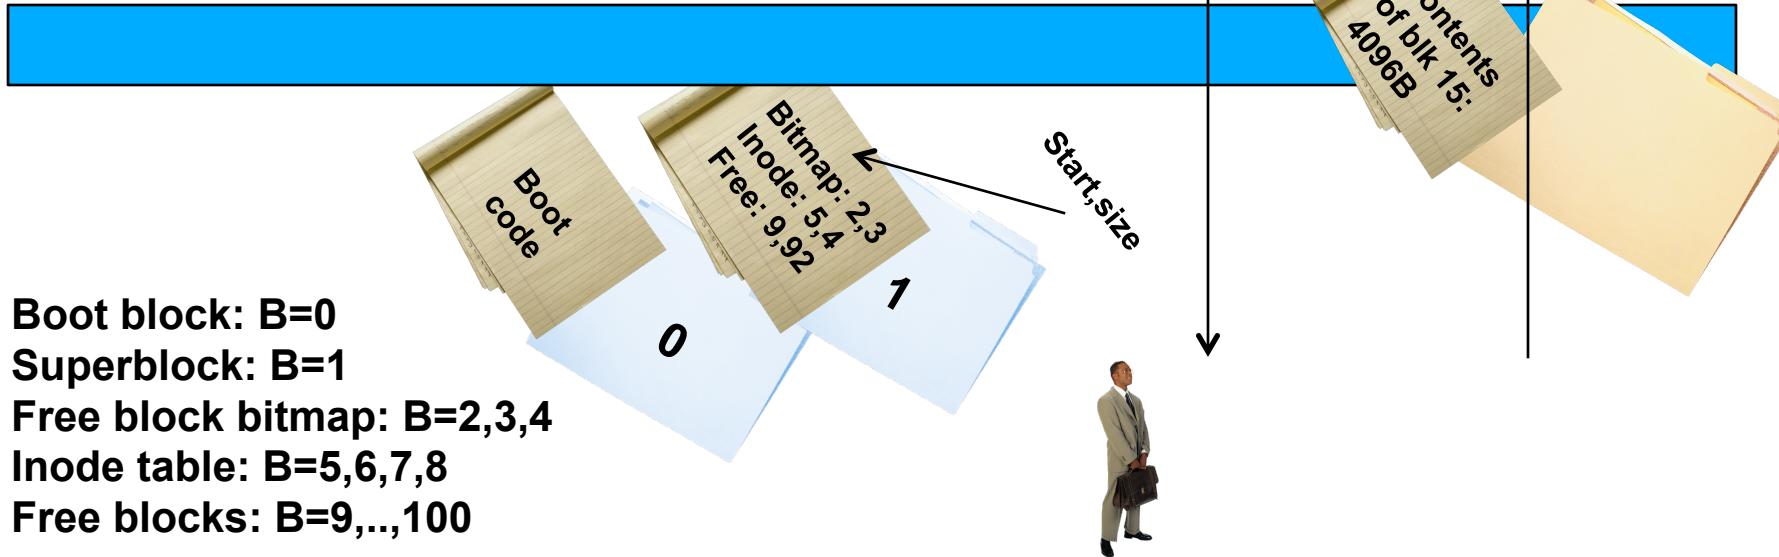
The block layer

- How to discover which blocks are in use?
 - Super-block – well-known name (1)
 - Size of file system, bitmap, inode table
 - Bitmap of free blocks
 - Bit i determines if block i is free or not
 - Well-known name – e.g. after the super-block
 - Boot block also a well-known name (0)



The Block Layer

BLOCK_NUMBER_TO_BLOCK (15)



Block layer


IO_Read(device[15])



Block size: 4096B

Resolve: scan, retrieve block 15

Physical storage layer
Array of blocks



The File Layer

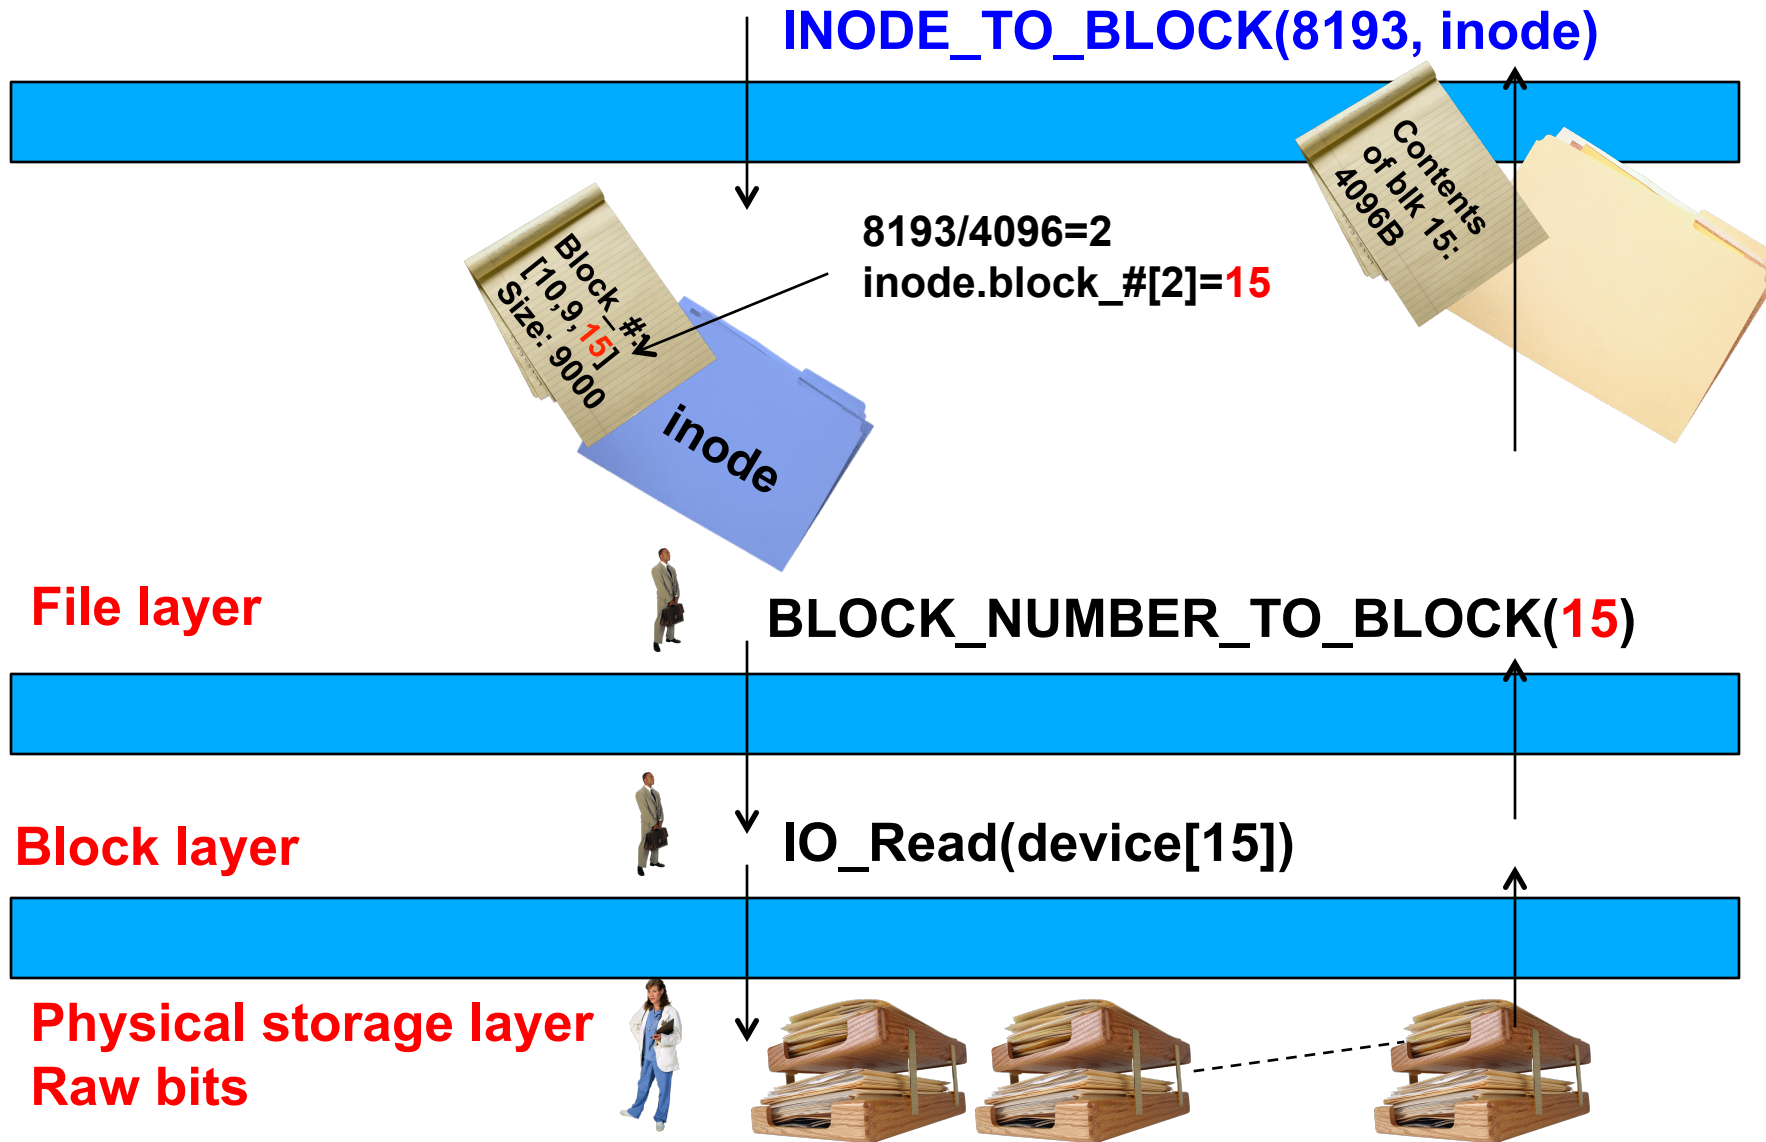
- Objects have 1+ blocks and can grow or shrink over time
- Next naming layer – files
 - File: linear array, arbitrary length
- Which blocks belong to file?
 - Information in index node (*inode*) object

structure *inode*

integer *block_numbers[N]* // which blocks in file

integer *size* // file size

The File Layer

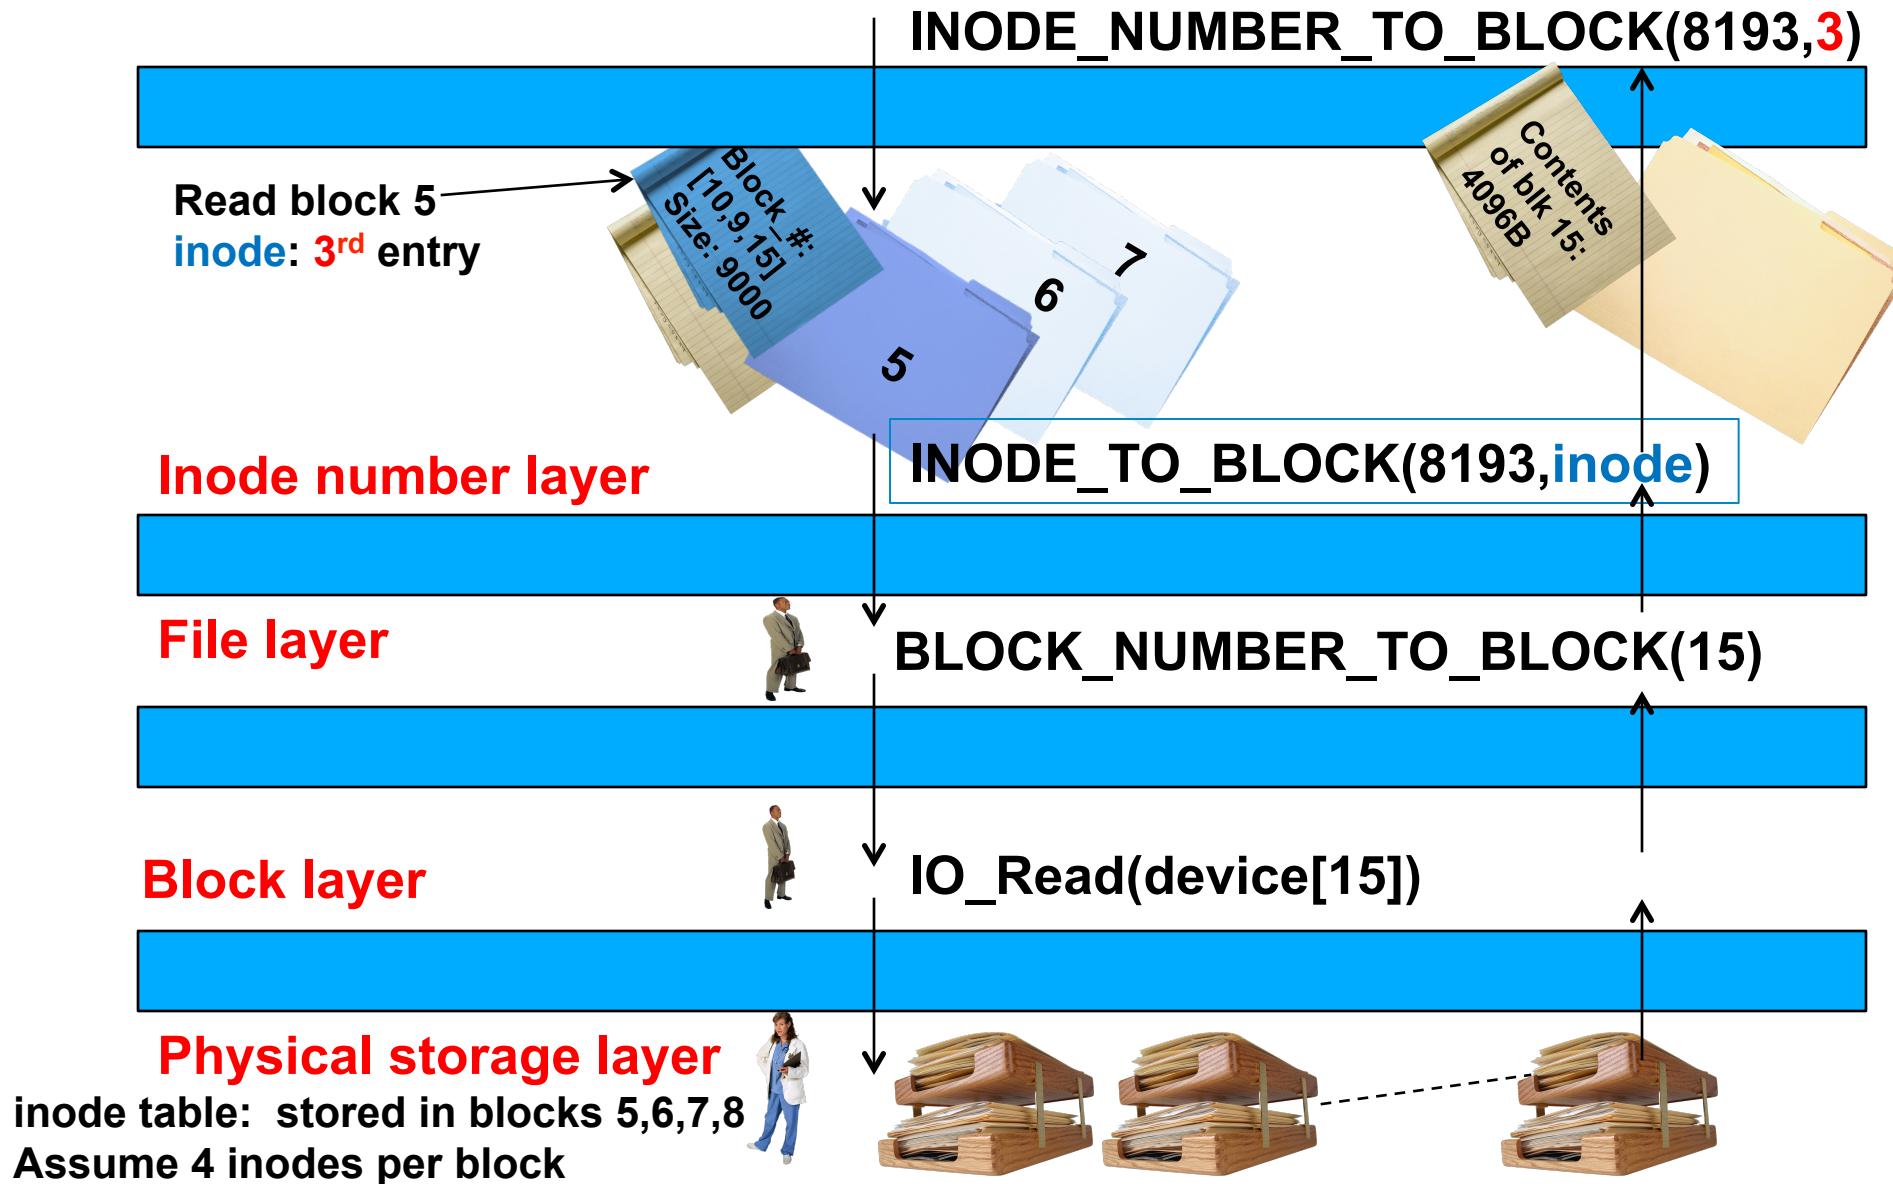


The Inode Number Layer

- Instead of passing inodes themselves, inodes are named
 - Table indexed by inode number
 - Putting it all together – returning the block which contains byte *offset* of file with *inode_number*

```
procedure INODE_NUMBER_TO_BLOCK (int offset, int
    inode_number) returns block
    inode instance i <- INODE_NUMBER_TO_INODE
    (inode_number)
    o <- offset / BLOCKSIZE
    b <- INDEX_TO_BLOCK_NUMBER (i, o)
    return BLOCK_NUMBER_TO_BLOCK (b)
```

The inode Number Layer



Resolving lower layers

```
procedure INODE_NUMBER_TO_INODE  
  (integer inode_number) returns inode  
  return inode_table[inode_number]
```

```
procedure INDEX_TO_BLOCK_NUMBER  
  (inode instance i, integer index) returns  
  integer  
  return i.block_numbers[index]
```

```
procedure BLOCK_NUMBER_TO_BLOCK  
  (integer b) returns block  
  return device[b]
```

The File Name Layer

- inode numbers clearly not convenient to users
 - Another problem: specify a fixed location
- Naming layer introduced to *hide the meta-data of file management*
- Directory
 - A *context* containing set of *bindings* between *character-string names* and *inode numbers*
 - Represented, itself, as a file
 - Extend inode to also have a “*type*”: regular file, or directory, (device, ...)

Directory example

- Each file name associated with an inode number
- Directory – context; represented as file

| File name | Inode number |
|-----------|--------------|
| program | 10 |
| paper | 12 |

A directory with two files

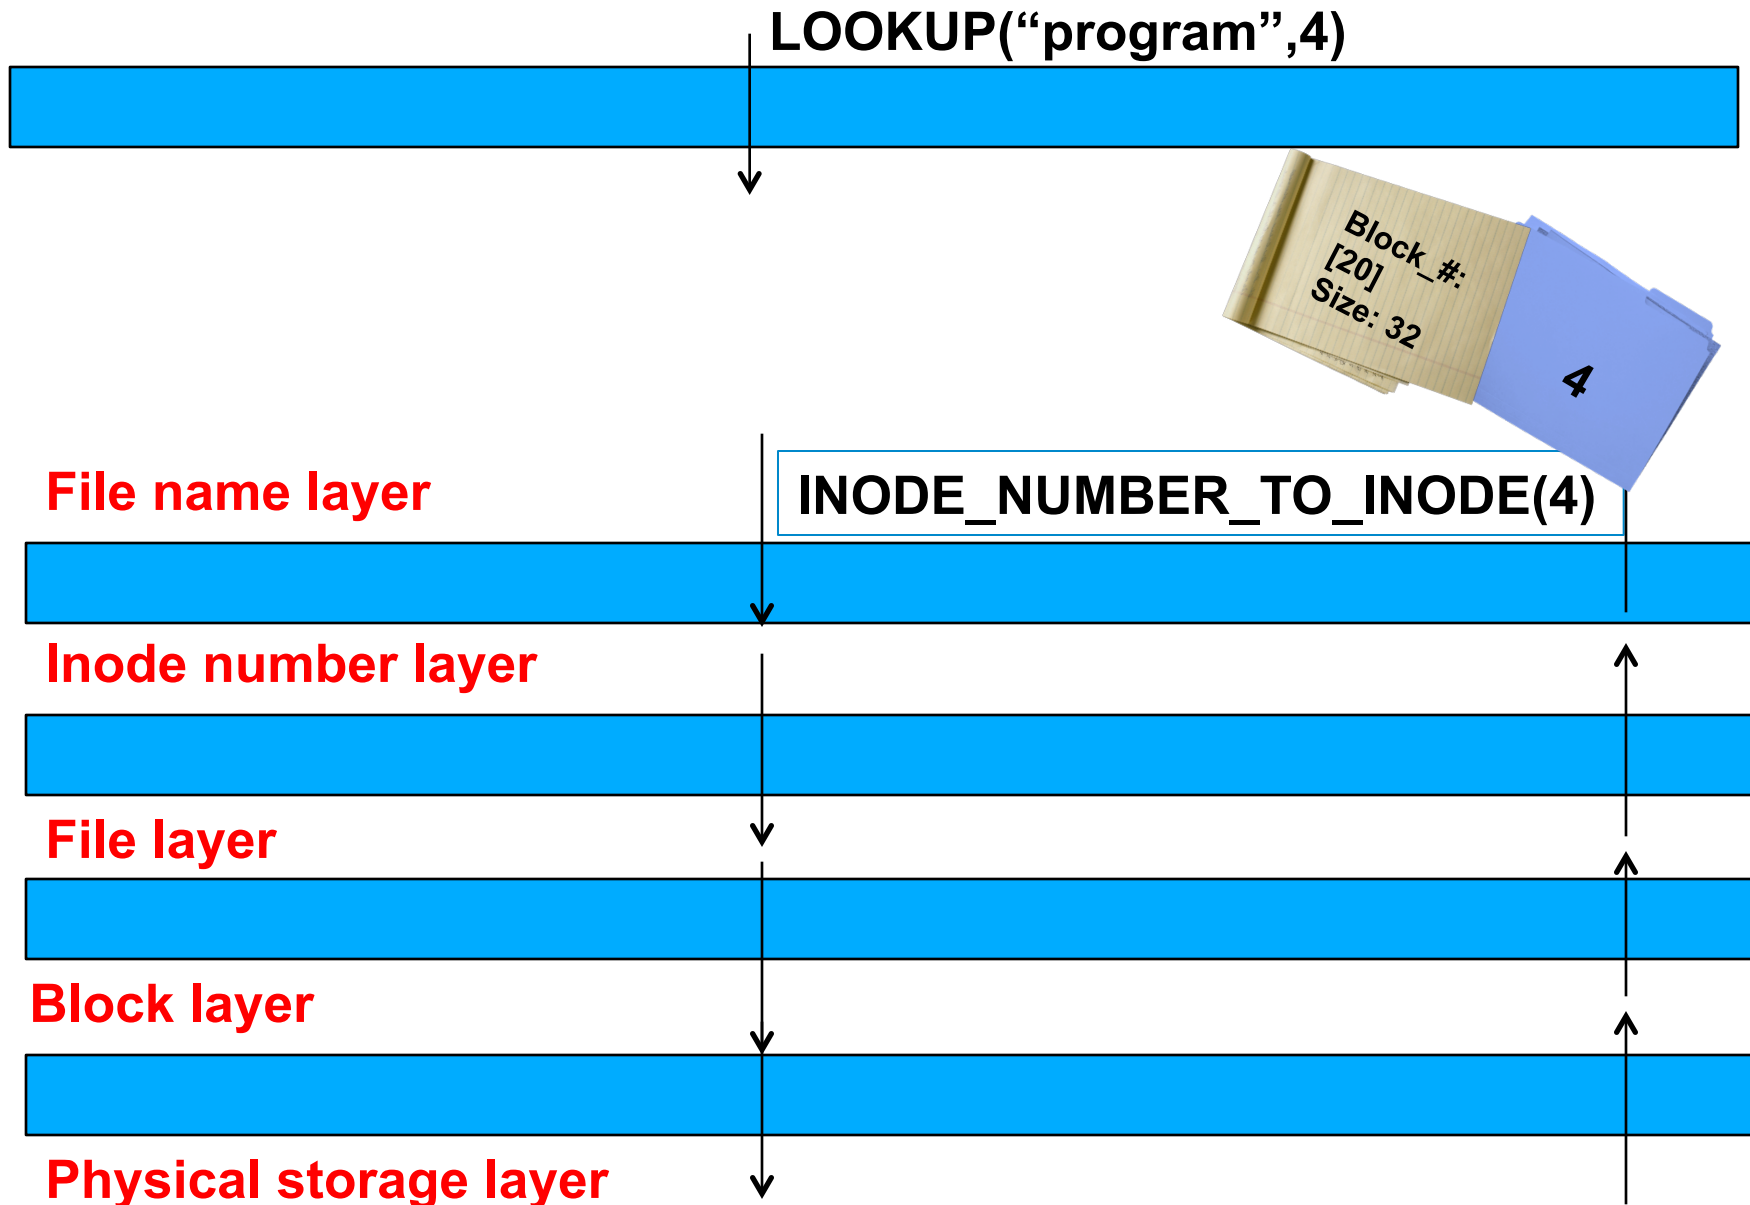
Directory has inode (e.g. inode 4)

Blocks in inode list store **table with context** instead of file data

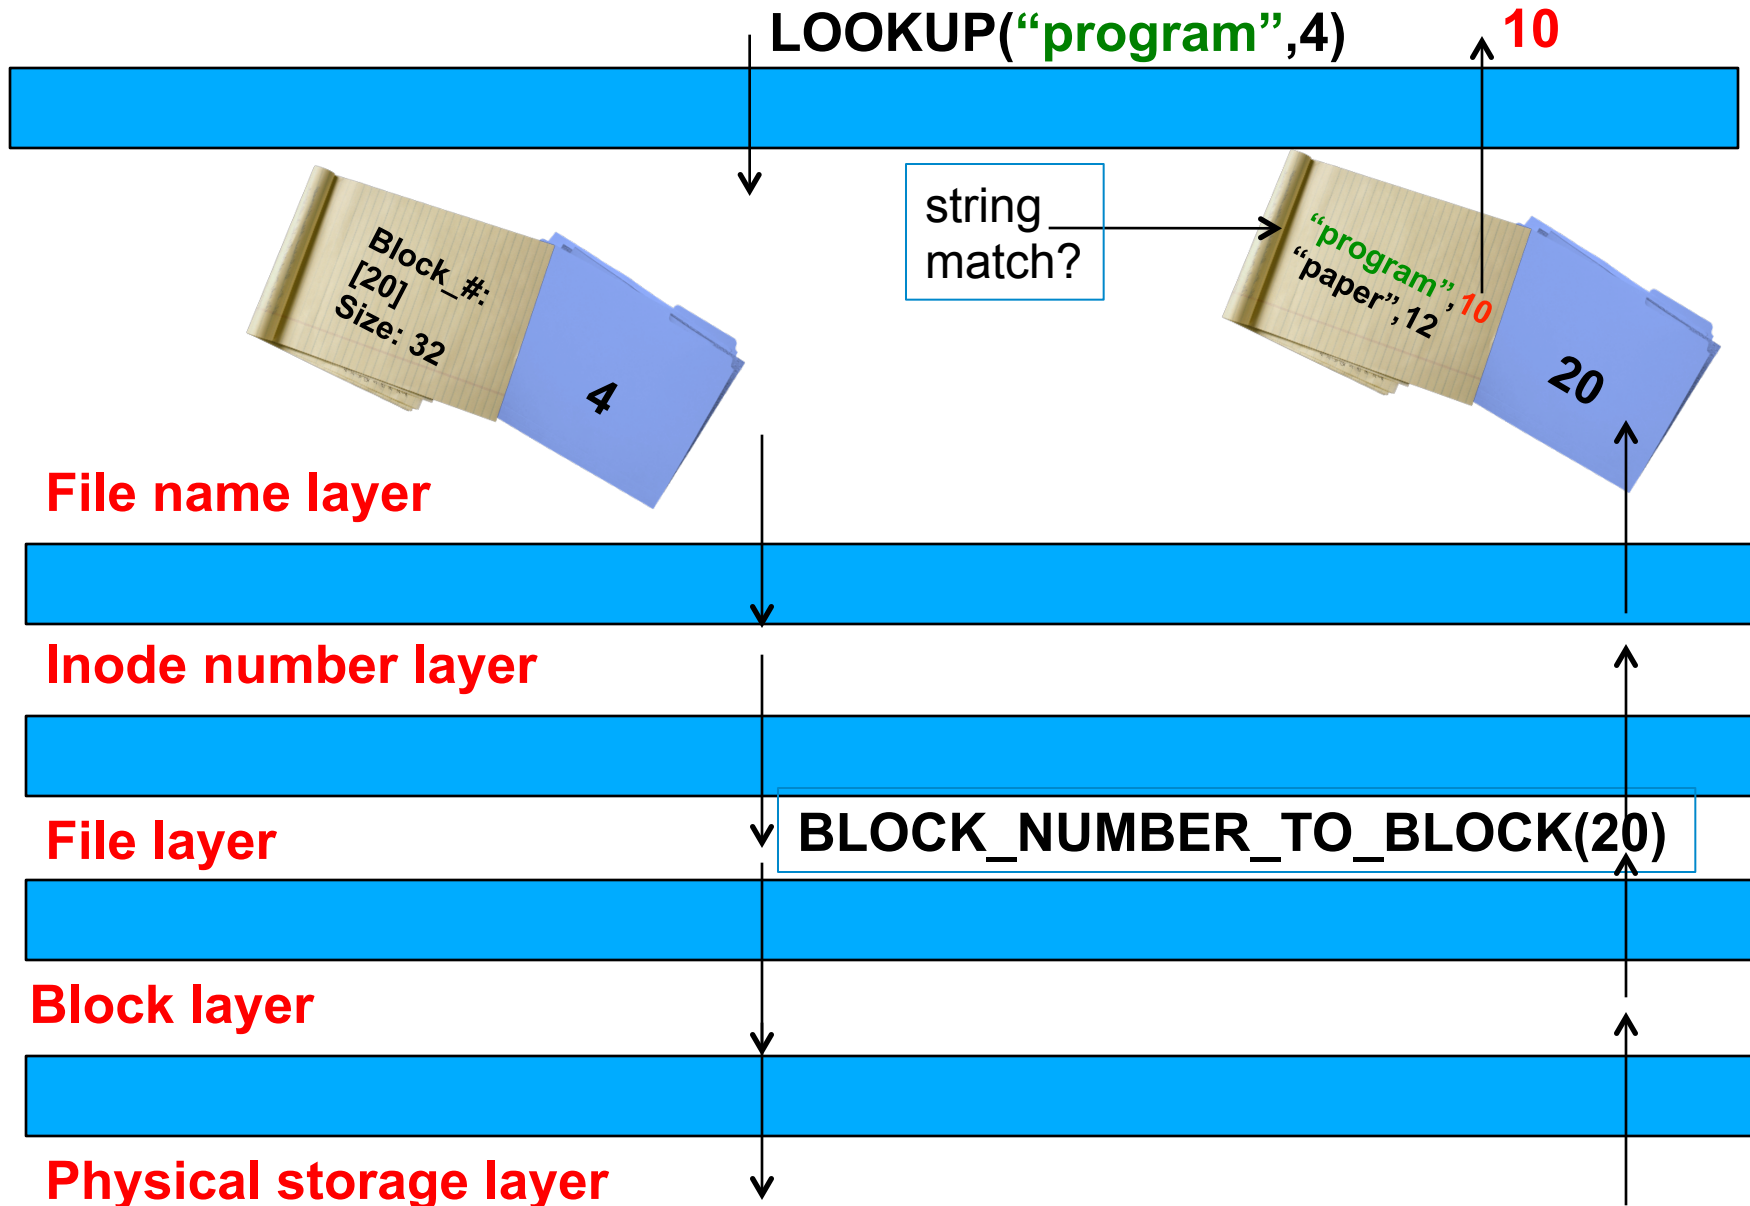
E.g. block #20:



The File Name Layer



The File Name Layer



Lookup

- Resolve *filename* to *inode* in the context of a *directory*

procedure LOOKUP (**string** *filename*, **int** *dir*) **returns integer**

block instance *b*

inode instance *i* **<-** INODE_NUMBER_TO_INODE(*dir*)

if (*i.type* **!=** DIR) **then return** FAILURE

for offset **from** 0 **to** *i.size*-1 **do**

b **<-** INODE_NUMBER_TO_BLOCK (offset, *dir*)

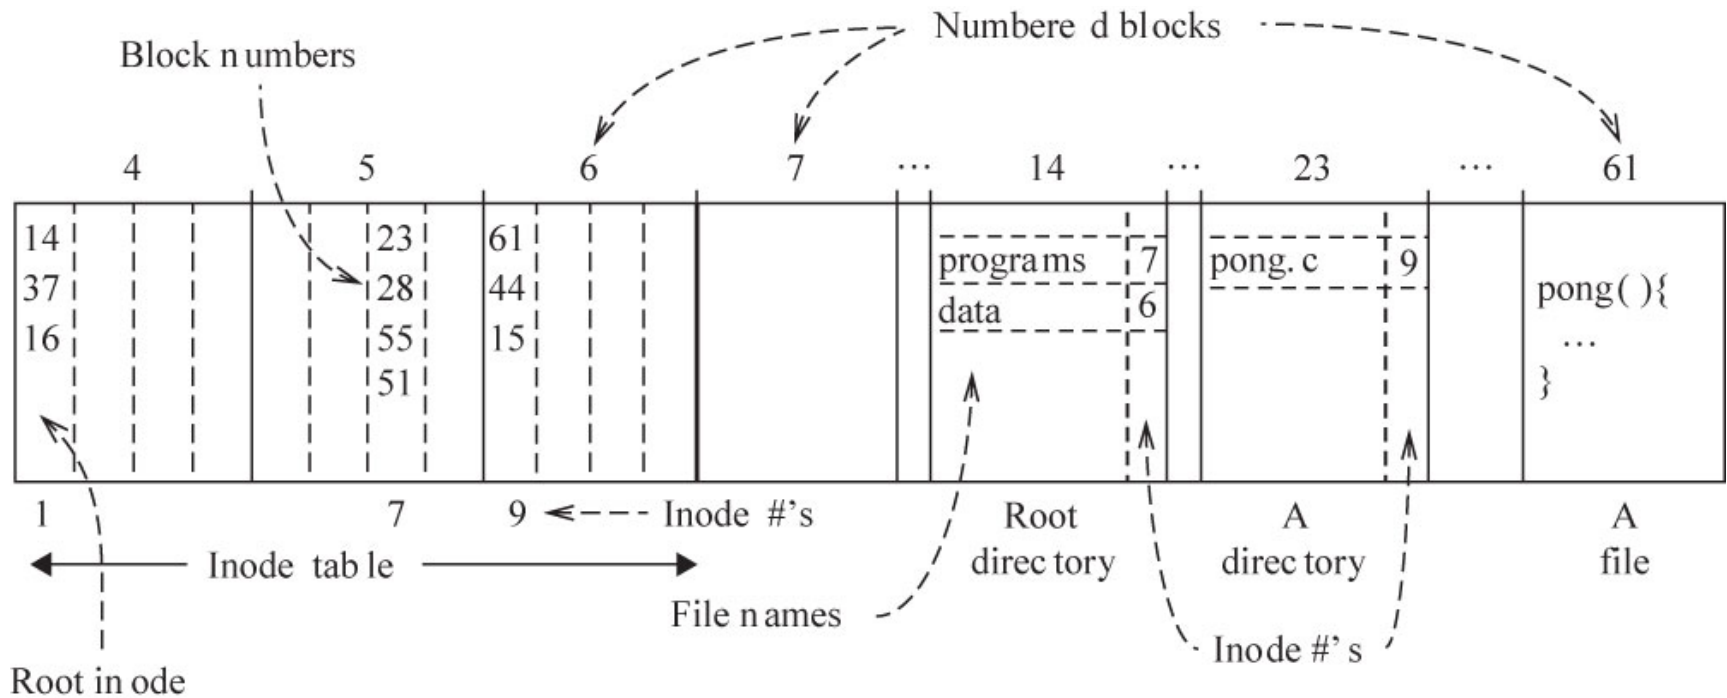
if STRING_MATCH(*filename*, *b*) **then**

return INODE_NUMBER(*filename*, *b*)

 offset = offset + BLOCKSIZE

return FAILURE

Example



The Path Name Layer

- Single directory - difficult to organize files
 - Path names and recursive lookup

procedure PATH_TO_INODE_NUMBER (**string** *path*, **int** *dir*) **returns integer**

if (*PLAIN_NAME(path)*) **return** NAME_TO_INODE_NUMBER(*path*,*dir*)

// plain name – no “/” separator

// NAME_TO_INODE_NUMBER is a LOOKUP

else

dir <- LOOKUP(FIRST(*path*), *dir*) // peel off and lookup first name

path <- REST(*path*) // pick rest of path name

return PATH_TO_INODE_NUMBER(*path*,*dir*) // recursion

The Path Name Layer

- Path names and recursive lookup
 - Result of a LOOKUP(name,dir) is an inode
 - Structured naming scheme provides ability to specify a path to name; recursive lookup
 - PATH_TO_INODE_NUMBER(“dir1/dir2/file”,inode_root) unfolds as:
 - LOOKUP (“dir1”, inode_root) -> inode_dir1
 - FIRST(“dir1/dir2/file”) -> “dir1”; REST -> “dir2/file”
 - PATH_TO_INODE_NUMBER(“dir2/file”,inode_dir1)
 - LOOKUP (“dir2”, inode_dir1) -> inode_dir2
 - PATH_TO_INODE_NUMBER(“file”,inode_dir2)
 - LOOKUP (“file”, inode_dir2) -> inode_file

Links

- Allow creation of *synonyms*
 - Bind a *name* in different *context* to the *same inode*
- E.g. creating a link to existing file in “projects”
 - LINK (“Mail/inbox/new-assignment”, “assignment”)
 - “assignment” is a synonym
 - Assume in context “Mail/inbox”, {“new-assignment”,481}
 - Now, in context “projects”, bind to same inode:
 - {“assignment”, 481)
 - Increment number of links for the inode 481
 - Extend inode data structure to add: **integer refcnt**
 - UNLINKing a file removes a binding
 - If the last binding (refcnt=0), mark inode and blocks as free

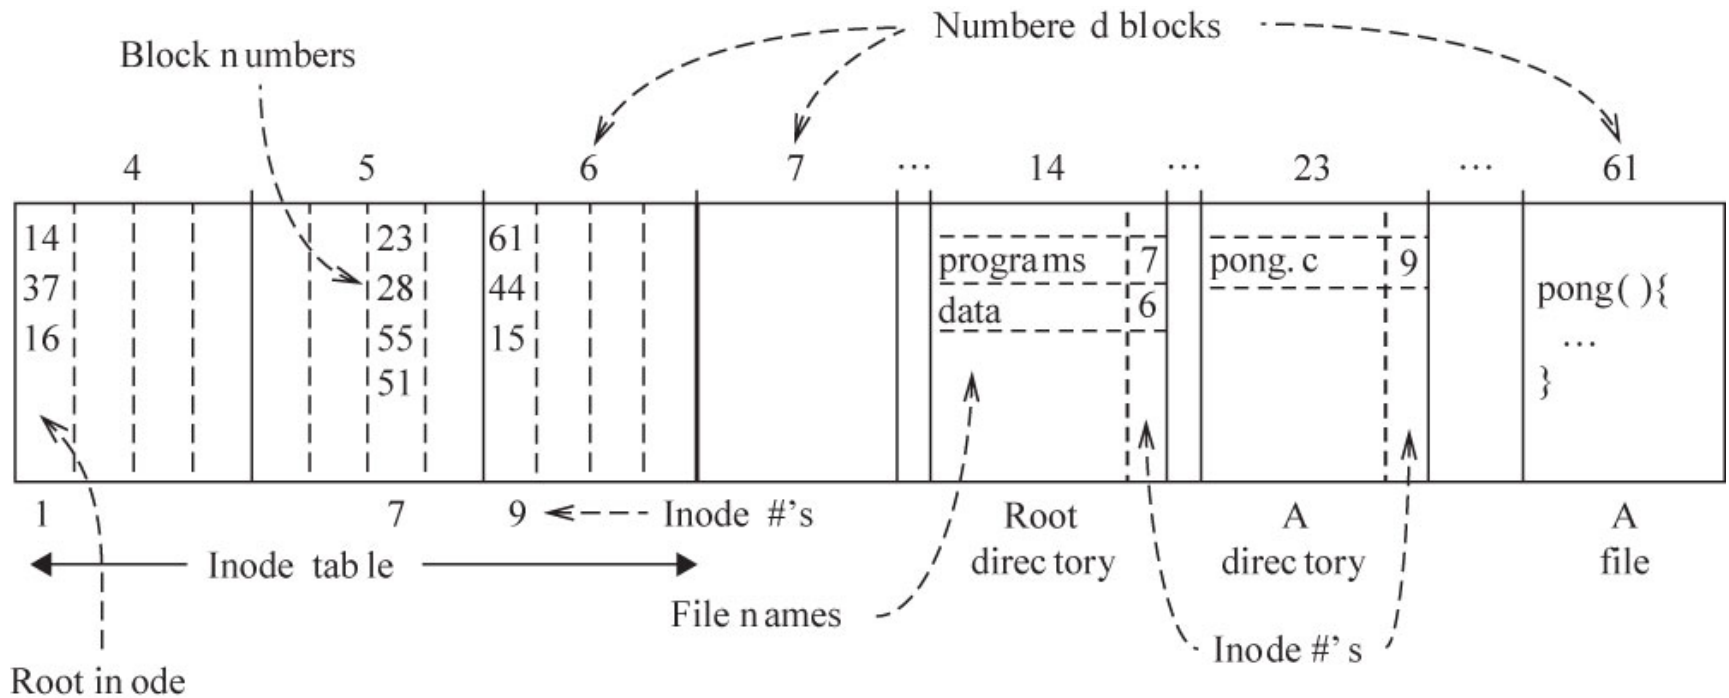
Renaming

- Can be achieved by adding a link, then removing a link:
 - LINK (from_name,to_name)
 - UNLINK (from_name)
- inode does not change; link count goes to 2, then back to 1

Absolute Path Name Layer

- Users often want both private and shared files
 - E.g. shared executable files/libraries with all; private documents; selectively share files with a group
- Need to provide a universal context
 - Root directory
 - Default: *well-known name*: “/”, inode 1
 - CHROOT: can specify different root inode
 - E.g. to constrain a program to a directory

Example



Symbolic Link Layer

- Allow users to create indirect names
- “Hard” links create additional bindings name -> inode
- Symbolic links create name -> name bindings
 - E.g. to name files on other devices
 - Can’t easily guarantee unique inode namespace across multiple devices, e.g. USB drive
- Inode `block_numbers[]` contain characters of a file name rather than block numbers for inodes of type symbolic link
 - “assignments”, “/other/disk/assignment3”

OPEN

- Open existing/create file; create entry in file table; return fd

```
procedure OPEN(string filename, int flags, mode)  
  inode_number <- PATH_TO_INODE_NUMBER(filename, wd)  
  if (inode_number = FAILURE) and flags=O_CREATE then  
    inode_number <- CREATE (filename, mode)  
  if (inode_number = FAILURE) then return FAILURE  
  inode <- INODE_NUMBER_TO_INODE (inode_number)  
  if (PERMITTED(mode,flags)) then  
    file_index <- INSERT (file_table, inode_number)  
    fd <- FIND_UNUSED_ENTRY(fd_table)  
    fd_table[fd] <- file_index  
    return fd  
  else return FAILURE
```

READ

- Read “n” bytes of data from an open file

procedure READ(*fd*, **char array reference** *buf*, **int** *n*)

file_index <- *fd_table*[*fd*]

cursor <- *file_table*[*file_index*].*cursor*

inode <- INODE_NUMBER_TO_INODE
(*file_table*[*file_index*].*inode_number*)

m <- MINIMUM(*inode.size*-*cursor*, *n*)

atime **of** *inode* <- NOW()

if (*m*=0) **then return** END_OF_FILE

for (*i* **from** 0 **to** *m*-1) **do** {

b <- INODE_NUMBER_TO_BLOCK(*i*+*cursor*,
 inode_number)

 COPY(*b*, *buf*, MINIMUM(*m*-*i*, BLOCKSIZE))

i <- *i* + MINIMUM(*m*-*i*, BLOCKSIZE)

file_table[*file_index*].*cursor* <- *cursor*+*m*

return *m*

Mounting a file system

- Example: MOUNT("/dev/fd1", "/flash")
 - “grafts” directory tree of file system in physical device /dev/fd1 onto /flash
 - This association is kept in volatile memory; does not persist after shutdown/reboot
 - Recorded in in-memory inode for “flash” that a file system has been mounted on it
 - Virtual file system layer – more on Chapter 4
 - LOOKUP under “/flash” uses the root file system of the mounted device

Shells – Implied Contexts

- A part of the UNIX system is the “shell”
 - A command interpreter
 - Where you type commands such as ls, cd, cat, echo, ...
 - Commands can reference files
 - E.g. “ls” is an executable file
 - How does the shell resolve names?
 - Absolute path references
 - /bin/ls
 - Search paths
 - Environment variable determines the search path – e.g. PATH = /bin:/usr/bin:/home/usr/bin

Reading

- We will now skip to chapter 4
- Start reading 4.1, 4.2