

EEL-4736/5737
**Principles of Computer System
Design**

Lecture Slides 6
Textbook Chapter 4
Client-service Organization – Enforcing
Modularity


Introduction

- Modular design does not imply modularity is enforced
 - “Soft” modularity can be bypassed due to programming error, or malicious intent
- Let us study approaches that have been used to enforce modularity
 - Clients and services
 - Virtualization (ch 5)

An example

```
procedure MEASURE (func)  
  start <- GET_TIME (SECONDS)  
  func()  
  end <- GET_TIME (SECONDS)  
  return end-start
```

MEASURE(func)



```
procedure GET_TIME (units)  
  time <- CLOCK  
  time <- CONVERT_TO_UNITS (time, units)  
  return time
```

Modularity

- *Why?*
 - From the perspective of measurement, you want to measure a difference in seconds, and support multiple systems
 - Which units? Which syscall API?
 - Split measurement from clock interactions
 - GET_TIME() can be different depending on hardware
- *How?*
 - Two procedures run in the same application process: same register set, same memory
 - Procedure call stack convention implemented by compiler

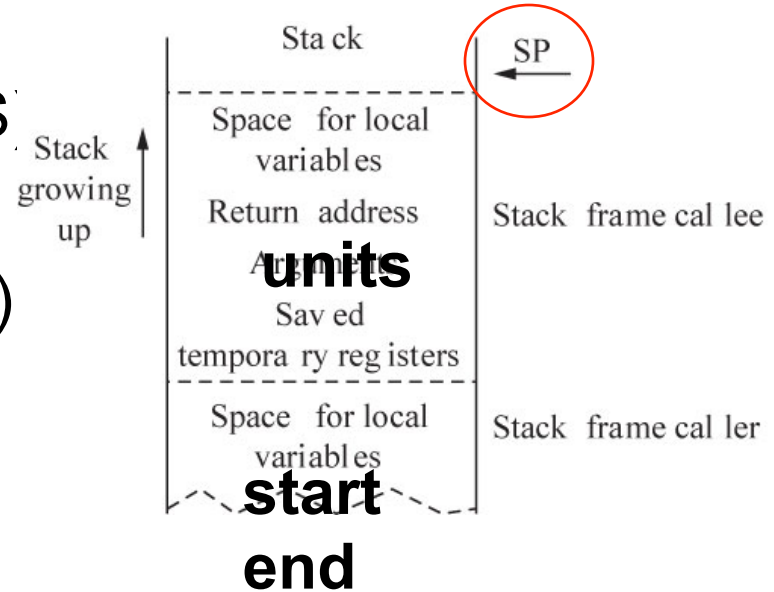
Stack discipline

- Typical way to support for procedure calls – stack memory abstraction
 - Push – add to top of stack
 - Pop – retrieve from top of stack
 - Different instruction sets implement this abstraction in different ways
 - PUSH/POP instructions
 - Stack pointers and loads/stores
 - Stack grow up or down
- At the end of invocation of a procedure
 - Should leave the stack as was found

An example

```
procedure MEASURE (func)  
  start <- GET_TIME (SECONDS  
  func())  
  end <- GET_TIME (SECONDS)  
  return end-start
```

```
procedure GET_TIME (units)  
  time <- CLOCK  
  time <- CONVERT_TO_UNITS (time, units)  
  return time
```



RISC-like machine code

MEASURE:

| | | |
|------|--------------|---|
| 100: | STORE R1, SP | Save registers that will use |
| 104: | ADD 4, SP | |
| 108: | STORE R2, SP | |
| 112: | ADD 4, SP | |

| | | |
|------|-----------------|----------------------------|
| 116: | MOV SECONDS, R1 | Argument for callee |
| 120: | STORE R1, SP | |
| 124: | ADD 4, SP | |


| | | |
|------|--------------|-----------------------|
| 128: | MOV 148, R1 | Return address |
| 132: | STORE R1, SP | |
| 136: | ADD 4, SP | |

| | | |
|------|-------------|----------------------|
| 140: | MOV 200, R1 | Call GET_TIME |
| 144: | JMP R1 | |

148: ...



Peeking into the stack



RISC-like machine code

GET_TIME:

| | | |
|------|-------------|---|
| 200: | MOV SP, R1 | Position R1 as index to read R2=argument |
| 204: | SUB 8, R1 | |
| 208: | LOAD R1, R2 | |

212: ... get_time body

| | | |
|------|--------------|--------------------------|
| 220: | MOV time, R0 | R0 = return value |
|------|--------------|--------------------------|

| | | |
|------|-------------|----------------------------|
| 224: | MOV SP, R1 | Load return address |
| 228: | SUB 4, R1 | |
| 232: | LOAD R1, PC | |

148: Back to MEASURE; adjust stack (sub 16),
restore registers R1, R2, use R0



Stack convention

- This example assumes that:
 - Callee leaves stack as it was found
 - Callee returns where caller told it to
 - Callee places return value in R0
 - Caller places values of temporary registers in stack not expecting callee to change them
- This is a **convention**
 - The compiler follows this convention when generating the machine code
- Suppose now that GET_TIME is linked to an object file

Stack convention

- This example assumes that:
 - Callee leaves stack as it was found
 - A mistake in GET_TIME may set the wrong value for SP
 - Cascades to MEASURE getting incorrect values for R1, R2; other procedures which called MEASURE?
 - Callee returns where caller told it to
 - A malicious programmer can try to return somewhere else so to gain control of sensitive information
 - Caller places values of temporary registers in stack not expecting callee to change them
 - Mistake may cause callee to overwrite stack and register values restored
- Caller will fail if convention is not followed
 - Specification; without enforcement

Other opportunities

- Errors can propagate when memory is improperly modified
 - E.g. a procedure that is passed a pointer to a data structure
 - Global variables
- *Procedures share the same memory name space and the same register set*
 - The same context – a process

Language and run-time support

- Languages have evolved to support techniques to ‘beef up’ soft modularity
 - “Strongly typed” – Java and C#
 - Language and runtime restricts that programs write only to memory locations that correspond to variables managed by the language
 - In accordance with their type
 - E.g. cannot use an Integer variable as an address of a memory location
- Still, many applications or modules written in languages that don’t enforce modularity



Client/Service Organization

- Limit interaction among modules to explicit messages
 - Contrast with the previous example, where both procedures are in the same memory name space and share registers

STORE R1,SP


R1

SP



R1

SP



Client/service organization

- Benefits
 - Isolation
 - Ability to run clients/services across two or more computers
 - Scalability, redundancy
- Disadvantage
 - Performance cost
 - Packing and sending messages has an overhead

Revisiting our example

procedure MEASURE (*func*)

 SEND_MESSAGE (*NameForTimeService*, {"Get Time",
 CONVERT2EXTERNAL(SECONDS)})

response <- RECEIVE_MESSAGE (*NameForClient*)

start <- CONVERT2INTERNAL (*response*)

func()

...

Revisiting our example

```
procedure TIME_SERVICE ()  
  do forever  
    request <- RECEIVE_MESSAGE (NameForTimeService)  
    opcode <- GET_OPCODE (request)  
    unit <- CONVERT2INTERNAL  
      (GET_ARGUMENT(request))  
    if opcode = "Get time" and (unit = SECONDS or unit =  
      MINUTES) then  
      time <- CONVERT_TO_UNITS (CLOCK, unit)  
      response <- {"OK", CONVERT2EXTERNAL(time)}  
    else  
      response <- {"Bad request"}  
    SEND_MESSAGE (NameForClient, response)
```

New things to worry about

- CONVERT2INTERNAL / EXTERNAL
 - Two computers may use different internal representations for data – big endian, little endian

| | | | | | | | | |
|-------|---------------------|---|-----|----------------|---------------------|---|-----|----------------|
| Words | 0 | | | | 1 | | | |
| Bytes | 0 | 1 | ... | 7 | 0 | 1 | ... | 7 |
| Bits | $2^0 2^1 2^2 \dots$ | | | $\dots 2^{63}$ | $2^0 2^1 2^2 \dots$ | | | $\dots 2^{63}$ |

| | | | | | | | | |
|-------|----------------|-----|---|---------------------|----------------|-----|---|---------------------|
| Words | n | | | | $n-1$ | | | |
| Bytes | 7 | ... | 1 | 0 | 7 | ... | 1 | 0 |
| Bits | $2^{63} \dots$ | | | $\dots 2^2 2^1 2^0$ | $2^{63} \dots$ | | | $\dots 2^2 2^1 2^0$ |

Send 0xABCDh, receive 0xDCBAh

New things to worry about

- *Marshaling*
 - Conversion of all data to be sent in a message into an array of bytes
 - Suitable for transmission and proper conversion to an object by the receiver
- Need to name and discover client and service endpoints
- Need to name and discover which services are exposed and their interface

Enforcing modularity

- Client/service don't rely on *shared state*
 - E.g. a stack or global variables
 - Failure in one does not directly corrupt data in the other
- Many errors cannot propagate
 - Client does not need to trust the service to return to appropriate return address
 - The service does not control that
 - Arguments can be checked when unmarshalled

Enforcing modularity

- Client can recover from service failure
 - Service goes into infinite loop?
 - Client can time out on waiting for response
 - Client may initiate a recovery process – e.g. try a backup, redundant service
 - Setting appropriate timers is not trivial
 - Added complexity that needs to be incorporated in the program (MEASURE example does not have it)

Summary

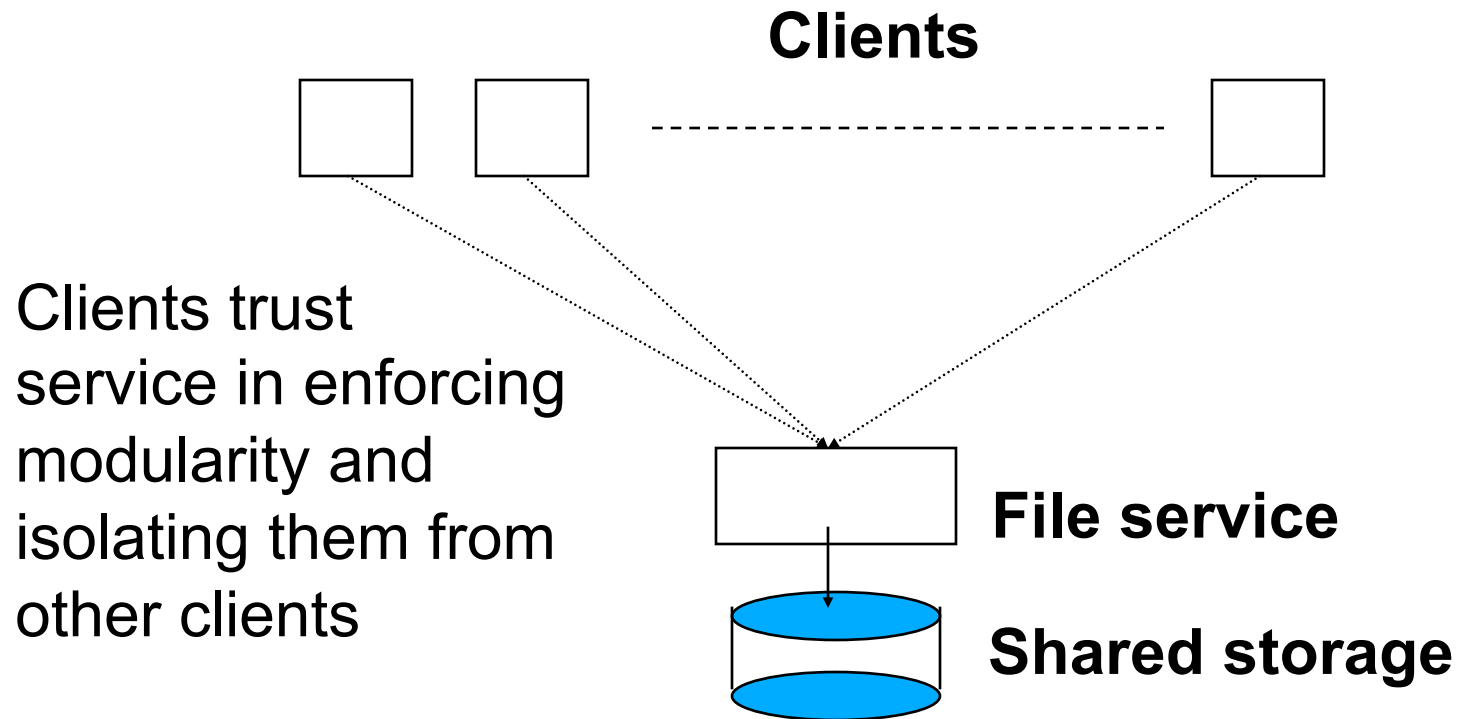
- Example of *sweeping simplification*
 - Remove all forms of interaction other than messages
- An improvement, but not a panacea
 - Service may allow messages to result in writes of value to any address in its space
- Trade-off
 - Shared data: ease of access within module
 - But, opportunity of error propagation
- Performance considerations put a bound on the size of modules

Additional flexibility

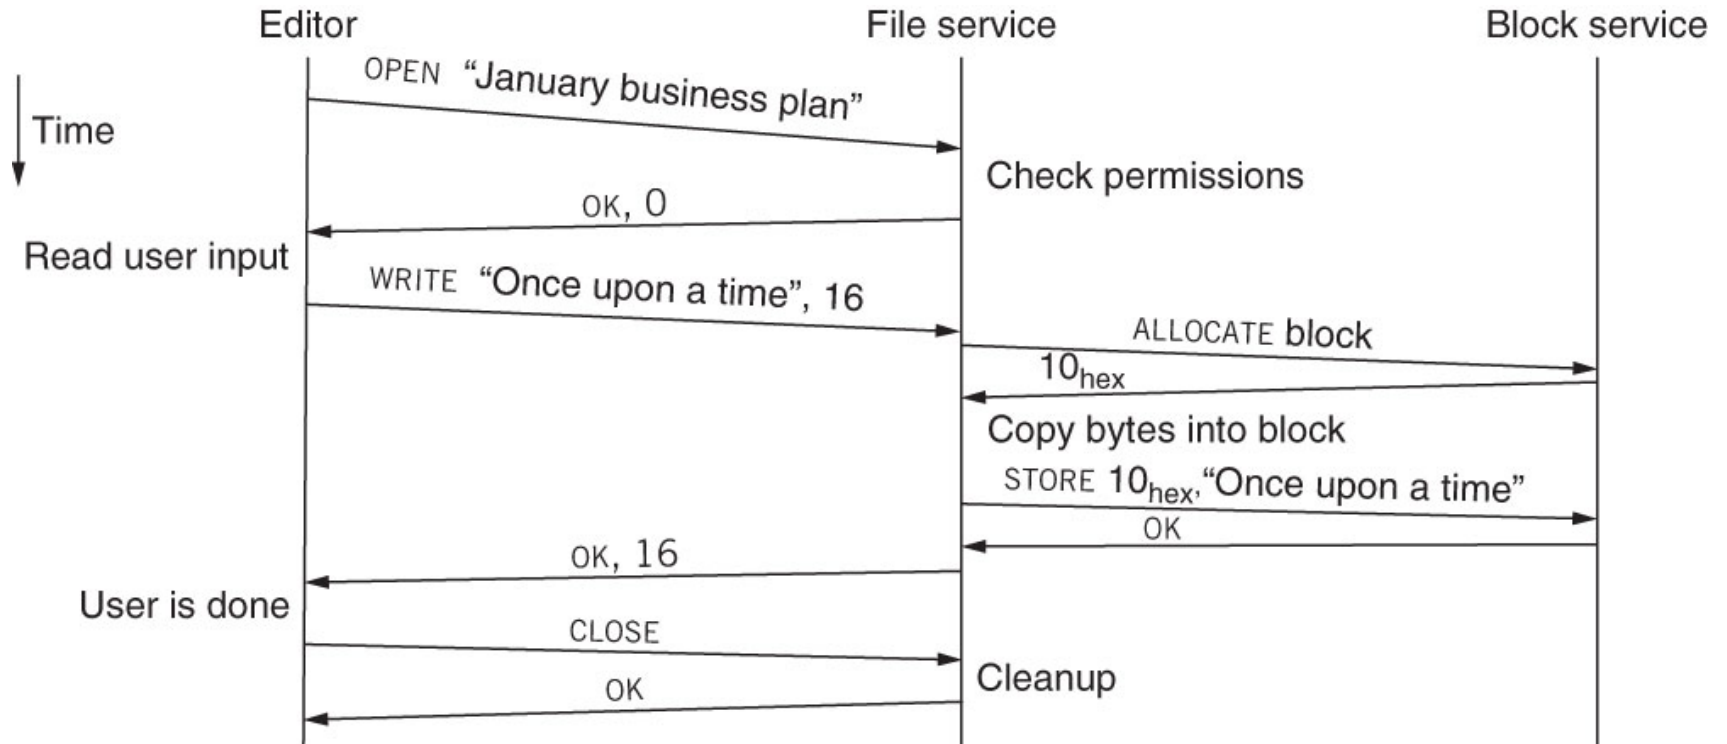
- One service can work for multiple clients
 - E.g. a printer or file server
- One client can use several services
 - E.g. a Web browser
- Single module can be both client and service
 - E.g. a proxy/cache

Trusted intermediaries

- Example: file service



File service



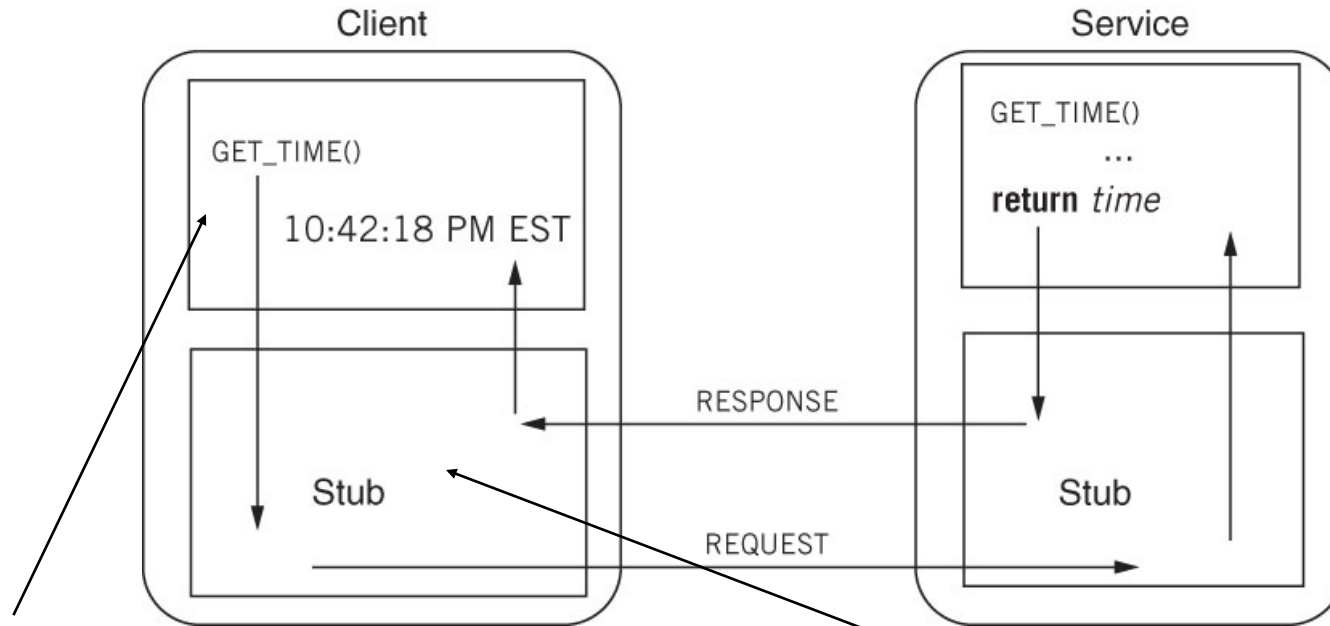
Trusted intermediaries

- There are general downsides
 - Vulnerability in the trusted intermediary can break the modularity enforced by the server
 - (Distributed) denial-of-service
 - Other clients bring service to a crawl
 - Vulnerability exploit
 - E.g. may gain access to other client's files
 - Must trust the intermediary
 - Eavesdropping? Privacy? Censorship?
- Peer-to-peer computing
 - Without intermediaries
 - Each peer a client and a service

Remote Procedure Call - RPC

- A typical client/service interaction has similar semantics to a procedure call
 - Client sends message to initiate a call to a named function with arguments
 - Service receives message, retrieves arguments, processes function
 - Service sends message with return value
 - Client receives message, retrieves return value
- Such a common pattern that it is useful to provide an abstraction and convenient APIs that hides implementation details

RPCs



procedure MEASURE (*func*)

start <- GET_TIME (SECONDS)

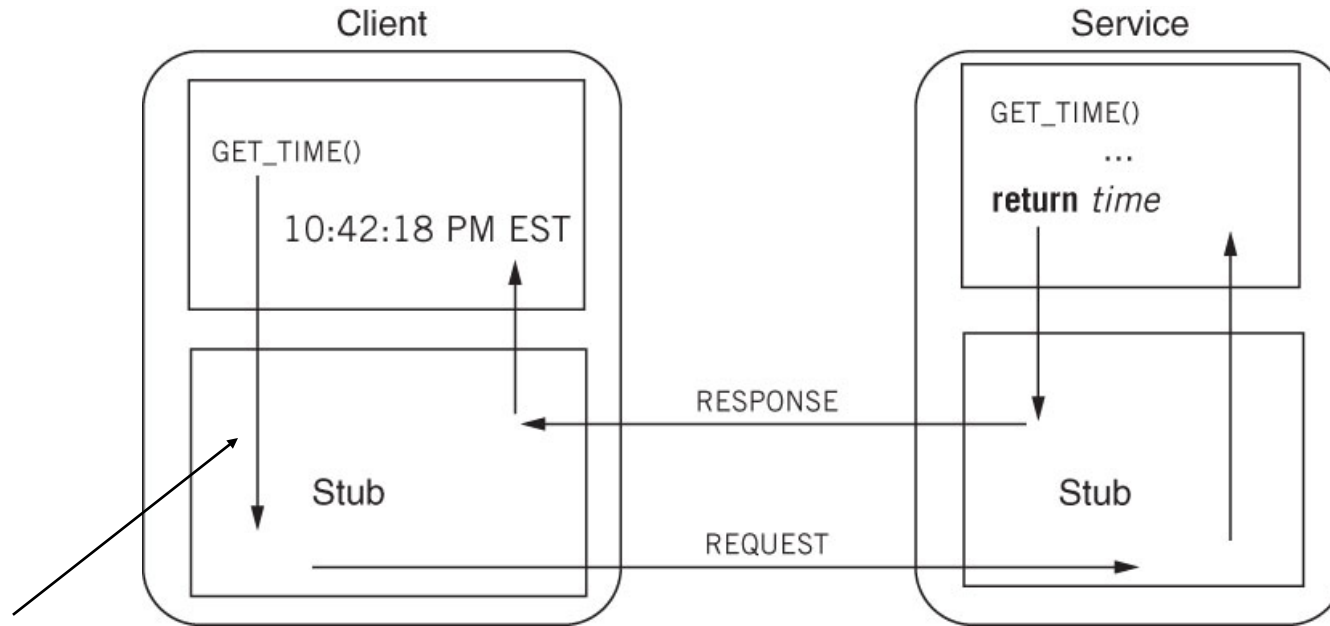
func()

end <- GET_TIME (SECONDS)

return *end-start*

Hides marshalling,
communication
details

RPCs



procedure `GET_TIME` (*unit*)

`SEND_MESSAGE` (`NameForTimeService`, {"Get Time", *unit*})

response \leftarrow `RECEIVE_MESSAGE` (`NameForClient`)

return *response*

Using RPCs

- Stubs can be generated automatically
- Need additional logic to handle faults

```
procedure MEASURE (func)  
  try  
    start <- GET_TIME(SECONDS)  
  catch (signal servicefailed)  
    return servicefailed  
  
  func ()  
  try  
    end <- GET_TIME(SECONDS)  
  catch (signal servicefailed)  
    return servicefailed  
  
  return end - start
```

RPC != PC

- RPCs take more time than PCs
- RPCs can reduce “fate sharing”
 - Callee failure causes caller to also fail
- RPCs introduce new failures that are not present in procedure calls
 - Service failure
 - *No response, timeout*
 - How to handle this? Not possible to determine if no response happened before or after service performed the action

Dealing with no-response

- At-least-once RPC:
 - Client stub re-sends as many times as needed until it gets a response
 - Client stub may still give up, after multiple re-tries – cannot provide guarantee implied by name
 - Service needs to be ready to execute many times the same request
 - Read? Write? Append?
 - Side-effect-free / idempotent operations

Dealing with no-response

- At-most-once RPC:
 - Client tries once; if no response, stub returns error to caller
 - More appropriate for requests with side effects
 - E.g. RPC to transfer funds from one account to another – at-most-once to avoid undesired transfers
 - Implementation needs to account for the underlying network – which may duplicate messages

Dealing with no-response

- Exactly-once RPC
 - Ideal semantics, but in principle impossible to guarantee if client and service are independent
 - What if service goes away forever?
 - Techniques can be used to implement an approximation of this semantics under certain assumptions

Reading

- Sections 4.2, 4.3, 4.5