

**EEL-4736/5737**  
**Principles of Computer System  
Design**


Lecture Slides 11  
Textbook Chapter 5  
Virtual Memory

# Introduction

- Memory domains restrict management
  - How to grow them if memory is limited?
    - Copying is expensive
    - Re-locating and re-computing address is cumbersome
- Virtual memory adds a layer of abstraction for better management
  - Applications use virtual addresses as names in the context of its own virtual address space
  - Bus sees physical address names
  - Re-mapping vs. copy/relocation


# Virtualizing addresses

- Threads use *virtual addresses*
- Bus sees *physical addresses*
- Memory manager: *permission checks and translation*
  - Decouple modules with indirection



# Example

- Large set of virtual addresses ( $2^{64}$ ) but a smaller set of physical addresses
- Still use domains, with virtual addresses
- What if need domain 1 to grow?



# Domains with virtual address


- Pros:
  - Relocation is application-transparent
- Cons:
  - Translation is in the path of every memory reference
  - Domain growth may still incur copying
  - Fragmentation
  - Need additional modules to implement memory manager functions

# Using a page map

- Mapping an entire domain
  - Coarse grain: fragmentation, expensive copy
- Mapping individual addresses
  - Possible conceptually, but keeping track of mapping would require huge tables
- Design approach: page maps
  - An array of page map entries
  - Each entry maps for a *fixed-size range of consecutive bytes of virtual addresses* (page) to a range of physical addresses (block, or frame)

# Page map

- Virtual addresses now is a name overloaded with a structure to translate on a per-page basis
  - Upper bits: page number
  - Lower bits: page offset




# Memory management

- One page map per allocated page
- Free to allocate/move any block in physical memory to hold a page
  - Insert/update appropriate map
  - Transparent to the thread
  - Unit of management: same size blocks instead of varying-size domains



# Implementing page maps

- Page tables - only one field of the virtual address needs translation



# Page tables

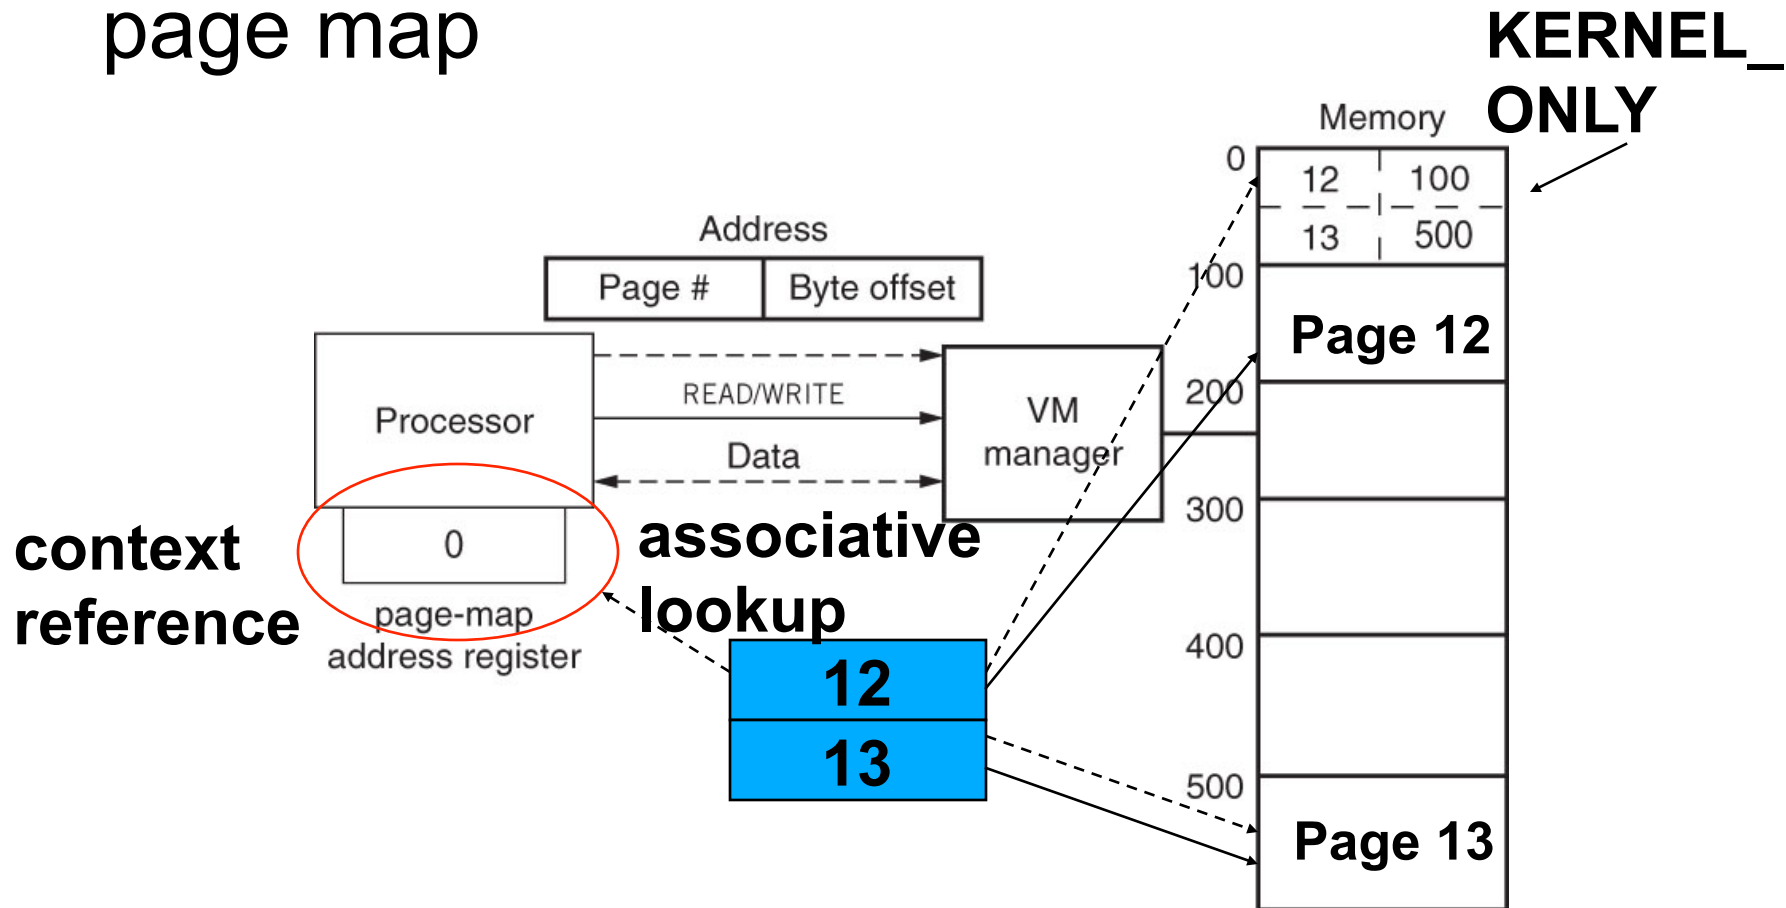
- Mapping table size still a concern
  - 4KBytes per page, 64-bit virtual addresses
    - $2^{52} * 52$  bits at a minimum
      - 52: 64 bits – 12 offset
  - Can be addressed with techniques such as multi-level page tables
  - We'll assume for simplicity a single-level table

# Implementing page tables

- Even with techniques to reduce size, page tables are larger than what's typically available within a processor
- Hence, they are themselves stored in main memory
- Implications:
  - Need to be able to locate the page table
  - Need to protect the page table from non-kernel threads
  - Performance optimizations are needed

# Implementing page maps

- Example – assume processor has a register pointing to the address of the page map



# Virtual address spaces

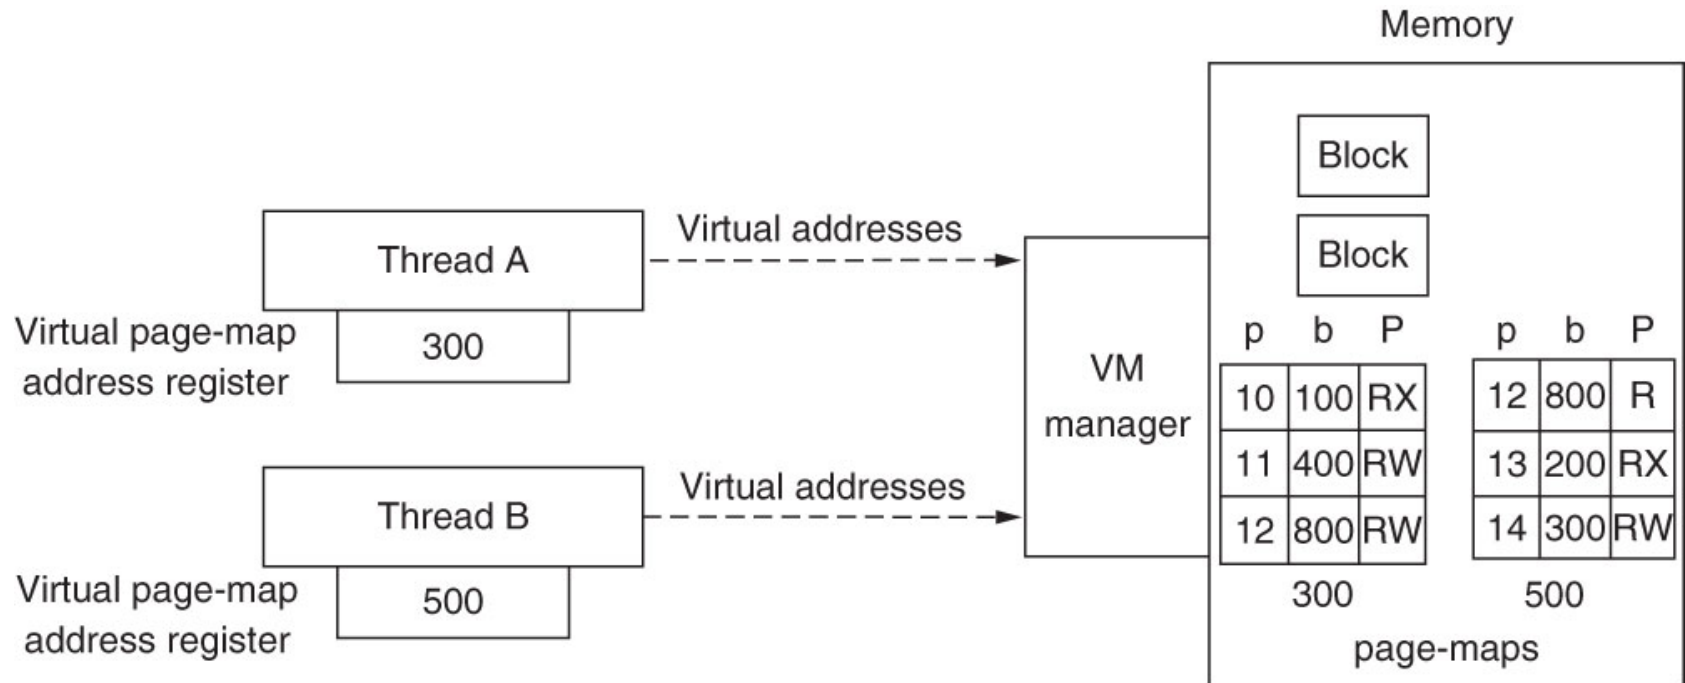
- A single, large set of virtual addresses could be shared by all threads
  - Practicality depends on address bits
    - E.g. 64-bit virtual addresses; might have 4 billion threads each using 4GB; 32-bit – 8 512MB threads?
    - Another problem: relocation needed depending on base address allocated
    - E.g. multiple instances of the same program
- Can eliminate this assumption with per-application virtual address spaces

# Virtual address spaces

- An application has the illusion that it has a complete address space to itself
  - 32-bit virtual address space: *each* thread can use  $2^{32}$  virtual addresses
- Key feature to support this:
  - *Each thread has its own page map*
- Often, operating system *process* is used to refer to 1+ threads that share a virtual address space

# Virtual address spaces

- Virtual memory manager now keeps a map per virtual address space



# Address translation

**procedure** TRANSLATE (integer virtual, perm\_needed)  
**returns** physical

page <- virtual[0..19]    // 32-bit address space

offset <- virtual[20-31]    // 4KB (12-bit offset) pages

page\_table <- PMAR

perm\_page <- page\_table[page].permissions

**if** PERMITTED (perm\_needed, perm\_page)

    block <- page\_table[page].address

    physical <- block + offset

**return** physical

**else return** error



# Address translation

**procedure** TRANSLATE (integer virtual, perm\_needed) **returns** physical

page <- virtual[0..19] // 32-bit address space

offset <- virtual[20-31] // 4KB (12-bit offset) pages

page\_table <- PMAR

perm\_page <- page\_table[page].permissions

**if** PERMITTED (perm\_needed, perm\_page)

block <- page\_table[page].address

physical <- block + offset

**return** physical

**else return** error

perm\_needed

LOAD: R


STORE: W

Exec: X

Kernel: K

**LOAD 0xF0F01020, R1**

page  
**F0F01**



page

offset

page\_table

AB123	RW

address


permissions

**return: 0xAB123020**

# Example

- What happens if:
  - A issues a read reference to VA within page 10?
  - B issues a read reference to VA within page 10?
  - A issues a write reference to VA within page 12?
  - B issues a write reference to VA within page 12?

**What module programs this?**  
**When does its value change?**



# Primitives

- `id <- CREATE_ADDRESS_SPACE()`
- `block <- ALLOCATE_BLOCK()`
- `MAP (id, block, page_number, permission)`
- `UNMAP (id, page_number)`
- `FREE_BLOCK (block)`
- `DELETE_ADDRESS_SPACE(id)`

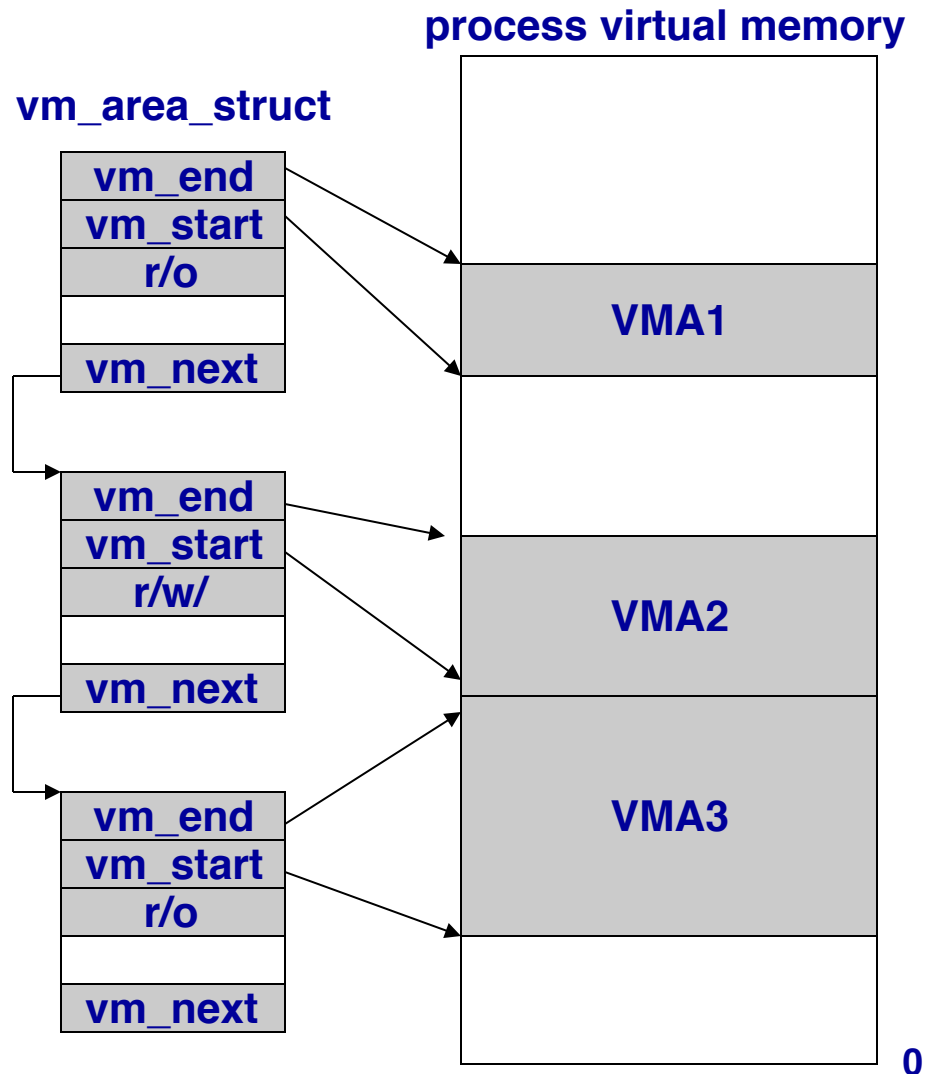
# SW/HW interface

- Virtual memory manager in the kernel
  - Allocate/map blocks
  - Maintain data structures for each virtual address space in memory
  - Program hardware structures to put appropriate address space “in context”
- Memory management unit in hardware
  - Translation/permission checks for individual memory accesses, at high speed

# Example: Linux VM “areas”

- Linked `vm_area_struct` structures
- A VM area: a part of the process virtual memory space that has a special rule for the page-fault handlers (i.e. a shared library, the executable area etc).
- These are specified in a `vm_area_struct`
  - Start and end VM address of area
  - Protection information, flags

# Linux fault handling

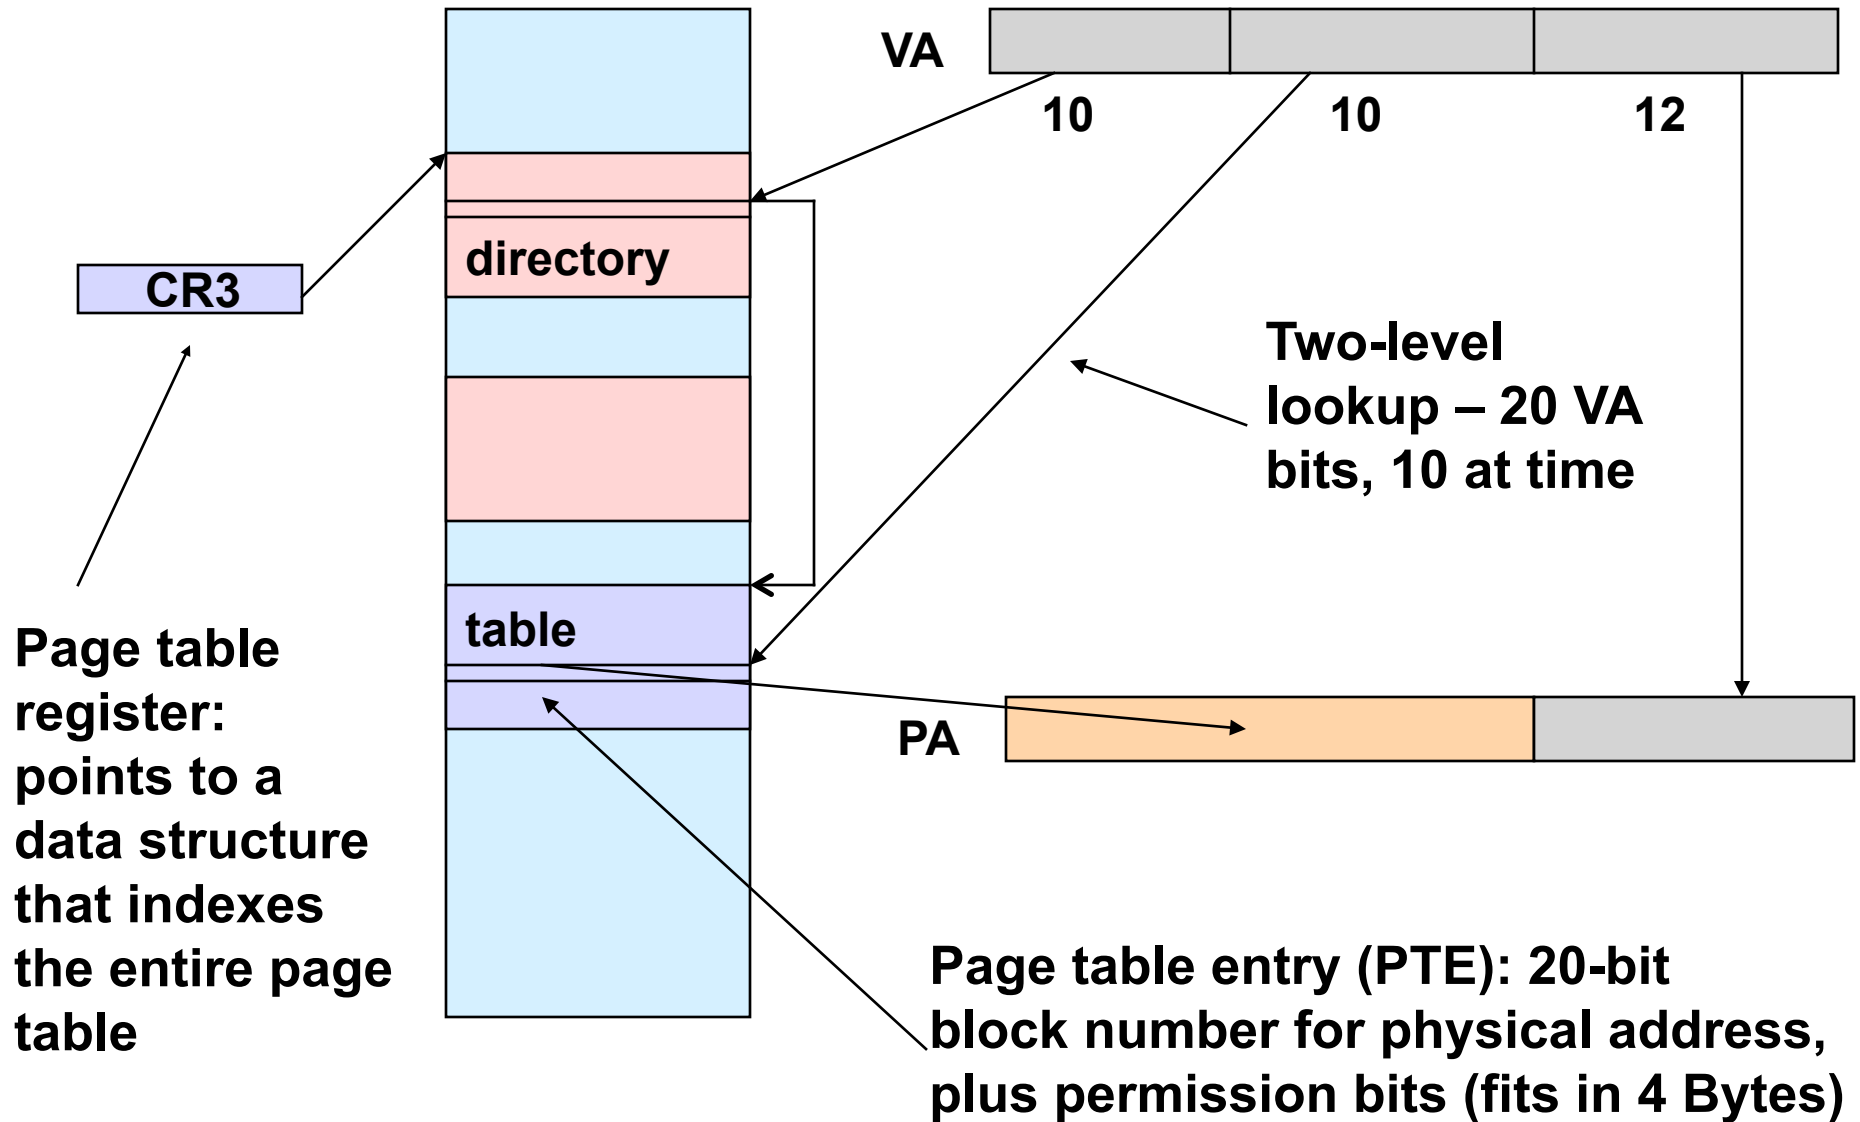


- Exception triggers O/S handling if address out of bounds or protection violated
- Traverse vm\_area list, check for bounds
  - If not mapped, it is a segmentation violation – signal to process
- If mapped, check protection of vm\_area\_struct
  - E.g. r/o, r/w
  - Signal protection violation to process if access not allowed
  - Otherwise, handle fault and bring page to memory

# Kernel and address spaces

- Can separate the kernel and user address spaces in different ways
  - They share the same address space, but kernel pages are `KERNEL_ONLY`
    - E.g. top bit of virtual address divides what's available to kernel vs. user
  - Or, kernel has different page map altogether
    - SVC instruction and return to user mode must also change the page map register

# E.g.: x86 2-level page tables



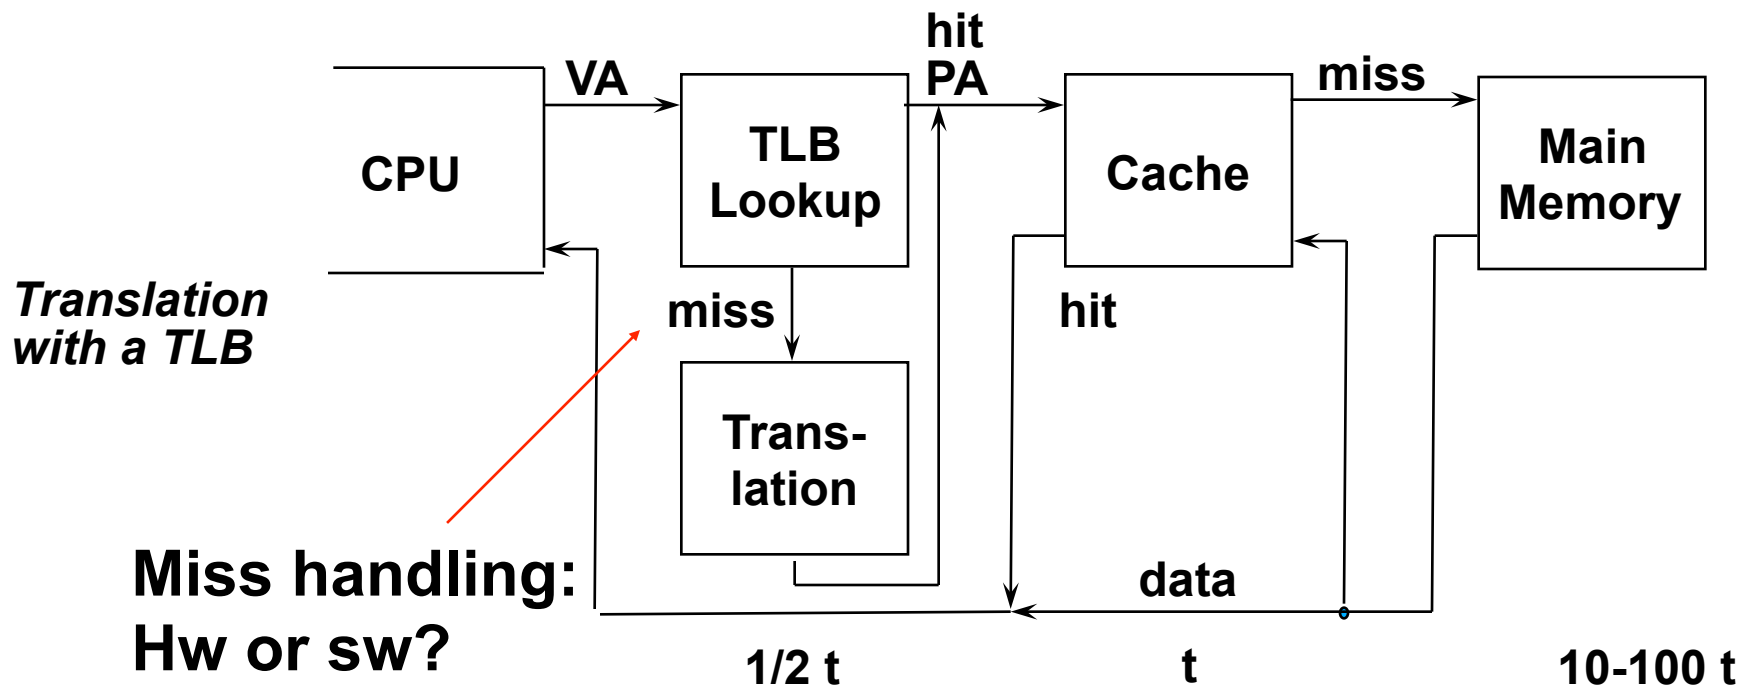


# HW/SW Tradeoffs

- Per-access translation needs to be done in hardware for high speed
  - Two memory lookups per access – too expensive
  - Translation look-aside buffer (TLB) – a specialized cache for page table entries

# TLBs

- Caches often indexed by physical address
- TLBs cache translations so they can be computed before going into the cache
  - Typically, set-associative



# HW/SW Tradeoffs

- How are TLB misses handled? Two approaches:
  - Software-managed TLB: raise exception and let software manager (kernel) handle it
    - MIPS, and other RISC-style processors
    - Pro: flexibility in data structures and algorithms to replace TLB entries
  - Hardware-managed TLB: hardware manager looks up the page table and handles miss without software involvement
    - x86 is a prime example
    - Pro: fast TLB miss handling

# Mapping pages to storage

- Additional flexibility given by virtual addresses
  - *page mappings do not need to be always to main memory*
  - A page may map to a block of a storage subsystem (e.g. hard disk)
- Benefit:
  - Allows the available virtual memory to exceed the main memory capacity
- Challenge:
  - Must be careful with performance – disk is orders of magnitude slower

# Mapping pages to storage

- Overall approach:
  - References to virtual addresses of pages that are mapped to memory translated by processor hardware (TLB, page table)
  - Pages mapped to storage blocks are flagged as not mapped, such that an exception is raised (*page fault*)
  - System software (kernel) responsible for bringing page from/to disk into memory, and letting the processor retry

# Reading

- Sections 5.5-5.6