

**EEL-4736/5737**  
**Principles of Computer System  
Design**

Lecture Slides 12  
Textbook Chapter 5  
Virtual Processors Using Threads

# Introduction

- Assumption so far
  - Each thread has own separate processor
- We have seen the importance of
  - privilege separation
  - the role of the kernel as a manager of shared resources
- We will now focus on how to share a processor by multiple threads

# Basic thread management

- Allocating and terminating threads
- Thread\_id <- ALLOCATE\_THREAD  
(starting\_procedure, address\_space\_id)
  - Allocate a new thread in given address space, start execution at given address
  - Return identifier that names this thread
- Thread manager:
  - Allocate a stack for thread, set SP
  - Assign a processor for execution, set its PC to the starting address

# Virtualizing the processor

- Time-sharing/multiplexing/multi-tasking
- Many threads are long-running
- At many instances, threads are waiting for events
  - Spin loop in SEND/RECEIVE
  - Waiting for a slow device to reply
- Inefficient to hog the processor up
- A primitive to *yield* the processor to another thread allows for time-sharing

# Yielding

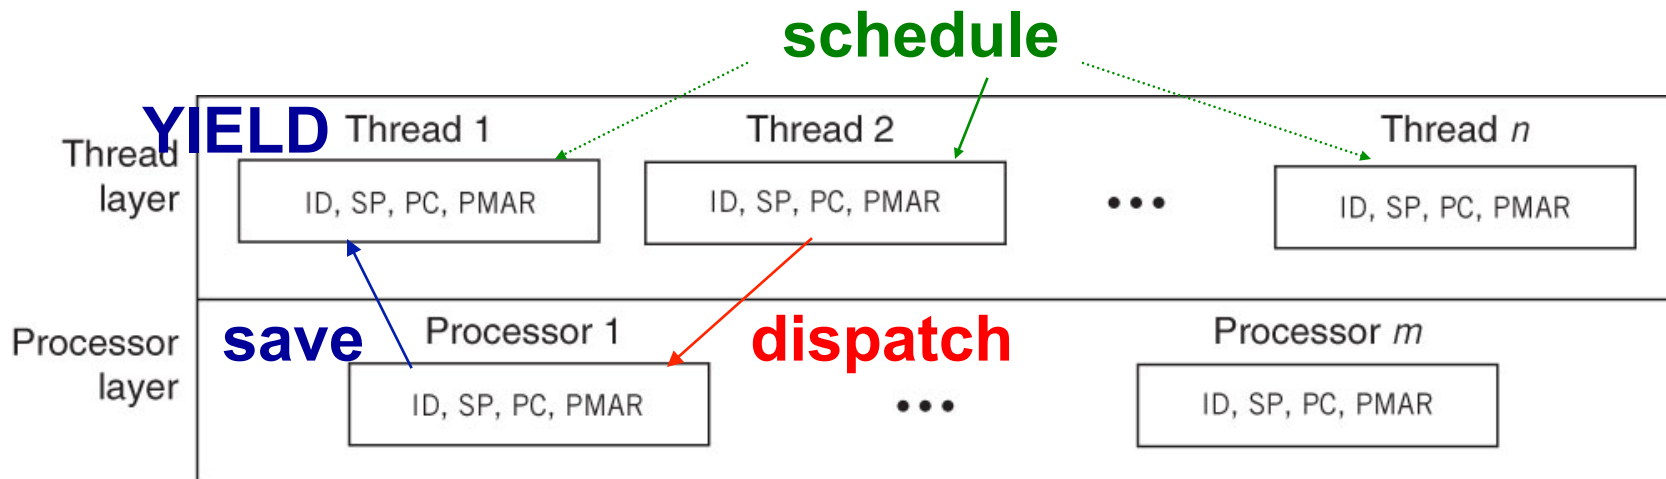
- Allow threads to call a procedure YIELD
  - An entry to the thread manager
  - Allows the thread manager to give the processor to a different thread
- Manager keeps thread state
  - Running
  - Runnable
- YIELD
  - Save yielding thread's state
  - Decide which other thread will run on the processor
  - Dispatch processor to chosen thread

# Spin loop with yield

```
procedure SEND (buffer reference  $p$ , message  
  instance  $msg$ )  
  ACQUIRE ( $p.buffer\_lock$ )  
  while  $p.in - p.out = N$  do  
    RELEASE ( $p.buffer\_lock$ )  
    YIELD ()  
    ACQUIRE ( $p.buffer\_lock$ )  
     $p.message[p.in \bmod N] \leftarrow msg$   
     $p.in \leq p.in + 1$   
  RELEASE ( $p.buffer\_lock$ )
```

# YIELD

- ID (identifier), SP (stack pointer), PC (program counter), PMAR (page map address register)



# YIELD implementation

- Requires processor support to implement the saving/restoring of state
  - In addition to the core registers from previous slide
  - Different instruction set architectures, different implementations - register by register, or all at once
- Implementation of YIELD belongs in the kernel, exposed through supervisor calls



# Interrupts and exceptions

- Interrupts occur when an event that is unrelated to the running thread occurs
  - E.g. an I/O device interrupt
  - Note, the thread may eventually use information related to interrupt (e.g. a keyboard character) but it does not know how to handle the interrupt directly
- Exceptions are events that pertain specifically to the running thread
  - E.g. divide by zero, access out of bounds
- “Fault”, “trap”, “signal” are terms also used

# YIELD version 1.0

- Assumptions
- Fixed number of threads (larger than number of processors)
- Assume they all run in the same address space
- Assume all threads are already running

# Shared data structures

**shared structure** processor\_table[2]

**integer** thread\_id // id of running  
thread

**shared structure** thread\_table[7]

**integer** topstack // stack pointer

**integer** state // runnable/running

**shared lock instance** *thread\_table\_lock*

# Example



**Thread table**



**Processor table**



**Thread\_table\_lock**



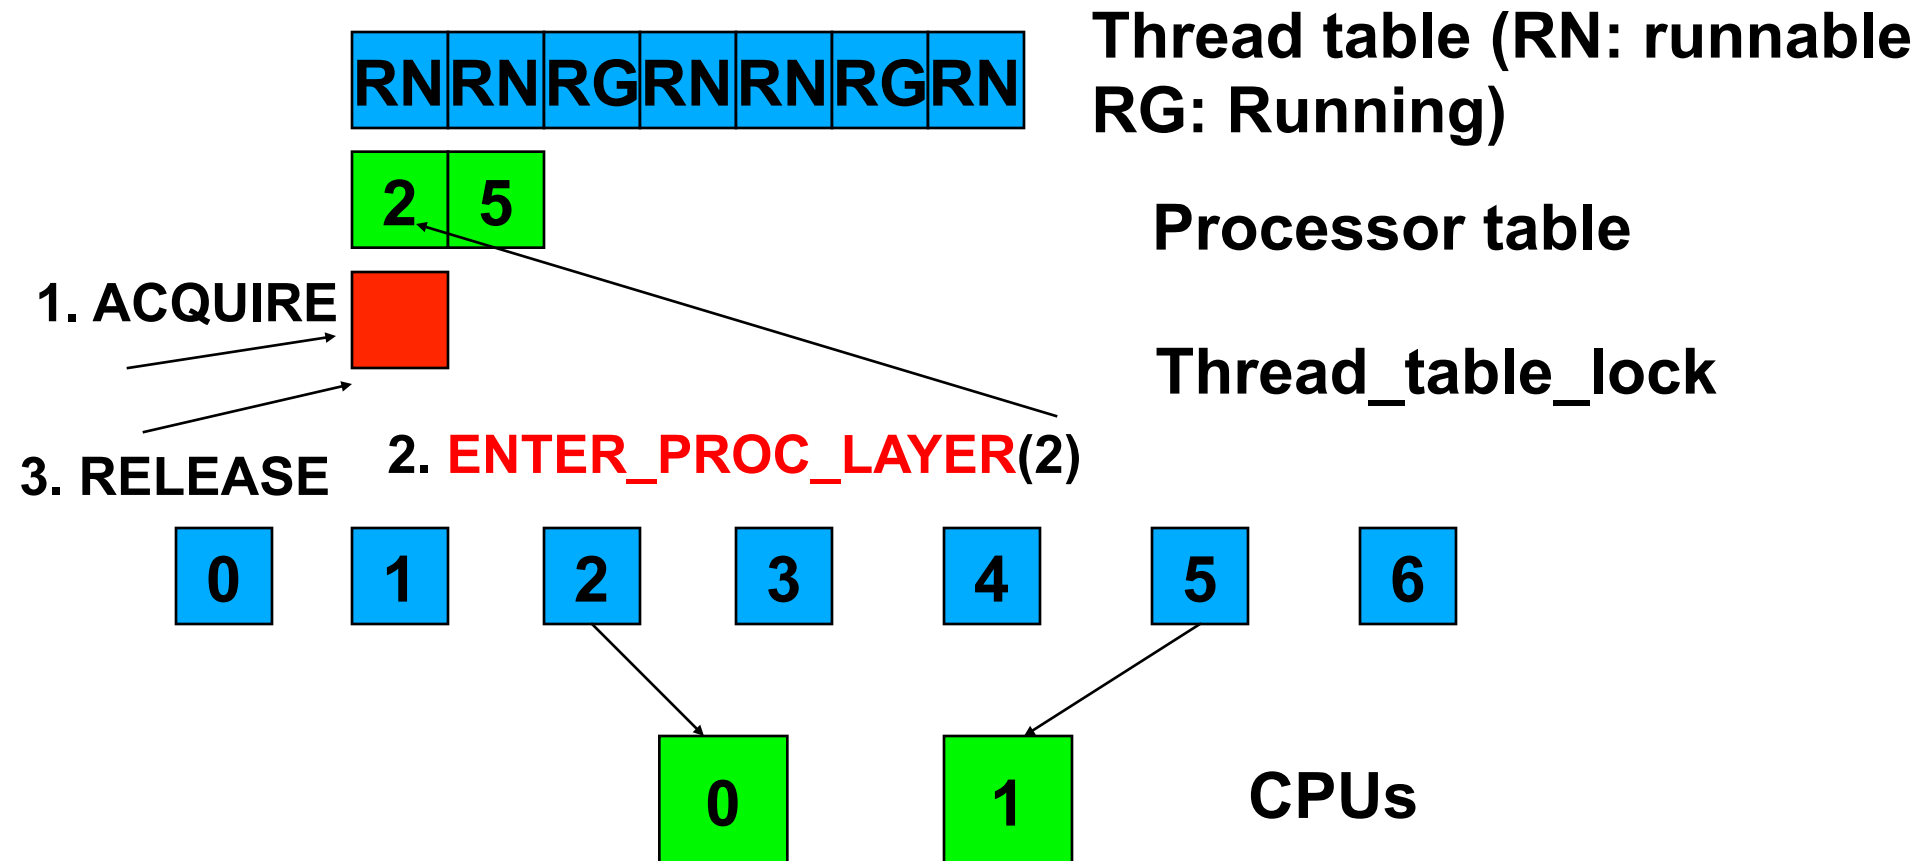
**CPUs**

# Main YIELD body

```
procedure GET_THREAD_ID() return  
    processor_table[CPUID].thread_id  
    // which thread is running on this CPU?  
procedure YIELD()  
    ACQUIRE (thread_table_lock)  
    ENTER_PROCESSOR_LAYER  
        (GET_THREAD_ID())  
    RELEASE (thread_table_lock)  
return
```

# Example

- Thread 2 YIELDS



# Unfolding ENTER\_PROC\_LAYER

```
procedure ENTER_PROC_LAYER(this_thread)  
  thread_table[this_thread].state <- RUNNABLE  
  thread_table[this_thread].topstack <- SP  
  SCHEDULER()  
return
```

```
procedure SCHEDULER()  
  j = GET_THREAD_ID()  
  do j <- (j+1) modulo 7  
  while thread_table[j].state != RUNNABLE  
  thread_table[j].state <- RUNNING  
  processor_table[CPUID].thread_id <- j  
  EXIT_PROCESSOR_LAYER(j)  
return
```

# Example

- Thread 2 YIELDS

2.1 Change to runnable; save SP

2.4  
Set proc.  
table

RNRNRNRGRNRGRNR

Thread table (RN: runnable  
RG: Running)

3 5

2.2 select j=3

Processor table

2.3 set as running

1. ACQUIRE

Thread\_table\_lock

2. ENTER\_PROC\_LAYER(2)

3. RELEASE

0

1

2

3

4

5

6

0

1

CPUs



# Exiting processor layer

```
procedure EXIT_PROCESSOR_LAYER(new)  
    SP <- thread_table[new].topstack  
    return
```

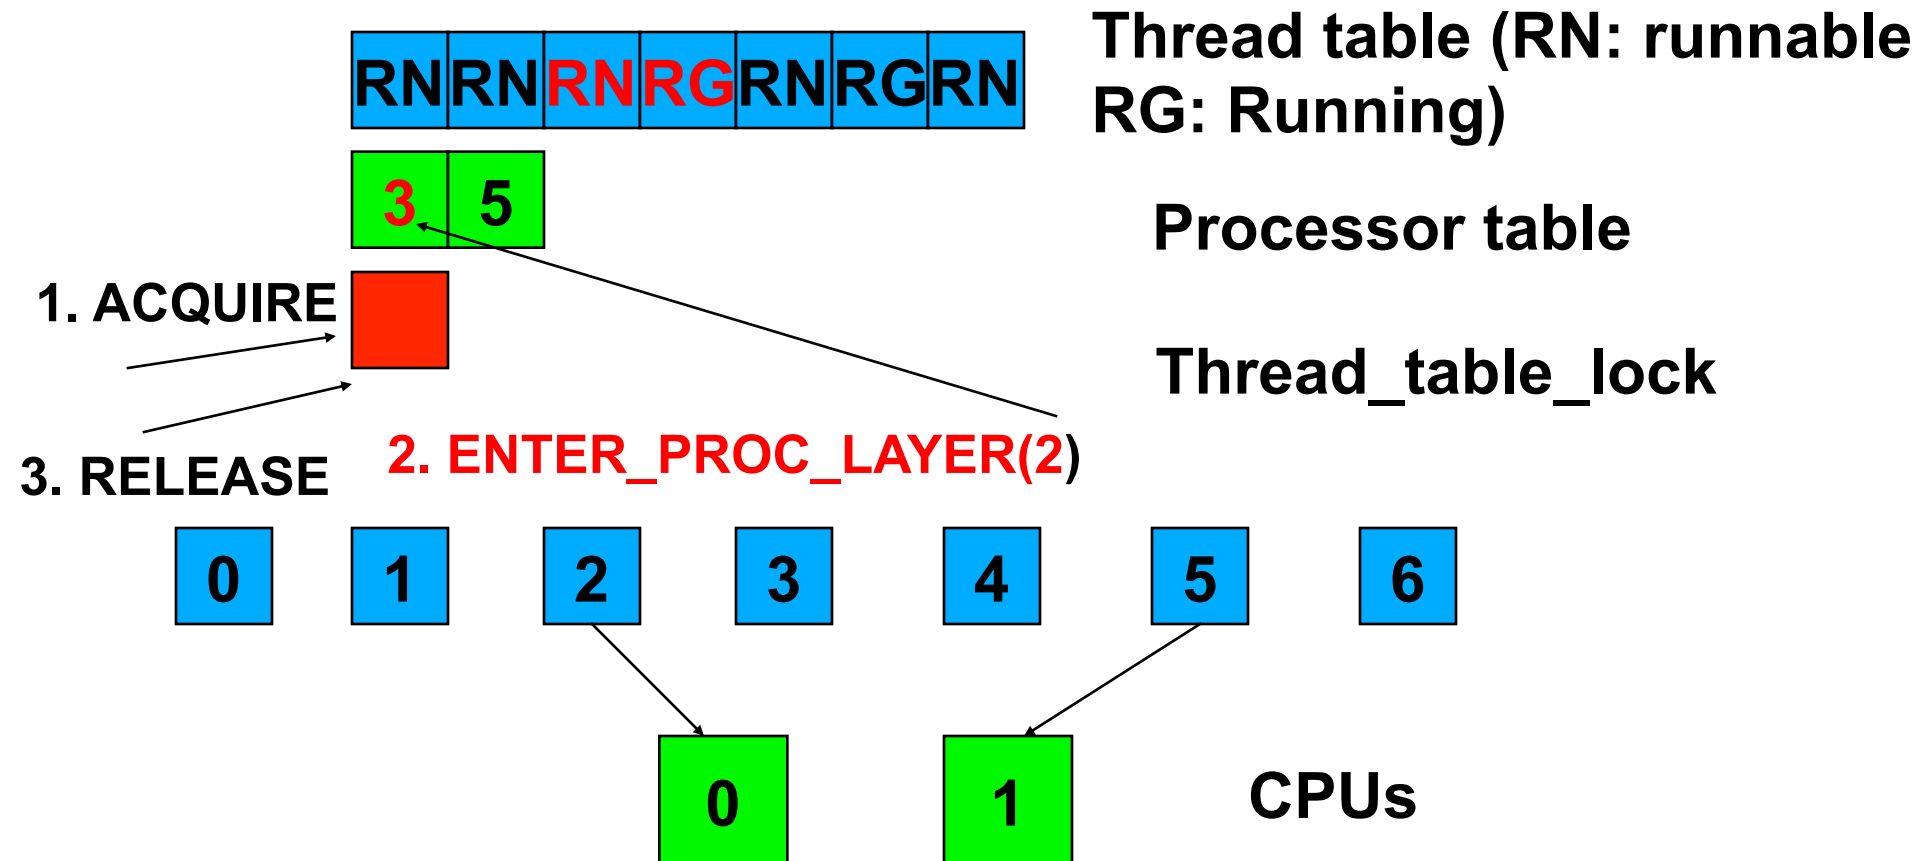
What happens when this return is called?

# Example

Thread ID 2 called YIELD

Threads assumed to have been running, so thread ID 3 had called YIELD previously

We want thread ID 3 to resume executing right after its ENTER



# Exiting processor layer

```
procedure EXIT_PROCESSOR_LAYER(new)  
    SP <- thread_table[new].topstack  
return
```

This **return** pops a return address **RA** from thread ID *new* (*new*=3 in our example)

**RA** is the return address that was pushed into the stack when thread ID 3 called ENTER\_PROCESSOR\_LAYER (3)

# Revisiting ENTER\_PROC\_LAYER

```
procedure YIELD()  
  ACQUIRE (thread_table_lock)  
  ENTER_PROCESSOR_LAYER (3)  
  RELEASE (thread_table_lock)  
return
```

```
procedure ENTER_PROC_LAYER(this_thread)  
  thread_table[this_thread].state <- RUNNABLE  
  thread_table[this_thread].topstack <- SP  
  SCHEDULER()  
return
```

When thread ID 3 called ENTER\_PROC\_LAYER,  
its **SP got saved** with **this return address RA**

# Control flow in our example

```
procedure YIELD()  
  ACQUIRE (thread_table_lock)  
  ENTER_PROCESSOR_LAYER (2)  
  RELEASE (thread_table_lock)  
return
```

```
procedure EXIT_PROCESSOR_LAYER(new)
```

```
  SP <- thread_table[new].topstack ← loads thread 3's saved stack into SP  
  return
```

```
procedure YIELD()  
  ACQUIRE (thread_table_lock)  
  ENTER_PROCESSOR_LAYER (3)  
  RELEASE (thread_table_lock)  
return
```

**Thread 2  
Acquires lock**

**Lock  
is passed  
between  
threads**

**Thread 3  
Releases lock**

RA from  
3's stack

# Thread creation/termination

- `id <- ALLOCATE_THREAD`
- `EXIT_THREAD ()`
- `DESTROY_THREAD (id)`
- Possible to have fewer threads than processors
- Create a “processor-layer thread” for each processor, which runs the SCHEDULER procedure
  - Multiple “Thread-layer” threads

# Processor-layer thread

- Gets created at each processor during kernel initialization
- Switching from a thread-layer thread to another thread-layer thread
  - Requires going through the processor-layer thread

# Initializing processor thread

**procedure** RUN\_PROCESSORS()


**for each** processor **do**

(allocate stack and set up processor thread)

shutdown <- FALSE

**SCHEDULER ()**

**Will only return from  
scheduler when  
shutdown=TRUE**



(deallocate processor thread stack)

(halt processor)



# Extended data structures

**shared structure** processor\_table[2]

**integer** topstack // stack pointer

**byte reference** stack // allocated stack (boot time)

**integer** thread\_id // id of running thread

**shared structure** thread\_table[7]

**integer** topstack // stack pointer

**integer** state // runnable/running/**free**

**boolean** kill\_or\_continue // initialized as “continue”

**byte reference** stack // allocated stack (create time)

**shared lock instance** thread\_table\_lock

# Extended YIELD

**procedure** YIELD()

ACQUIRE (*thread\_table\_lock*)

ENTER\_PROCESSOR\_LAYER

(GET\_THREAD\_ID(), **CPUID**)

RELEASE (*thread\_table\_lock*)


**return**



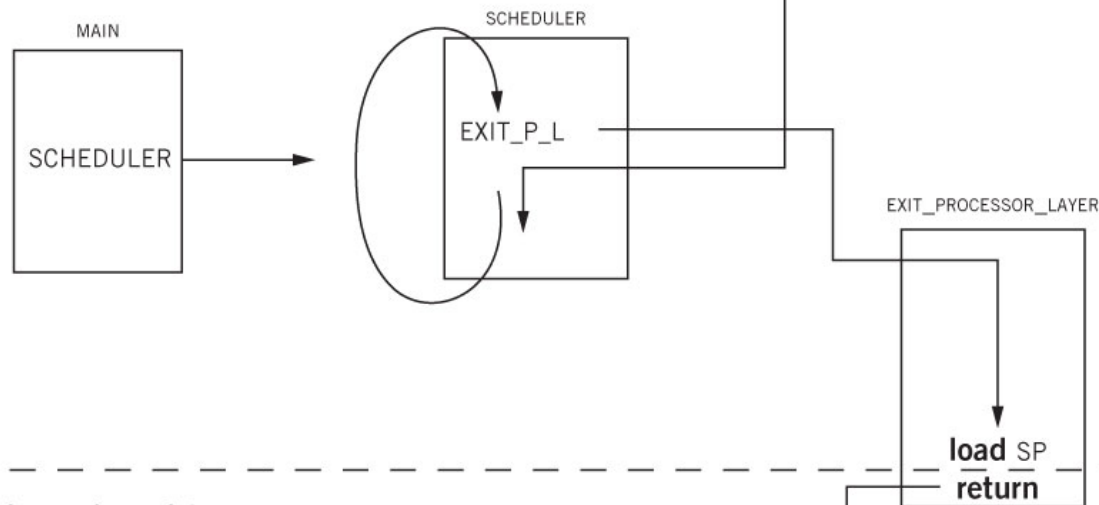
**Set up to switch  
Into processor-layer  
thread**

# Example – 0 yields to 6

User layer thread 0



Processor layer thread



User layer thread 6

# Entering processor layer

**procedure**

ENTER\_PROC\_LAYER(*tid*,*procid*)

*thread\_table*[*tid*].*state* <- RUNNABLE

*thread\_table*[*tid*].*topstack* <- SP

*SP* <- *processor\_table*[*procid*].*topstack*

~~SCHEDULER()~~

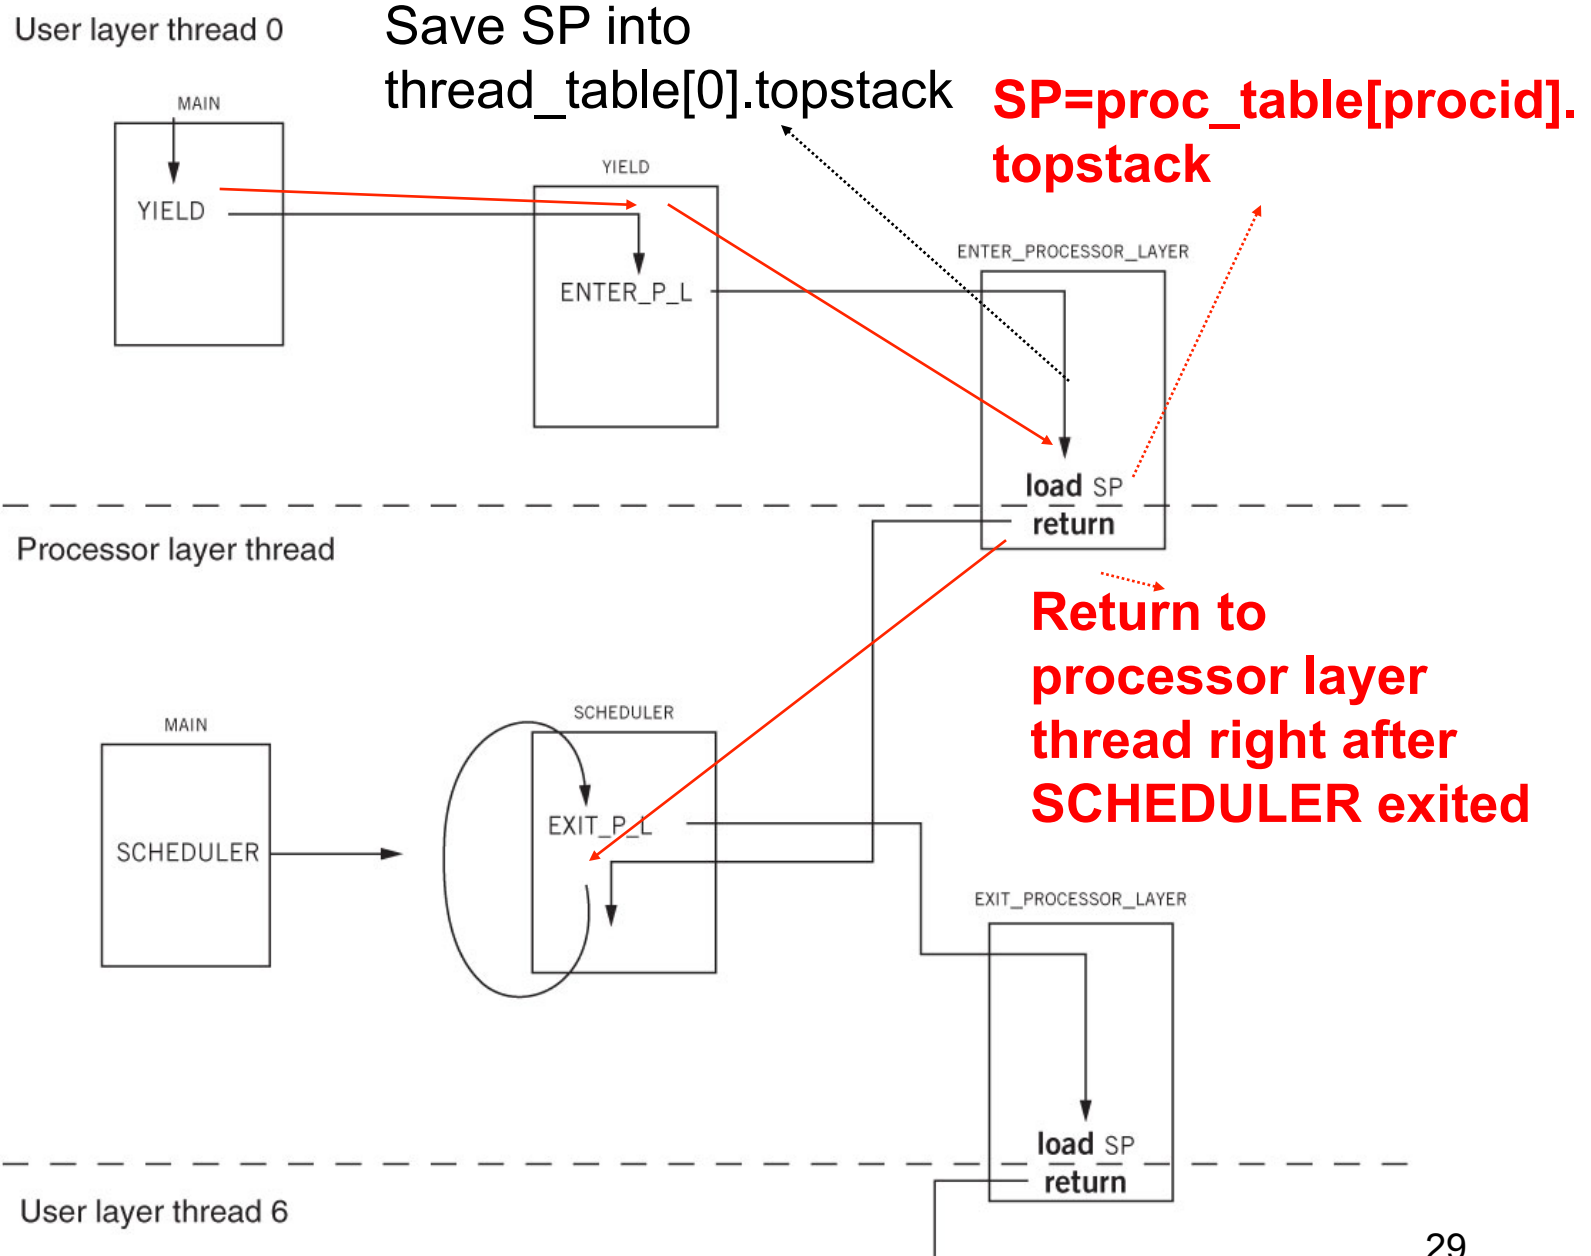
**return**

Enter processor  
layer

No need  
for this call  
anymore

Set up stack so  
Return will take  
It to processor thread  
(i.e. SCHEDULER)

# Switching threads

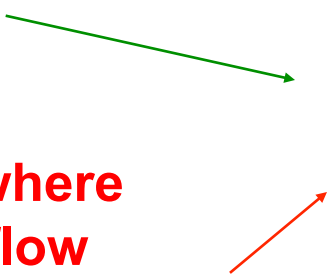


# Processor thread - SCHEDULER

```
procedure SCHEDULER()
  while shutdown = FALSE do
    ACQUIRE (thread_table_lock)
    for i from 0 until 7 do
      if thread_table[i].state = RUNNABLE then
        thread_table[i].state <- RUNNING
        processor_table[CPUID].thread_id <- i
        EXIT_PROCESSOR_LAYER(CPUID, i)
        if thread_table[i].kill_or_continue = KILL
          thread_table[i].state <- FREE
          DEALLOCATE(thread_table[i].stack)
          thread_table[i].kill_or_continue = CONT
        RELEASE (thread_table_lock)
    return // shutdown this processor
```

exit to thread layer

This is where control flow returns to when a thread ENTERs




**procedure**

EXIT\_PROCESSOR\_LAYER(*procid*, *tid*)

*processor\_table*[*procid*].topstack <- SP

SP <- *thread\_table*[*tid*].topstack


**return**




**Save stack pointer  
to processor  
layer thread's  
topstack**

# Switching threads

User layer thread 0



Processor layer thread




Save SP into  
processor\_table[procid].topstack

SP <-  
thread\_table[6].  
topstack

User layer thread 6




Processor layer thread



User layer thread 6

Control flow now  
right after YIELD  
of thread 6

YIELD releases lock and  
returns



Return puts us right after  
the ENTER\_PROC\_LAYER  
call of thread 6

# Initializing processor thread

**procedure** RUN\_PROCESSORS()


**for each** processor **do**

(allocate stack and set up processor thread)

shutdown <- FALSE

**SCHEDULER ()**

**Will only return from  
scheduler when  
shutdown=TRUE**



(deallocate processor thread stack)

(halt processor)

# Processor thread - SCHEDULER

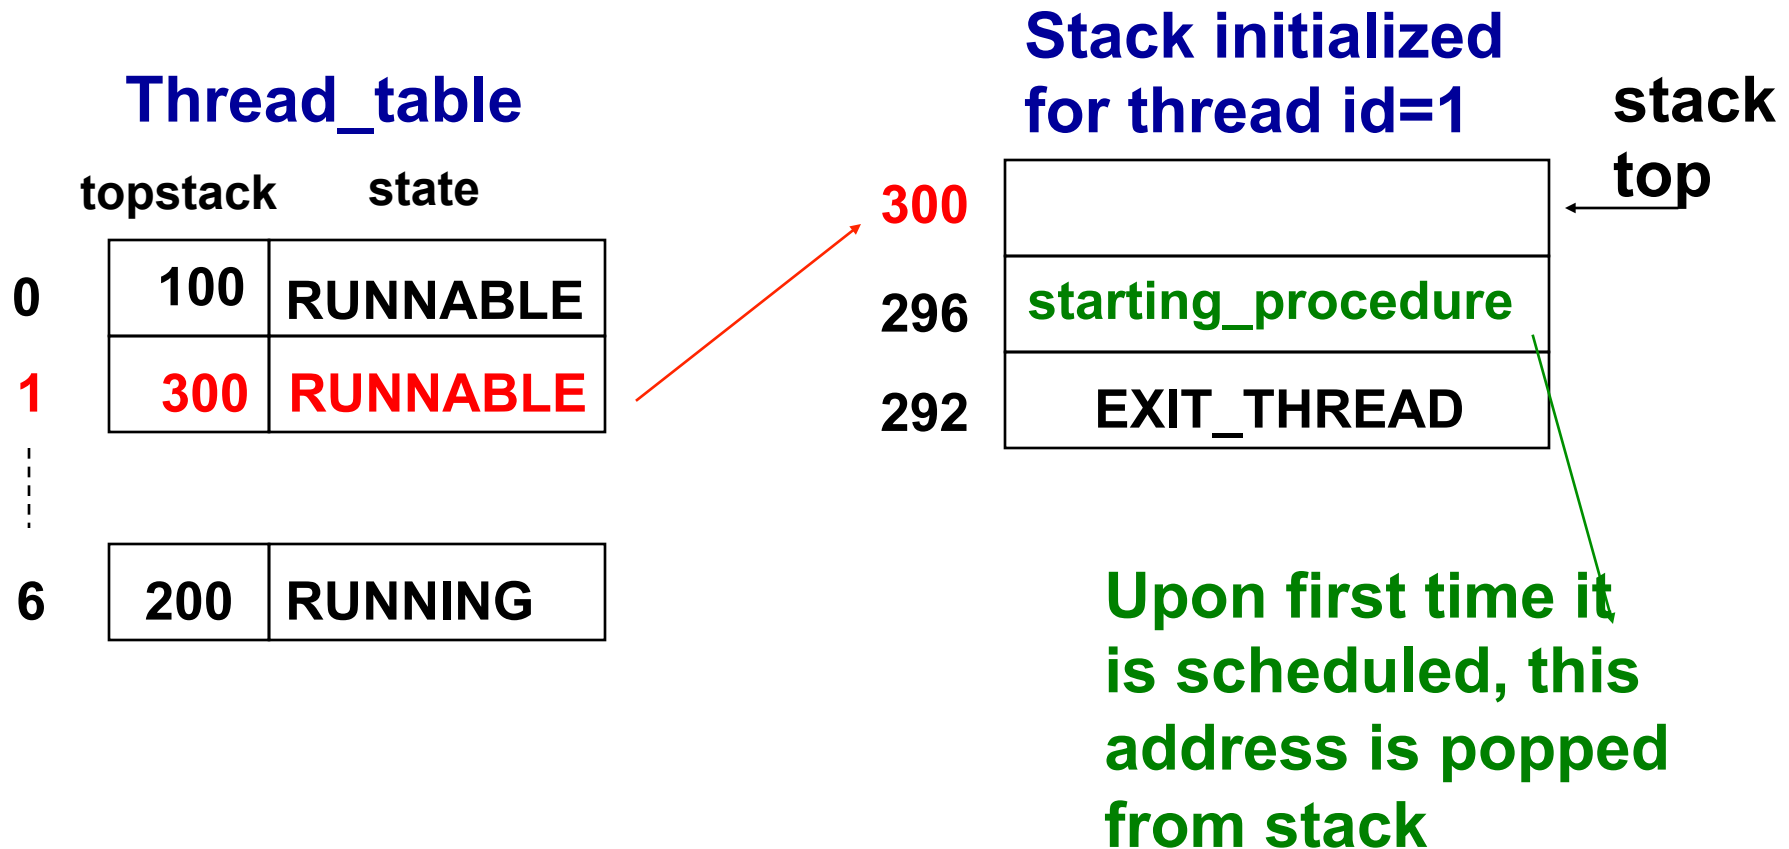
```
procedure SCHEDULER()  
  while shutdown = FALSE do  
    ACQUIRE (thread_table_lock)  
    for i from 0 until 7 do  
      if thread_table[i].state = RUNNABLE then  
        thread_table[i].state <- RUNNING  
        processor_table[CPUID].thread_id <- i  
        EXIT_PROCESSOR_LAYER(CPUID, i)  
        if thread_table[i].kill_or_continue = KILL  
          thread_table[i].state <- FREE  
          DEALLOCATE(thread_table[i].stack)  
          thread_table[i].kill_or_continue = CONT  
        RELEASE (thread_table_lock)  
      return // return to RUN_PROCESSORS  
    // and shutdown this processor
```

# Allocating threads

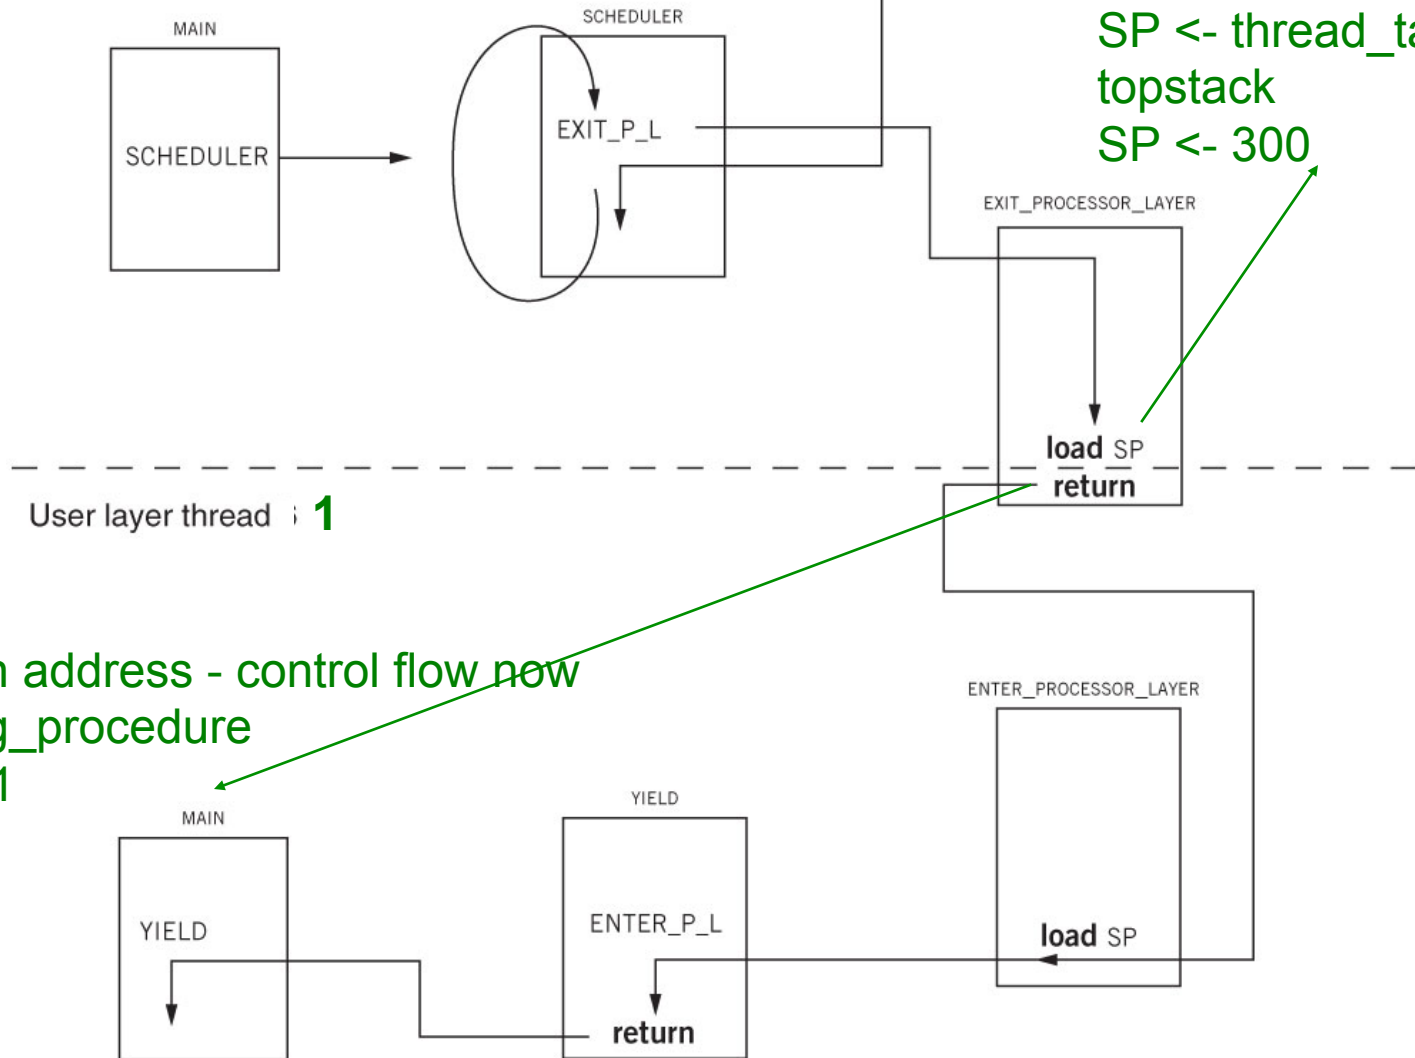
- Allocate space in memory for new stack
- Initialize this stack with return addresses of `EXIT_THREAD` and `starting_procedure`
- Find an entry in thread table that is `FREE` and initialize for the new thread by storing the top of the new stack
- Set state of newly created thread as `RUNNABLE`

# Example

- `id <- ALLOCATE_THREAD`  
(starting\_procedure)



# Example



# Exiting a thread

- Initialized thread stack
  - Constructed to appear as if EXIT\_THREAD called starting\_procedure
  - When thread finishes its execution and returns from starting\_procedure, control flow goes to EXIT\_THREAD

**procedure** EXIT\_THREAD()

ACQUIRE (thread\_table\_lock)

thread\_table[tid].kill\_or\_continue <- KILL

ENTER\_PROCESSOR\_LAYER(tid, CPUID)

# Processor thread - SCHEDULER

```
procedure SCHEDULER()  
  while shutdown = FALSE do  
    ACQUIRE (thread_table_lock)  
    for i from 0 until 7 do  
      if thread_table[i].state = RUNNABLE then  
        thread_table[i].state <- RUNNING  
        processor_table[CPUID].thread_id <- i  
        EXIT_PROCESSOR_LAYER(CPUID, i)  
        if thread_table[i].kill_or_continue = KILL  
          thread_table[i].state <- FREE  
          DEALLOCATE(thread_table[i].stack)  
          thread_table[i].kill_or_continue = CONT  
    RELEASE (thread_table_lock)  
  return // return to RUN_PROCESSORS  
          // and shutdown this processor
```



# Shortcomings

- YIELD – expects threads to be cooperative in sharing the processor
  - Non-preemptive scheduling
- Issues
  - Length of time a thread holds processor is up to the thread itself to decide
  - Errors can propagate from one module to another
    - E.g. infinite loop, fate sharing

# Enforcing Modularity

- In order to enforce modularity, provide all threads with access to the processor
  - Pre-emptive scheduling
- Thread manager invoked not only when a thread YIELDS, but also re-scheduled on a periodic basis
  - Forces a thread to yield
- Core requirements:
  - A timer device that generates interrupts
  - Separation of privilege: Kernel vs. user

# Pre-emptive scheduling

- Time interval – raise interrupt
- Interrupt handler
  - Runs in processor layer, in privileged kernel mode
  - Invokes exception handler in thread layer, which forces thread to YIELD
- Need to deal with concurrency
  - Interrupt handler thread and processor layer thread

# Pre-emptive scheduling

- Example:
  - Thread 0 calls YIELD and acquires thread\_table\_lock
  - Hardware timer interrupt arrives and interrupt handler also calls YIELD
    - Another attempt to acquire thread\_table\_lock – deadlock
- An often-used solution relies on processor support to disable/re-enable interrupts
  - Disable interrupts before acquiring lock in YIELD
  - Re-enable after releasing the lock
  - Interrupt handling is disabled in between

# Modularity in memory


- Pre-emptive scheduling needs to be coupled with an approach to enforce modularity in memory
- Make thread manager aware of virtual address spaces of threads
  - When switching to another thread, also switch the virtual address space
    - Thread state also holds PMAR
    - ENTER\_PROCESSOR\_LAYER: save PMAR into thread state (thread\_table[] extended to store PMAR)
    - EXIT\_PROCESSOR\_LAYER: load PMAR from thread state

# Address space switch

- Once the PMAR register is switched, all further memory references will be translated with the new PMAR
- Additional complexity to deal with
  - Map thread manager text/data address space in a subset of every thread's address space (with `KERNEL_ONLY`)
  - If hardware supports loading of PMAR, SP, PC in a single atomic action, may switch the entire context in one shot
    - Add processor-layer PMAR to `processor_state`

# Address space mappings

- Reduces availability of virtual addresses to user-level threads
- Avoids the need to invalidate (flush) hardware-based TLB when entering the kernel
  - Flush still needed if another thread is scheduled



# Reading

- Sections 6.1 and 6.3