

**EEL-4736/5737**  
**Principles of Computer System  
Design**

Lecture Slides 22

Textbook Chapter 9


Atomicity Logs – Before-or-After and Locks

# Introduction

- Before-or-after atomicity
  - Mark-point discipline
    - Commits to journal
    - Values not visible to other concurrent actions until committed
      - READ\_CURRENT\_VALUE gets data from journal when it is committed
  - Logs
    - Installs to cell storage
    - Cell storage updates are visible immediately to other threads
      - READ\_CURRENT\_VALUE gets data from cell storage when installed

# Introduction

- Multiple threads installing to cell storage
  - A commit to the log is a single atomic operation – append
  - Installs can require multiple writes
  - Ok if single action; but concurrent actions, possible to see intermediate requests



# System-wide lock

- Simple serialization
  - One lock used to serialize all actions
- Simple, correct – but conservative

Begin transaction

ACQUIRE(system-lock)

...

RELEASE(system-lock)

End transaction

# Simple locking discipline

- Similar in spirit to mark-point discipline
- Rule #1:
  - Transaction must acquire a lock for *every* object it intends to *read or write*, *before* actually reading or writing
    - A “lock set”
- Rule #2:
  - It may release a lock only *after* it has installed its *last update* and committed (or completely restored its data and aborted)

# Simple locking discipline

- Intuition:
  - Acquire locks of *all* objects before proceeding to *any* reads or writes
    - At this point, guaranteed that no other action will update any of these objects
  - Release *first* lock only after all updates are installed
    - Any other action that uses any of the objects in the “lock set” is blocked because it has not been able to acquire at least one lock
    - Hence, no other action using anything on lock set will “see” intermediate installs on cell storage – only after full set is installed

# Simple locking discipline

- A “lock manager” can expose lock set and enforce discipline without relying on application to write locks
  - Name all objects that will be used at the beginning of a transaction
  - Lock manager ensures `begin_transaction` starts by acquiring all locks in set (thus blocking until all acquired)
  - Use logging of `END` to release locks

# Simple locking discipline

- Improved opportunities for concurrency compared to system-wide locking
  - E.g. consider transfer transactions working on completely disjoint sets of accounts
- However, there are still opportunities for concurrency that may be lost
  - E.g. consider two actions A1, A2 that read from the same object X, while no concurrent action intends to write to X
  - A1 or A2 will block waiting to acquire lock for X, but no real need for lock

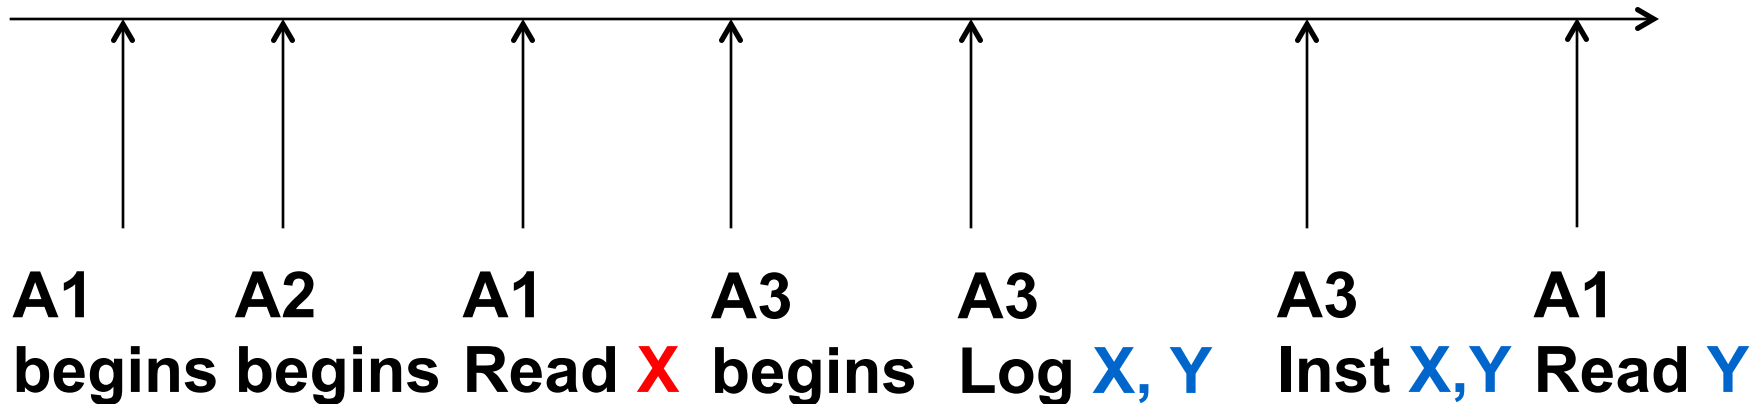


# Why locking of reads?

**A1: read X, Y; write Z**

**A2: read X, Y; write W**

**A3: write X, Y**




**A1 reads “red” version of X, “blue” version of Y**  
**Intermediate steps of A3 exposed**

# Why locking of reads?

**A1: read X, Y; write Z**

**A2: read X, Y; write W**

**A3: write X, Y**



**Lock of X, Y blocks A2, A3 from reading/writing  
A1 reads “red” version of Y**

# Simple locking discipline

- Need to know in advance all data that may be read or written so they can be locked
  - If not known in advance, it is a problem
  - What to read may depend on value that itself is read from the shared storage
  - E.g. consider traversing a pointer-based linked list – next pointer to read not known until current is read

# Two-phase locking discipline

- Avoids requirement that transaction knows in advance which locks it needs to acquire
- Allows transaction to acquire locks as it progresses – before it uses an object
- May read/write an object as soon as it acquires its lock
- Primary constraint:
  - May only release *any locks* until it passes its *lock point*
    - First instance when it has acquired *all locks*

# Two-phase locking discipline

- Second constraint:
  - May only release a lock (after lock point) for an object that it only reads *if* it will never need to read the object again (even if it needs to abort)
- First phase
  - Number of locks monotonically increase as they are acquired until reaching lock point
- Second phase
  - Number of locks monotonically decrease as they are released

# Two-phase locking discipline

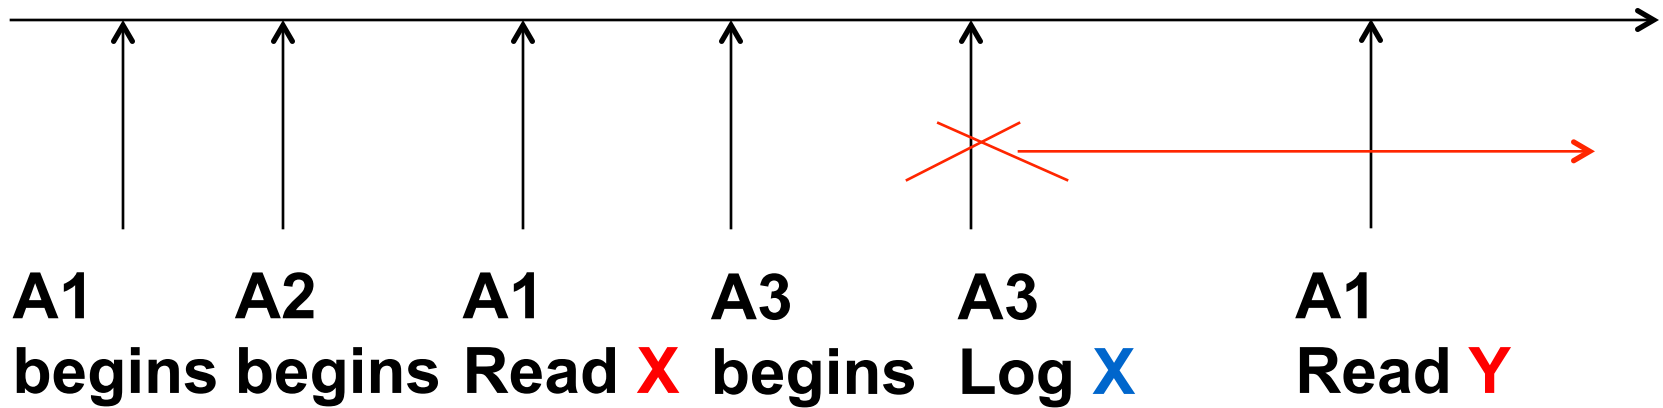
- Lock manager can simplify use and enforce correctness
  - Interpose reads/writes
    - First use of any object: precede by a lock acquire
  - Interpose logging the end of transaction
    - Releases all locks

# Two-phase locking

**A1: read X, Y; write Z**

**A2: read X, Y; write W**

**A3: write X then Y**



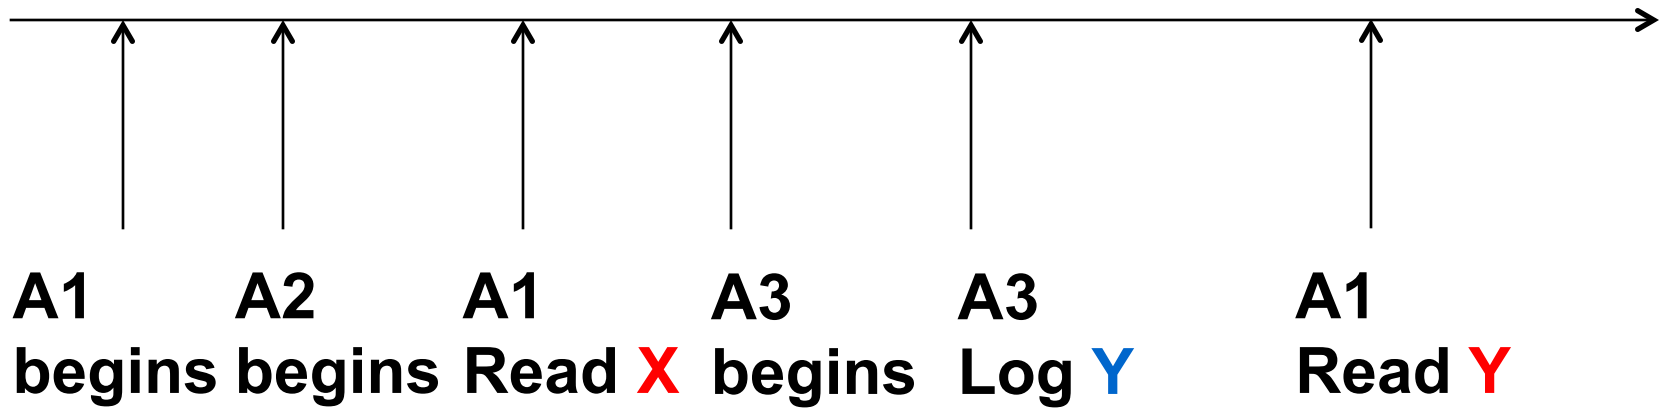
If A1 reads “red” version of X, it has locked X;  
Thus, A3 cannot write “blue” version of X until  
after A1 releases lock; A1 only releases lock after  
lock point, hence it also has a lock for Y – so it must  
Also read “red” Y

# Two-phase locking

**A1: read X, Y; write Z**

**A2: read X, Y; write W**

**A3: write Y then X**



If A1 reads “red” version of X, it has locked X;  
A3 may log “blue” version of Y, but it cannot  
install Y and commit until it has reached its  
“lock point”: it’d need X’s lock to reach lock point  
(and X’s lock is still held by A1)

**Deadlocks can happen; detect, abort victim thread**