

EEL-4736/5737
**Principles of Computer System
Design**

Lecture Slides 4

Textbook Chapter 2

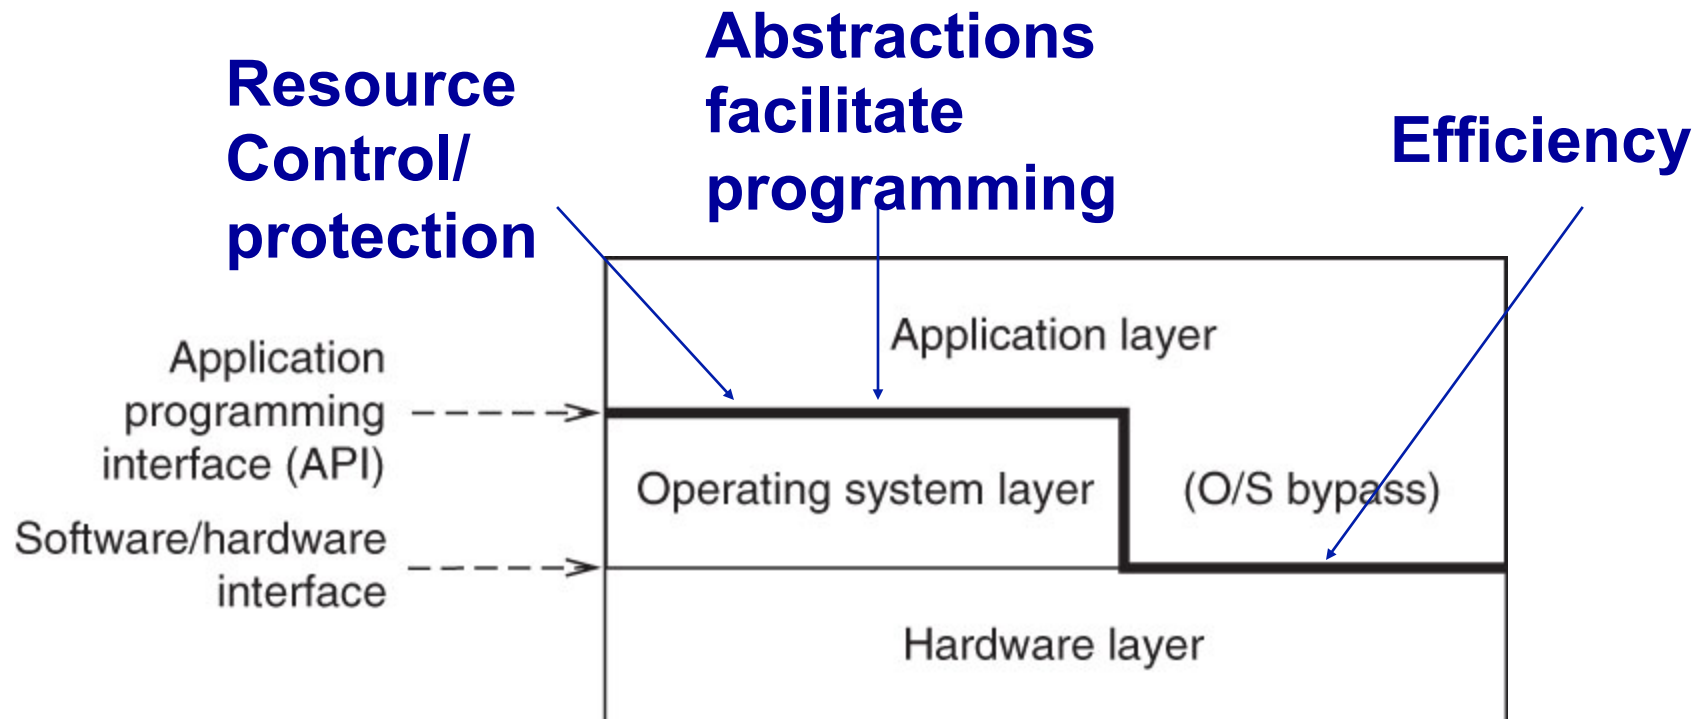
Organizing Computer Systems with Names
and Layers

Introduction

- Example of a layered system: O/S
- What constitutes an operating system?
- Set of software (kernel, applications, libraries) that facilitates programmers and users to do their job
 - Provide convenient abstractions to facilitate programming
 - Handles interface with, allocation and scheduling of hardware devices
 - We'll study O/Ss in more depth from a system design perspective in chapter 5

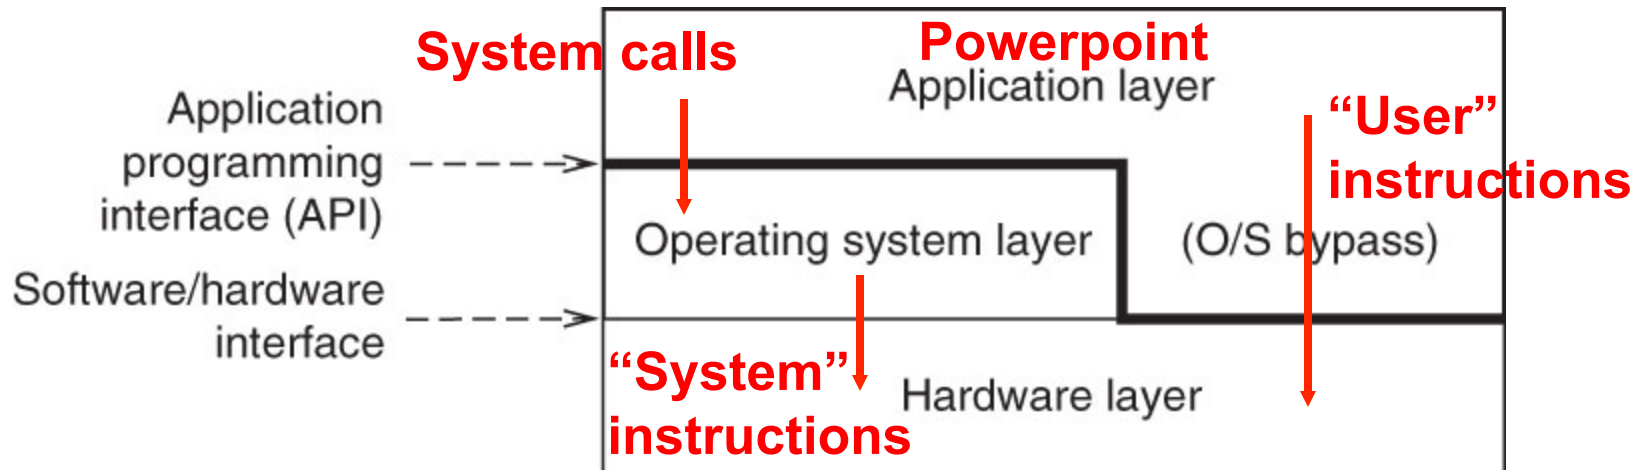
Operating system and layers

- Coarse-grain layers in typical computer
 - Several layers within each

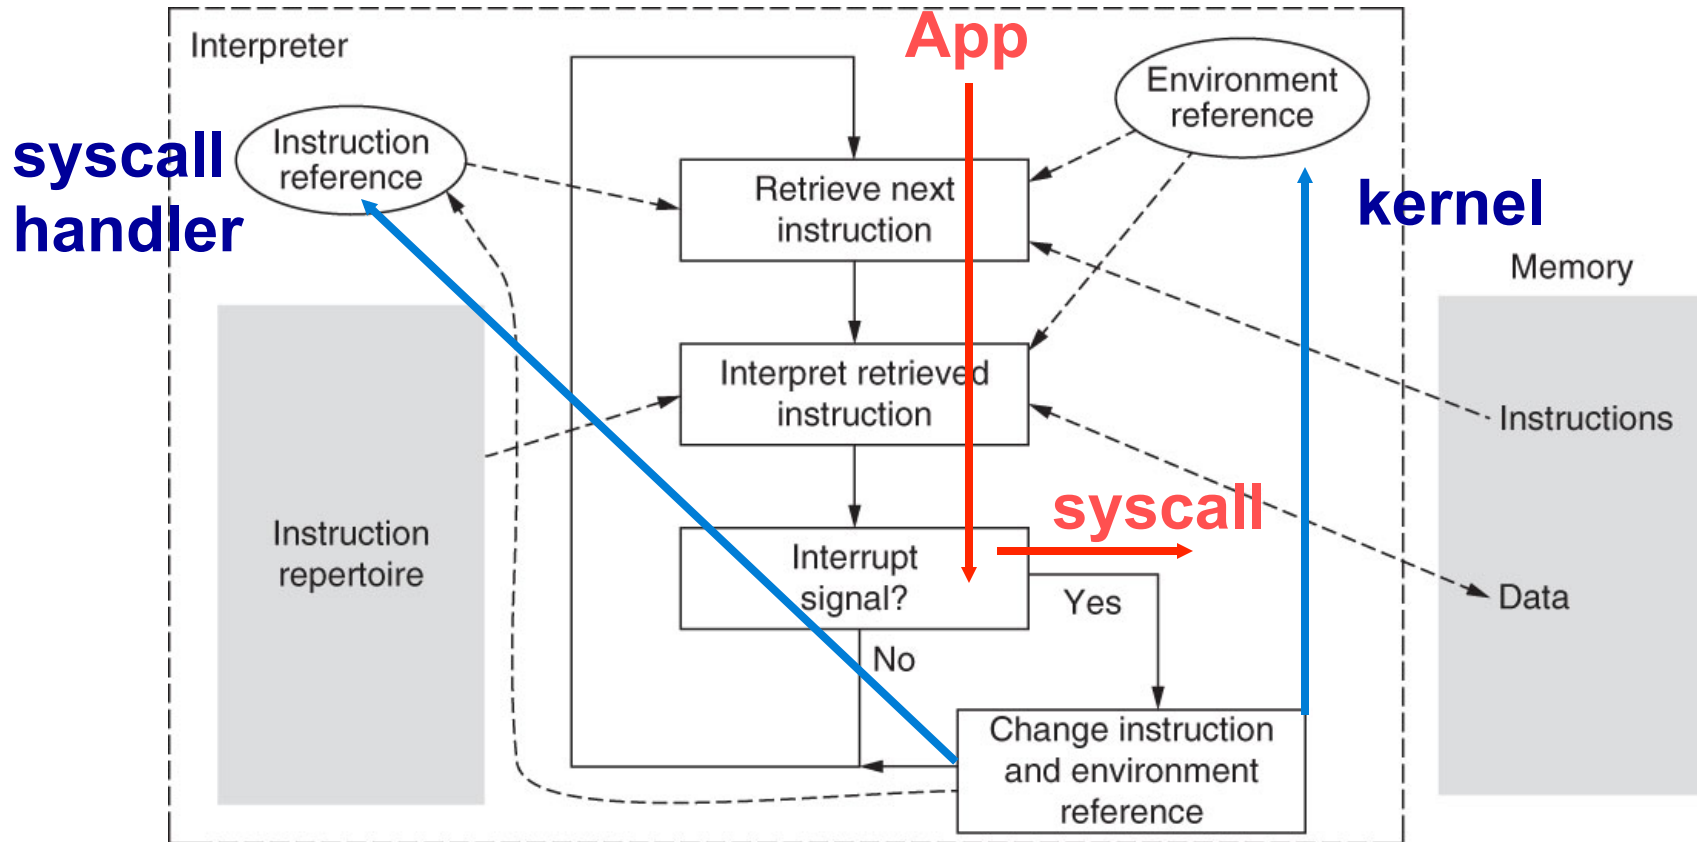


Operating system and layers

- Example



System call – synchronous intr



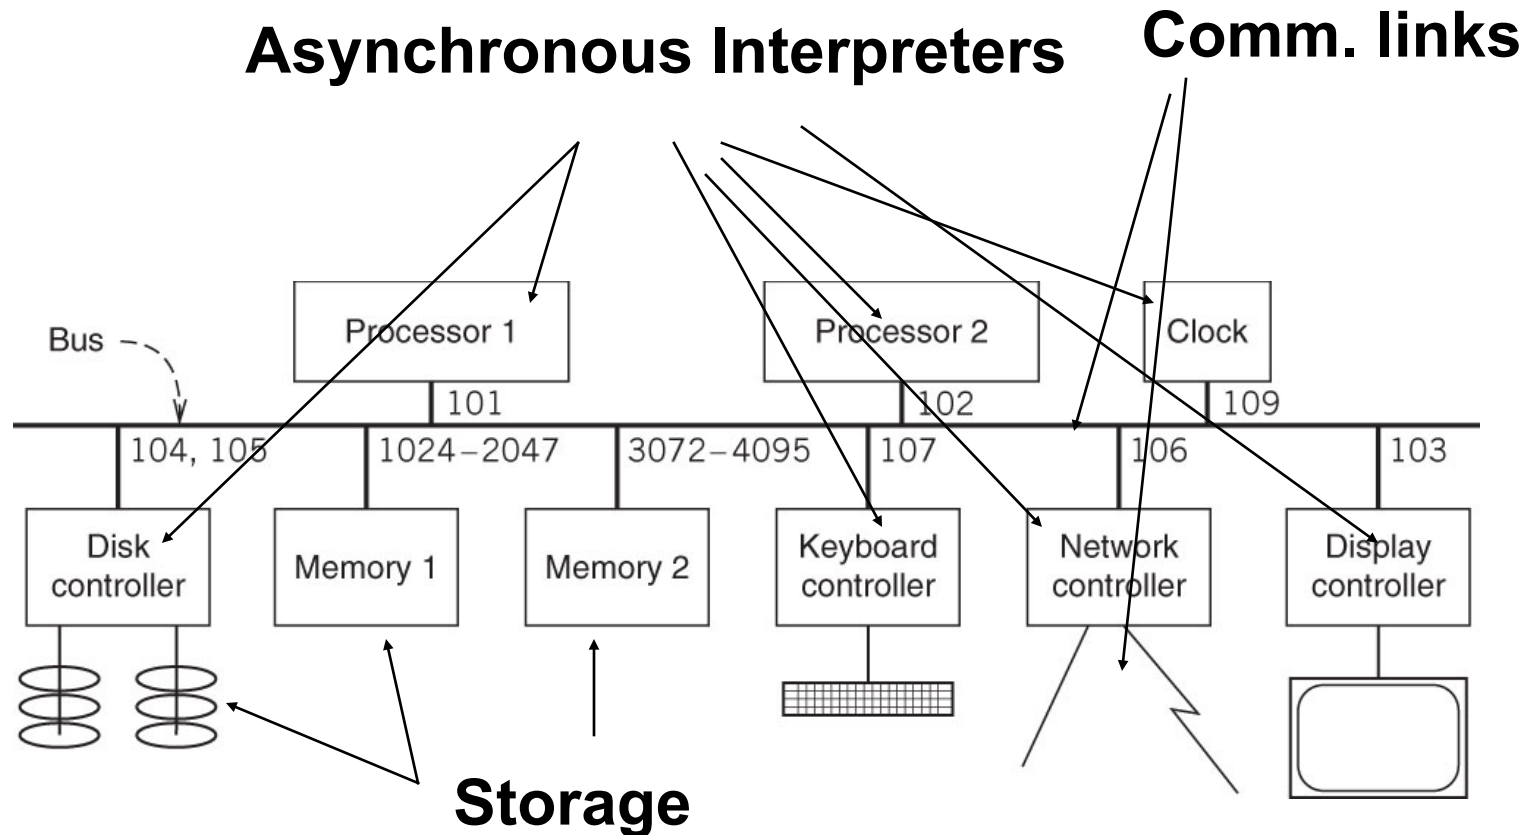
```

1  procedure INTERPRET()
2    do forever
3      instruction ← READ (instruction_reference)
4      perform instruction in the context of environment_reference
5      if interrupt_signal = TRUE then
6        instruction_reference ← entry point of INTERRUPT_HANDLER
7        environment_reference ← environment ref of INTERRUPT_HANDLER

```

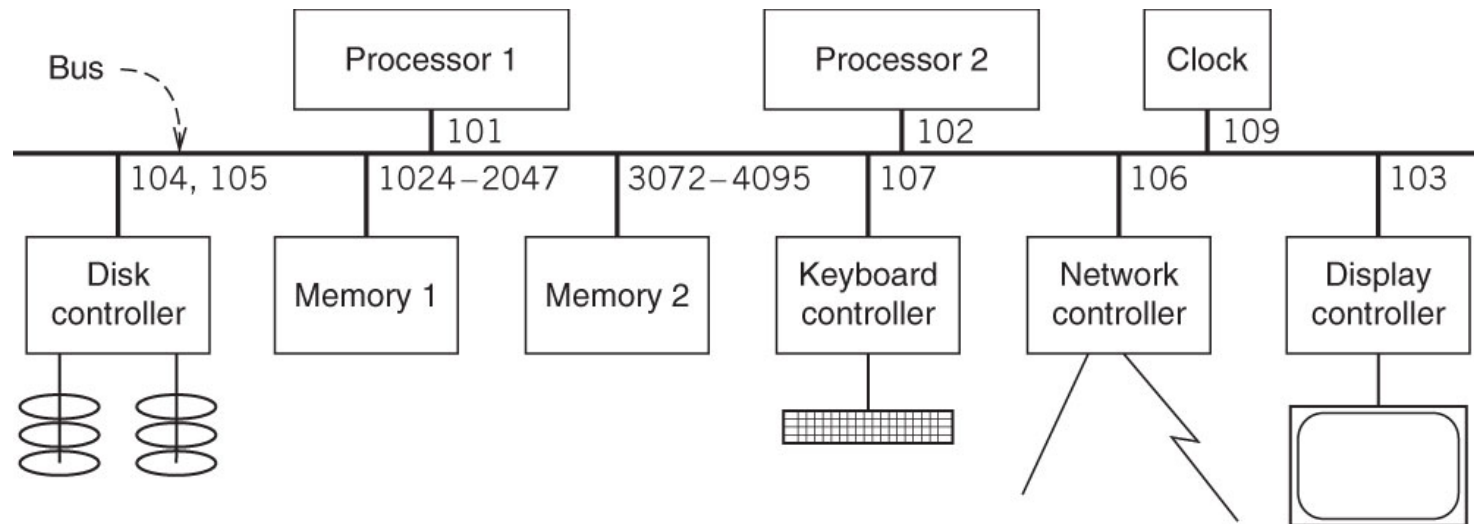
A hardware layer - Bus

- Implement low-level versions of the three fundamental abstractions




Examples

- Each device has an interface to the bus with a unique name (address)
- Broadcast communication link
 - SEND reaches all devices; RECEIVE sees all messages; no need to route



Examples

- Name space partitioning
 - N-bits: integer 0 to $2^n - 1$
 - In the example, memories use a range of address space

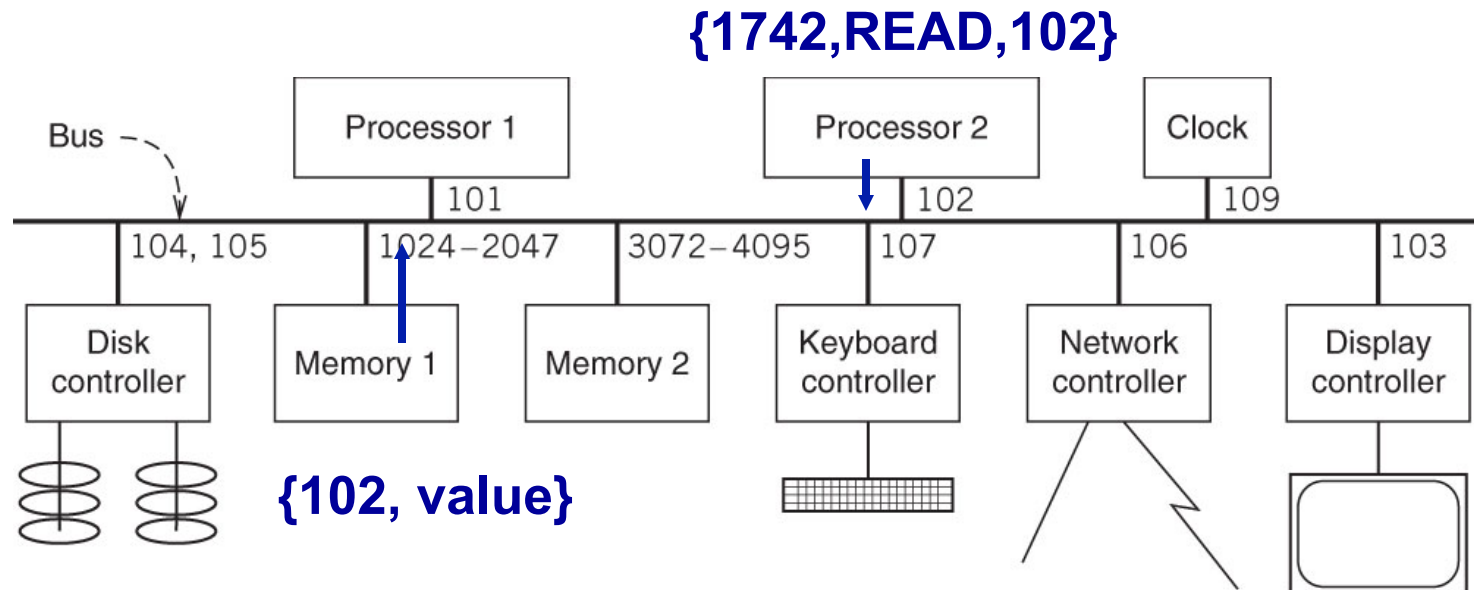


Pseudo-code convention

- $\text{total} \leftarrow a + b$
 - Assignment
- $a(11..15)$
 - Bits 11..15 from string a
- $x \Rightarrow y: \{M\}$
 - Message with contents M sent from x to y
 - $\{a,b,c\}$ – message contains named fields marshaled in some way the recipient understands

Examples

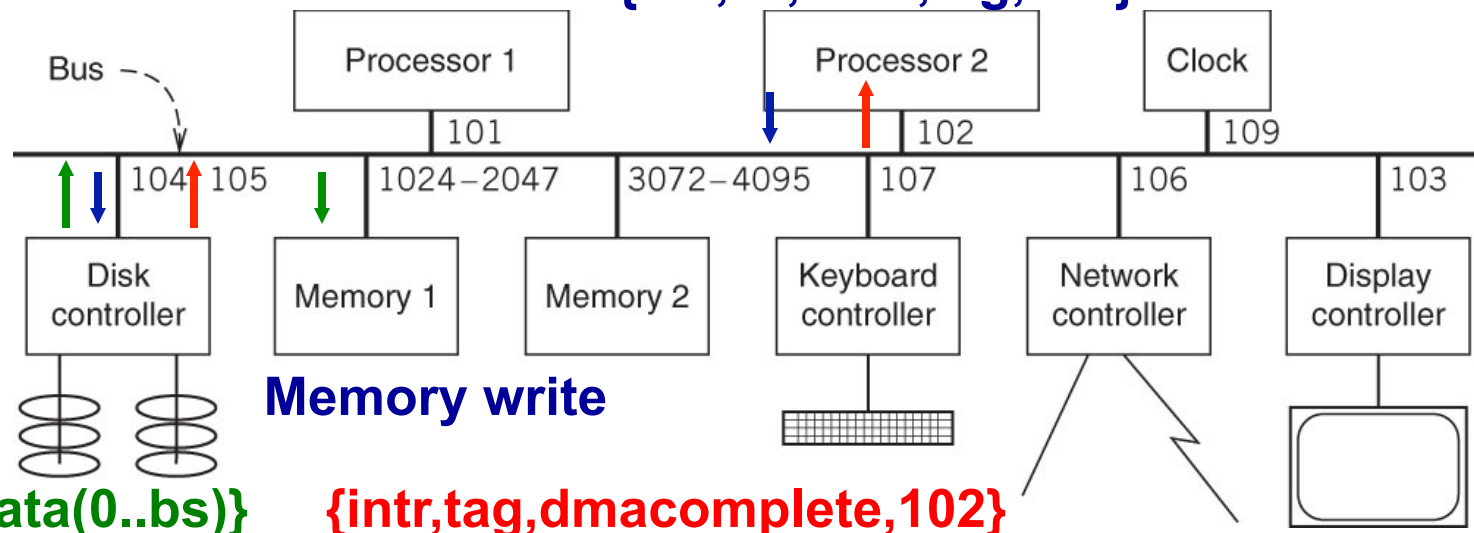
- Processor 2 interprets LOAD 1742, R1
 - P2 => all bus modules: {1742, READ, 102}
- Memory module 1
 - value <- READ(1742) (internally)
 - M1 => all bus modules: {102, value}



Examples

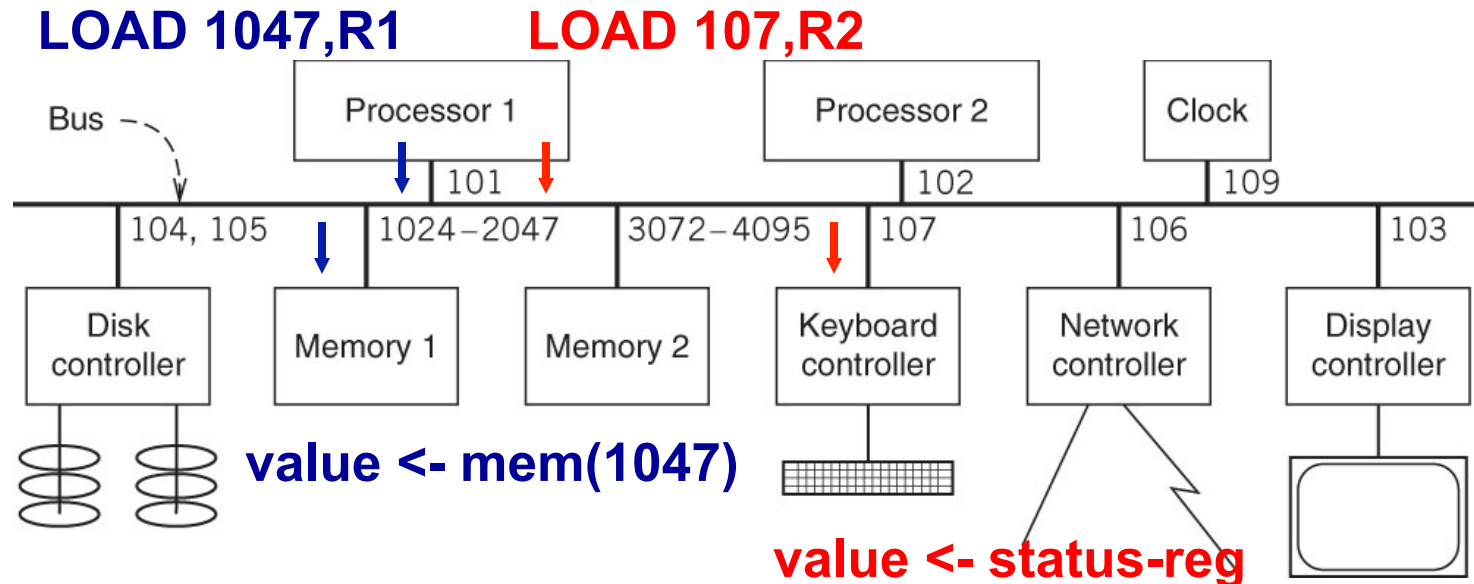
- Direct memory access – decoupled transfers
 - E.g. disk read: P2 sends a message to device 104:
 - Which block to read, and size
 - Reference to starting memory address
 - Disk controller reads block from disk
 - Sends *data* message directly to memory module
 - Sends *interrupt* message to P2

{blk,sz,start,tag,102}



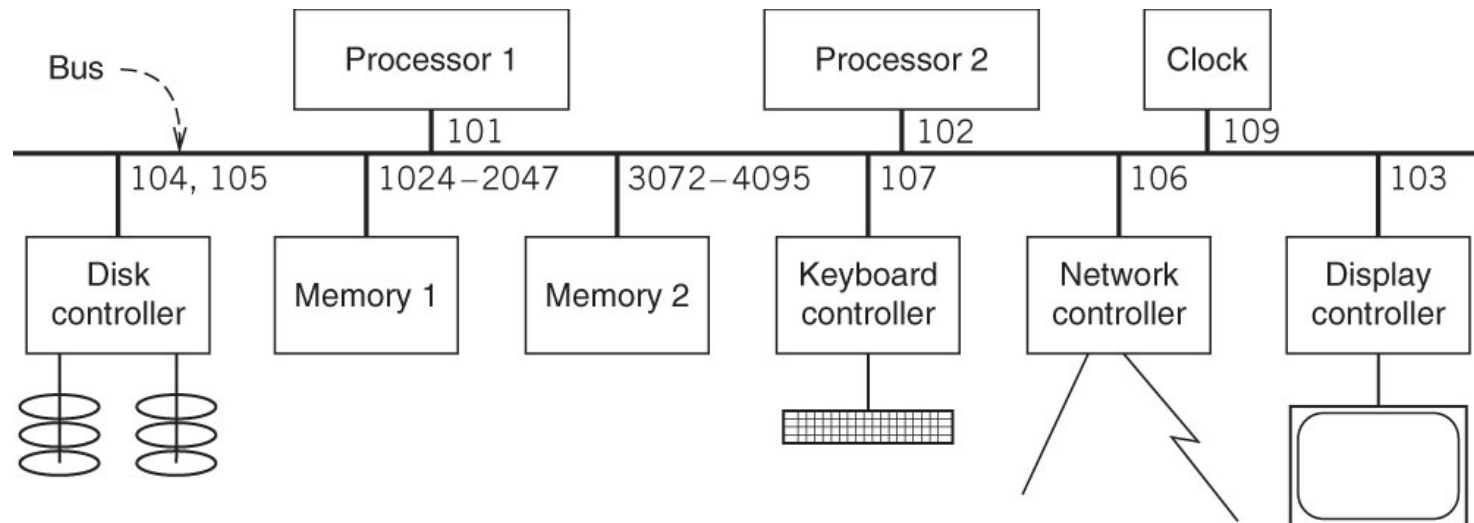
Memory-mapped I/O

- Memory-mapped I/O
 - Processor-issued loads/stores – messages on the bus
 - Subset of addresses bound to devices
 - E.g. DMA controller exposes addresses for device control registers



Addressing devices

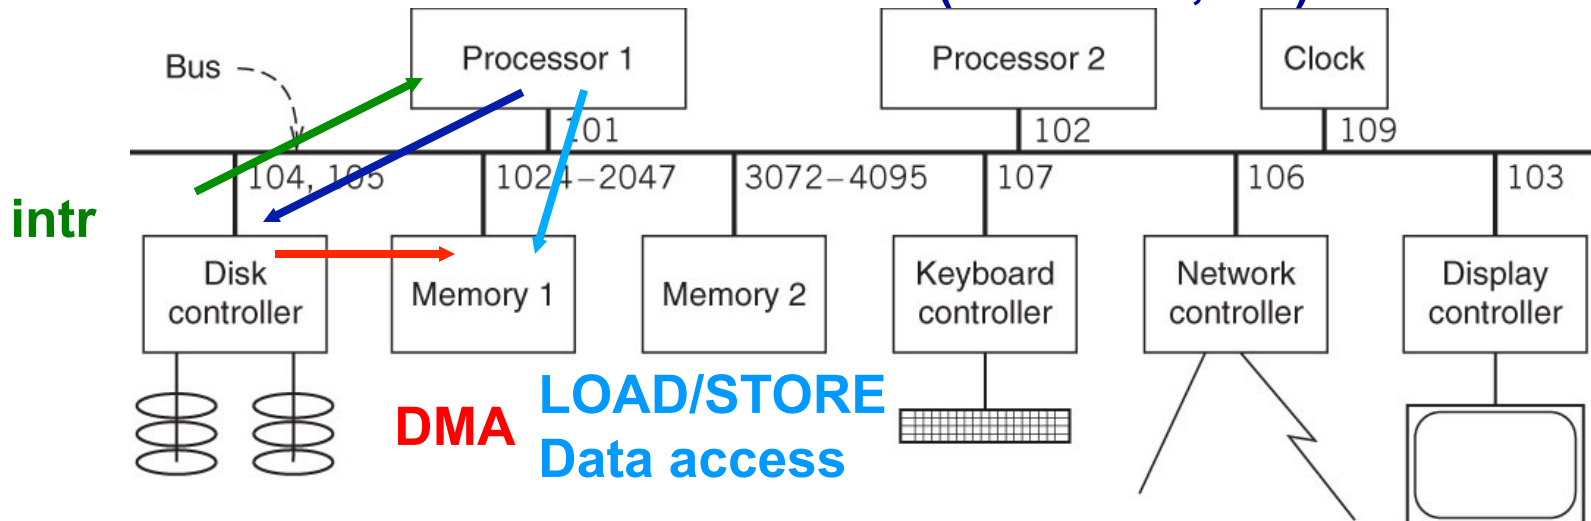
- Disk – storage
 - Should disk storage be exposed as memory-mapped I/O?
- Remote main memory



Virtual memory

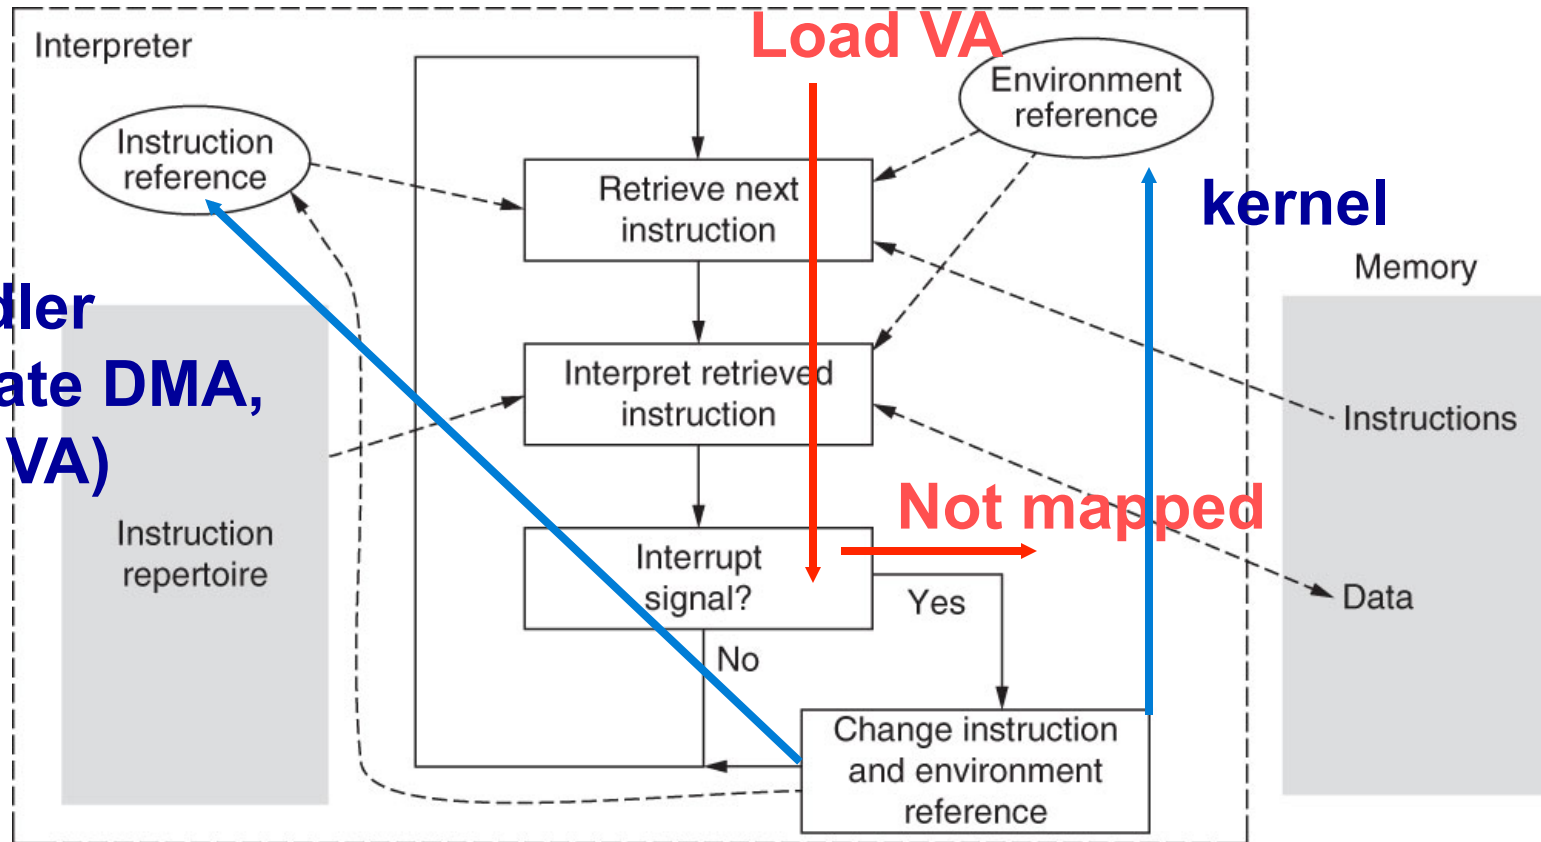
- A widely-used approach - virtual memory
 - Memory-mapped I/O used to program devices – disk controllers, DMA
 - Transfers at a large granularity (blocks, pages vs. words – efficient)

LOAD/STORE: initiate transfer (addr: 104,105)



Virtual memory and interpreter

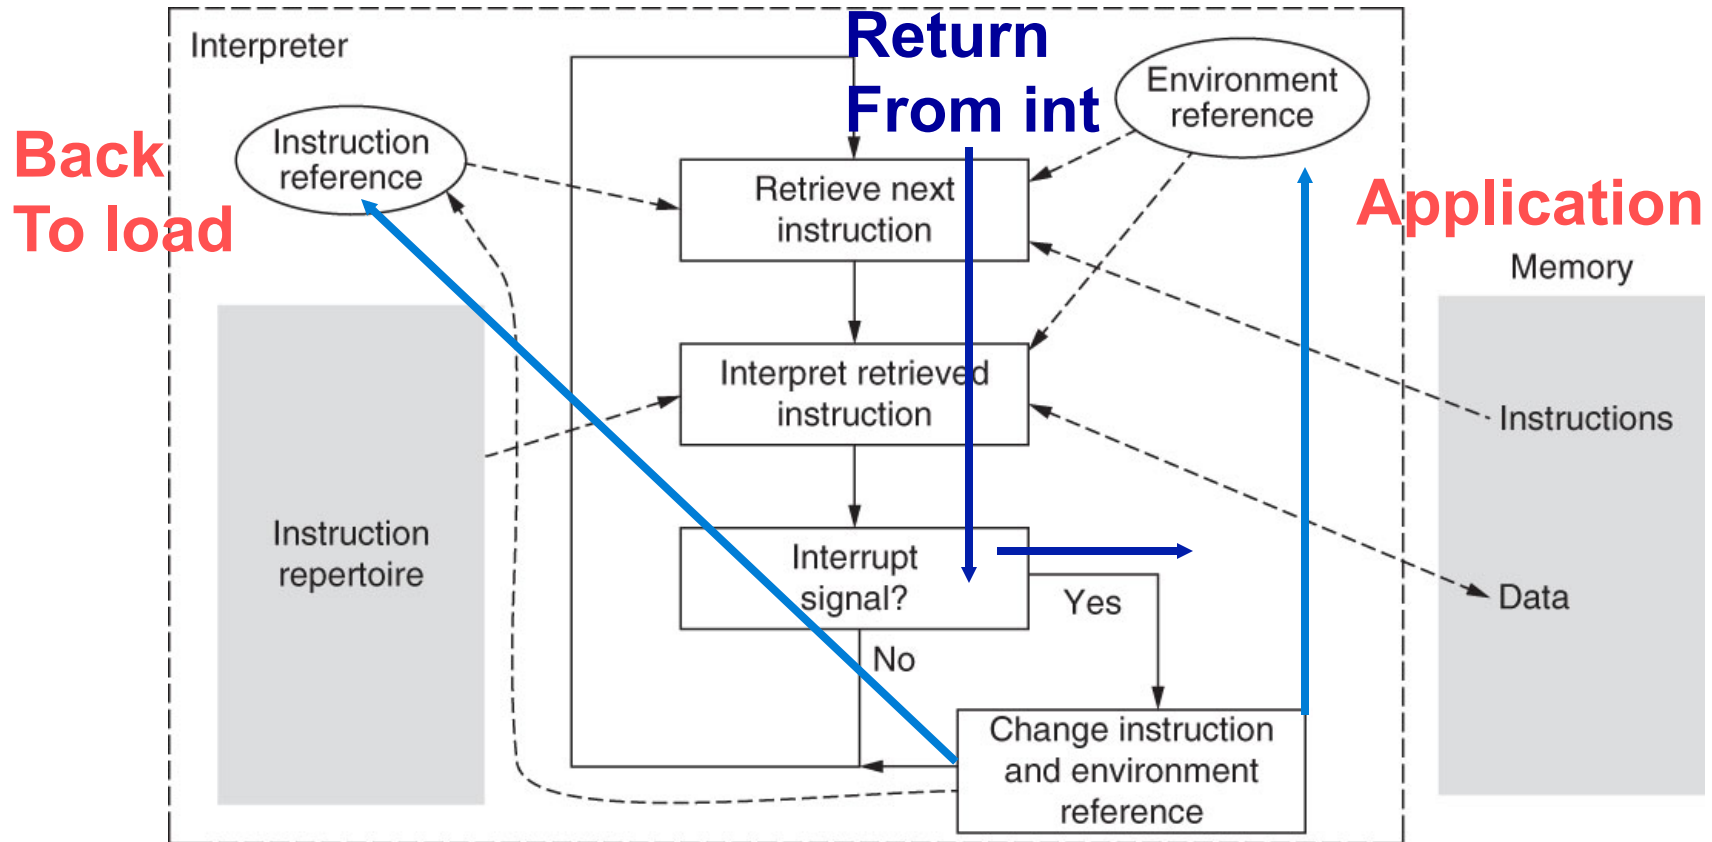
**Virt.
Mem
Handler
(Initiate DMA,
map VA)**



```

1  procedure INTERPRET()
2    do forever
3      instruction ← READ (instruction_reference)
4      perform instruction in the context of environment_reference
5      if interrupt_signal = TRUE then
6        instruction_reference ← entry point of INTERRUPT_HANDLER
7        environment_reference ← environment ref of INTERRUPT_HANDLER
    
```

Virtual memory and interpreter

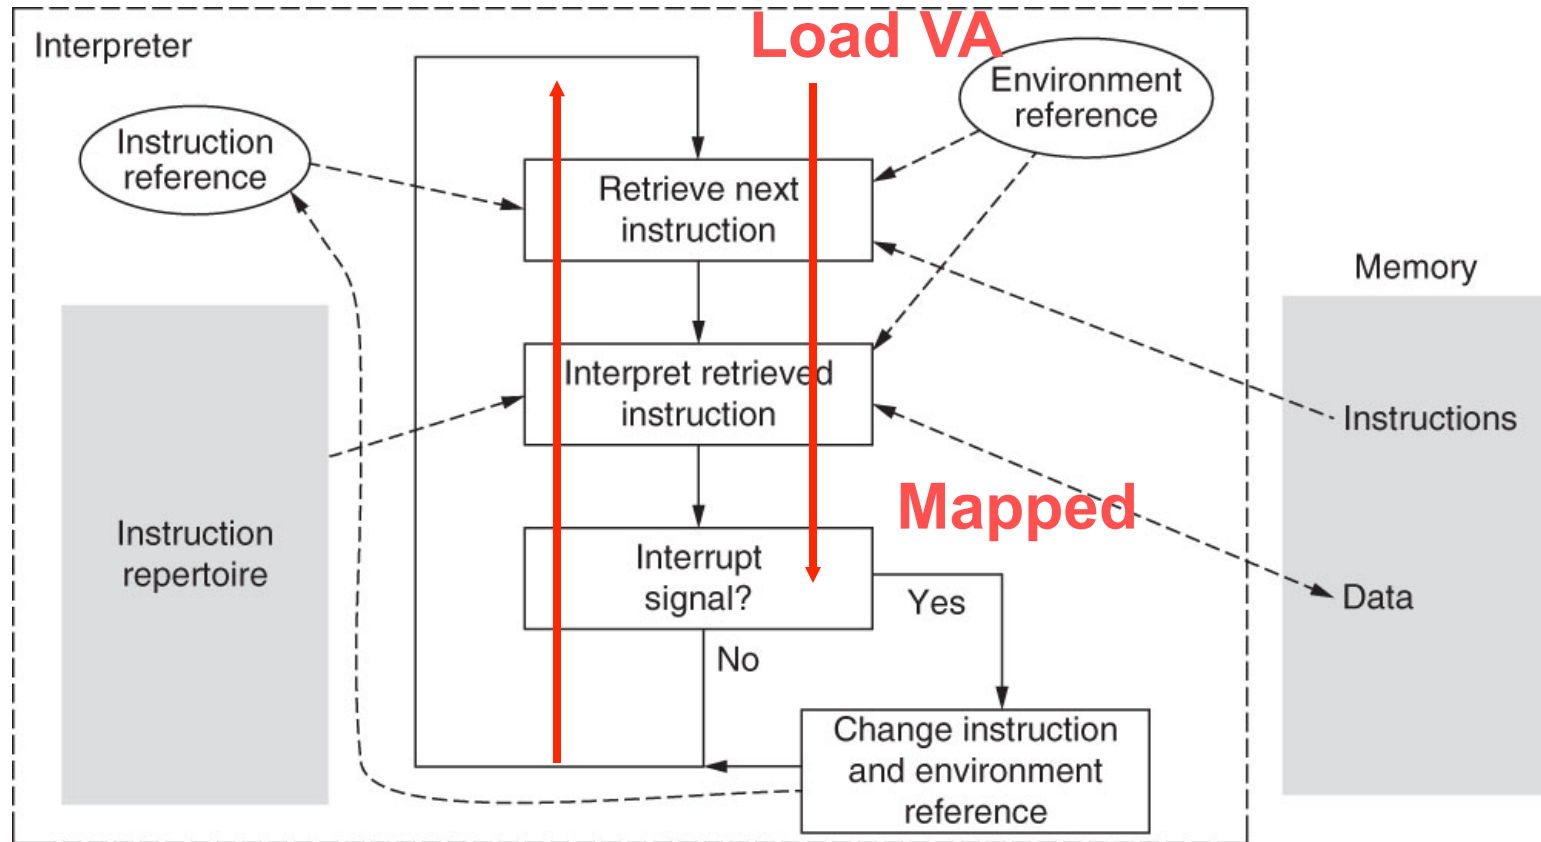


```

1  procedure INTERPRET()
2    do forever
3      instruction ← READ (instruction_reference)
4      perform instruction in the context of environment_reference
5      if interrupt_signal = TRUE then
6        instruction_reference ← entry point of INTERRUPT_HANDLER
7        environment_reference ← environment ref of INTERRUPT_HANDLER

```


Virtual memory and interpreter



```
1  procedure INTERPRET()
2    do forever
3      instruction ← READ (instruction_reference)
4      perform instruction in the context of environment_reference
5      if interrupt_signal = TRUE then
6        instruction_reference ← entry point of INTERRUPT_HANDLER
7        environment_reference ← environment ref of INTERRUPT_HANDLER
```

A software layer: File abstraction

- High-level version of memory abstraction
 - A file holds an array of bytes
 - It is durable
 - How long depends on many factors
 - It has a name
 - Can refer to it later;
 - Can be used to share information
- The abstraction can be used in ways other than storing data with durability
 - E.g. a device as a file

File abstraction – core API

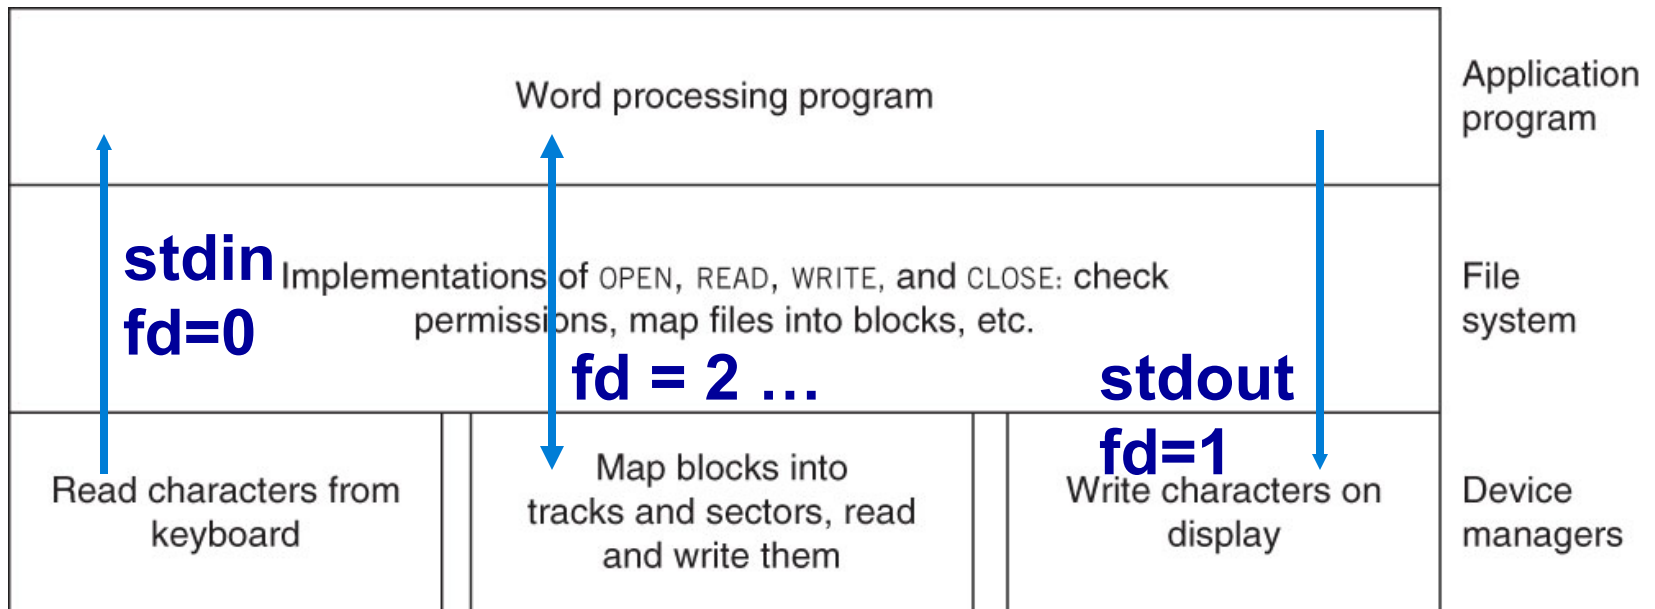
- OPEN
 - Translate file name to name within a local name space to be used by subsequent operations
 - E.g. an integer file descriptor
- READ, WRITE
 - Retrieve/store data from/to an open file
- CLOSE
 - Unbinds file name from descriptor; may also store data on stable storage
- Why OPEN and CLOSE?

File abstraction - APIs

- Need for OPEN/CLOSE
 - Efficiency
 - Replace (long) filename with short integer
 - Saves space – system call stack
 - Saves time – name resolution, permission check
 - Atomicity and concurrency control
 - Can determine which objects share a file at a given point in time
 - E.g. you can block an OPEN if file is already opened for before-after atomicity
 - Can be used to deal with faults
 - Crash before a CLOSE?

File abstraction uses

- E.g. UNIX standard in/standard out



Reading

- Section 2.5