

EEL-4736/5737
**Principles of Computer System
Design**

Lecture Slides 16
Textbook Chapter 7
Network layers

Introduction

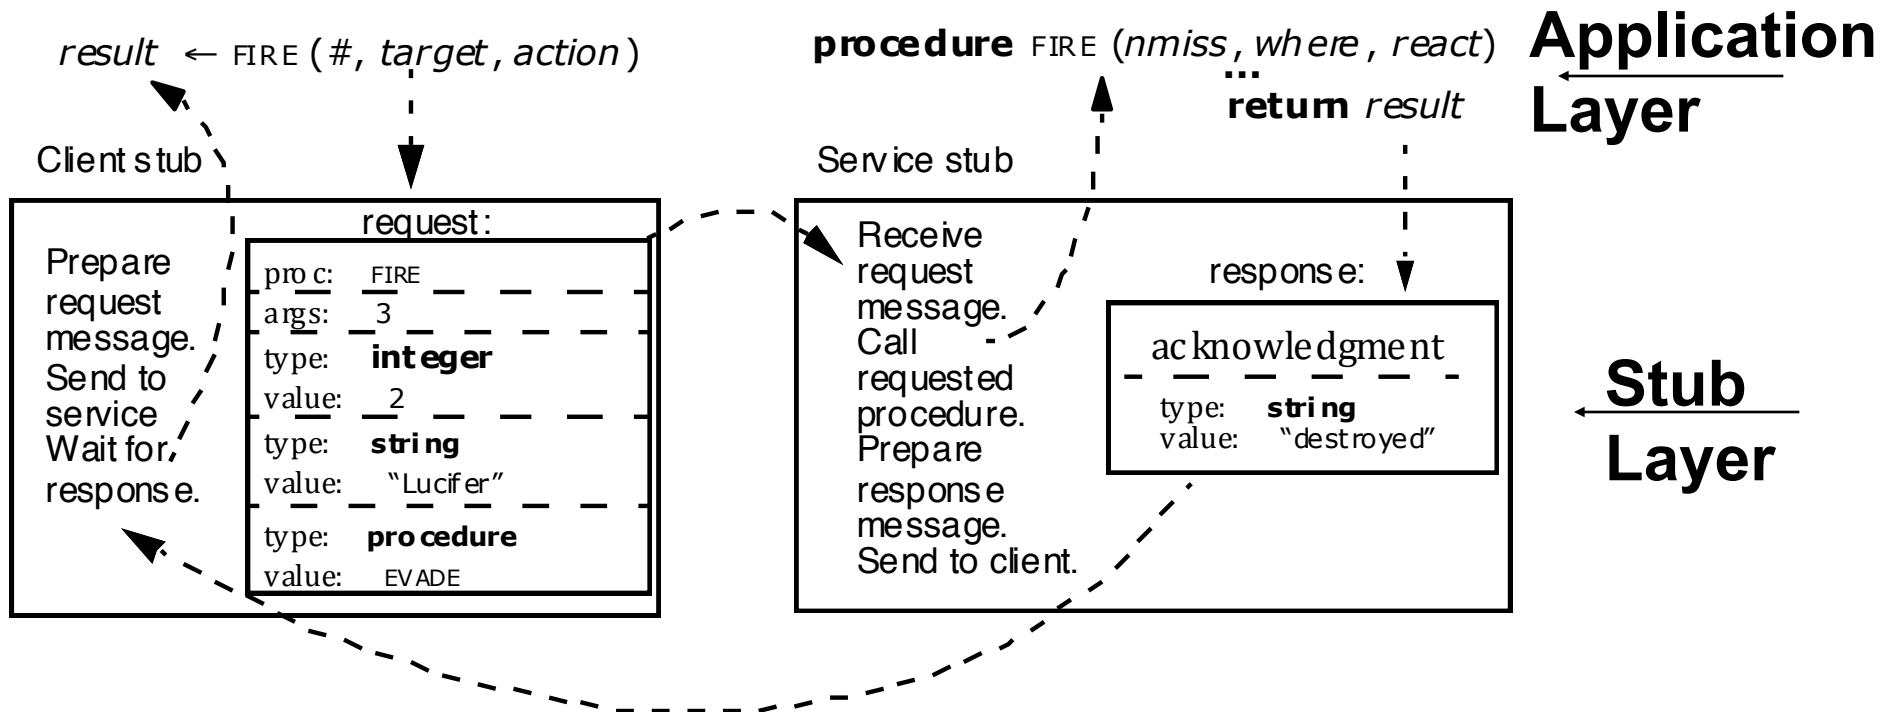
- How to deal with the properties of networks
- 1. Networks encounter a vast range of
 - Data rates
 - Propagation, transmission, queuing, and processing delays.
 - Loads
 - Numbers of users
- 2. Networks traverse hostile environments
 - Noise damages data
 - Links stop working
- 3. Best-effort networks have
 - Variable delays
 - Variable transmission rates
 - Discarded packets
 - Duplicate packets
 - Maximum packet length
 - Reordered delivery

Divide-and-conquer

- Rather than addressing all these problems at once, networks apply the *layering* design principle extensively
- Each layer follows a *protocol*, and responsibility to address different challenges we have described are assigned to appropriate layer(s)

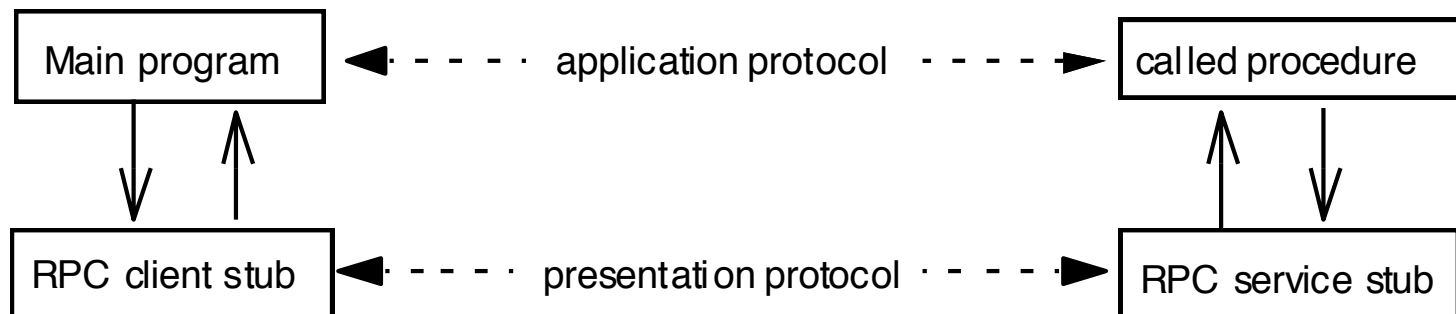
Protocols

- Specify all possible kinds of interactions between entities in the system
 - Examples: RPC, Network File System (NFS)



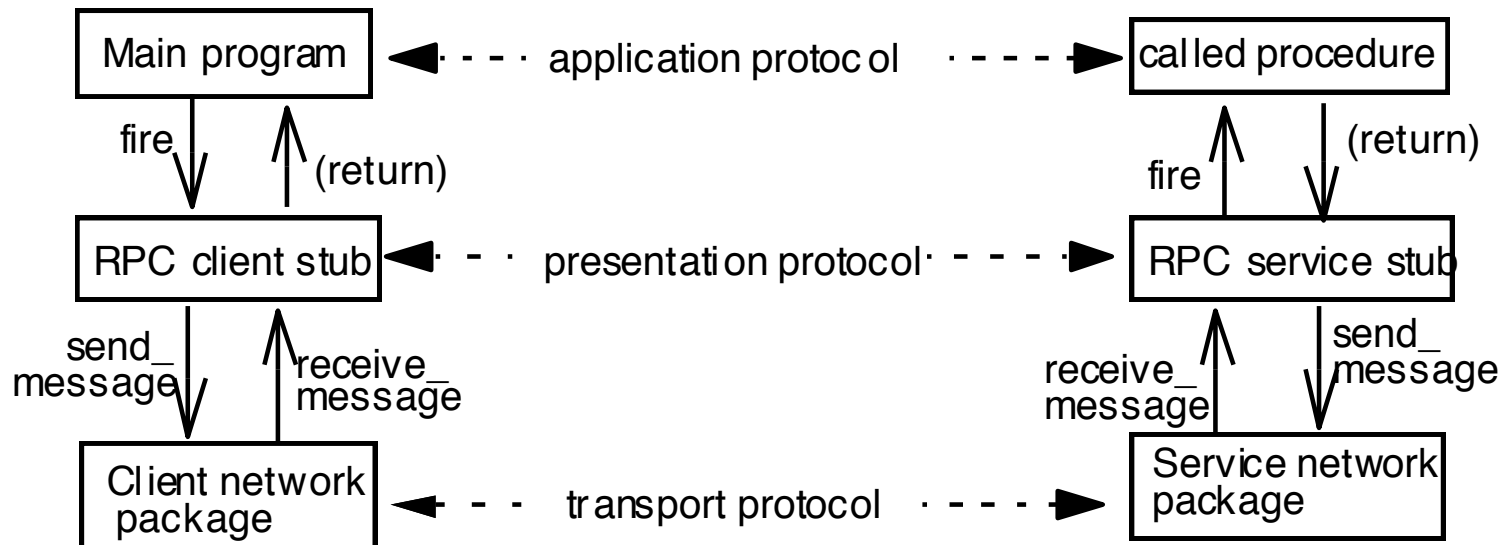
Layering protocols

- Application protocol: not concerned with how RPC arguments are marshaled, timeout/retry
 - It is concerned with what each procedure implements
 - E.g. NFS READs vs REMOVEs, idempotency, etc
- Presentation protocol: not concerned with what the procedure implements



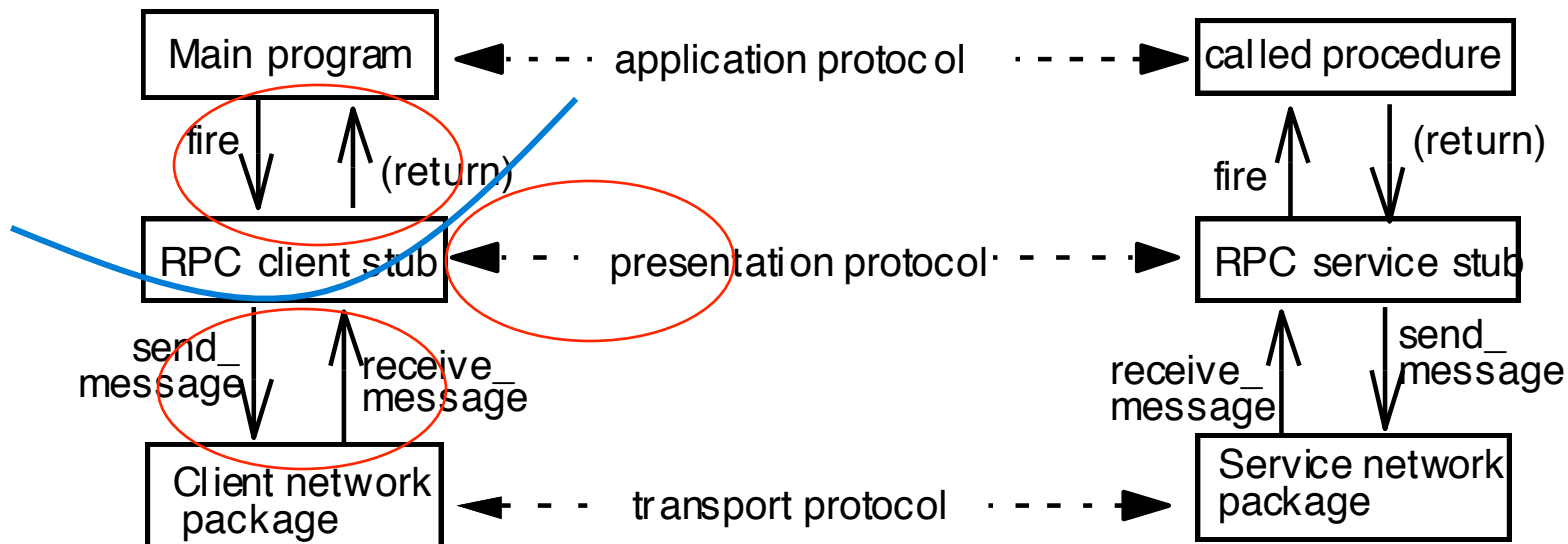
Layering protocols

- Client/service stub protocol needs the primitives to send/receive messages over the network
 - Transport protocol



Layering protocols

- Modules interact with three interfaces
 - Upper layer and lower layer
 - Same layer, on remote end
- Module *hides* interfaces from upper layer



Networking Layers

- Protocols can be layered in many different ways, and different networks will implement layers differently
- One well-known reference model (OSI, Open System Interconnect) establishes seven layers
 - Physical, data link, network, transport, session, presentation, application
- We will focus our discussion on a simpler model that helps highlight key system design challenges and principles

3-layer model

- From bottom up:
 - Data link layer:
 - Responsible for moving data from one point to another through a link
 - Network layer:
 - Responsible for forwarding data through intermediaries between source and destination
 - End-to-end layer:
 - Everything else at the endpoints necessary to support an application's functionality

3-layer model

- From top down:
 - End-to-end layer:
 - Applications need to transmit *messages* or establish *streams* with other endpoints
 - Enforces communication semantics
 - E.g. RPC, at-least-once
 - Requests for messages/streams, breaks them into smaller *segments*; call upon lower layer to send each segment
 - Copes with lost or duplicate segments; places segments in order
 - Applications can be thought as 4th layer, and end-to-end thought as multiple layers, but interfaces are not always clear-cut
 - Example: NFS, RPC over TCP or UDP

3-layer model

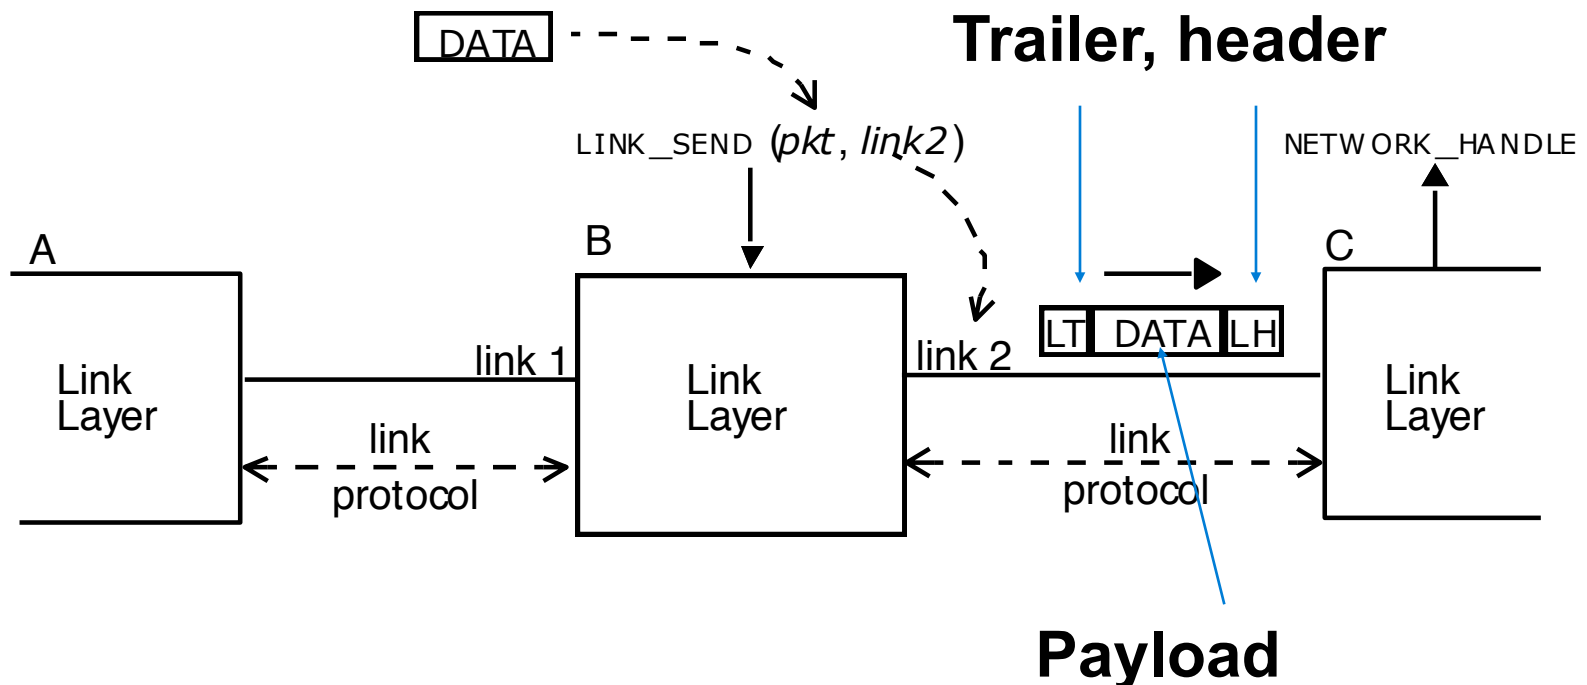
- From top down:
 - Network layer
 - Accepts *segments* from the end-to-end layer
 - Creates packets and transmits them over the network, choosing which links to follow, from source to destination
 - Example: IP
 - Link layer
 - Accepts *packets* from the network layer
 - Creates and transmits *frames* along a single link (between forwarders and/or network attachment devices)
 - Example: Ethernet

Link Layer

- Handles linking devices in a network segment
 - Node to switch, switch to switch
 - Move bits of a packet through a link, hiding physical mechanisms involved
- Simplified interface:
 - LINK_SEND (data_buffer, link_identifier)
 - Data_buffer: reference to memory containing data to be transmitted
 - Link_identifier: names one out of multiple possible links a device is connected to

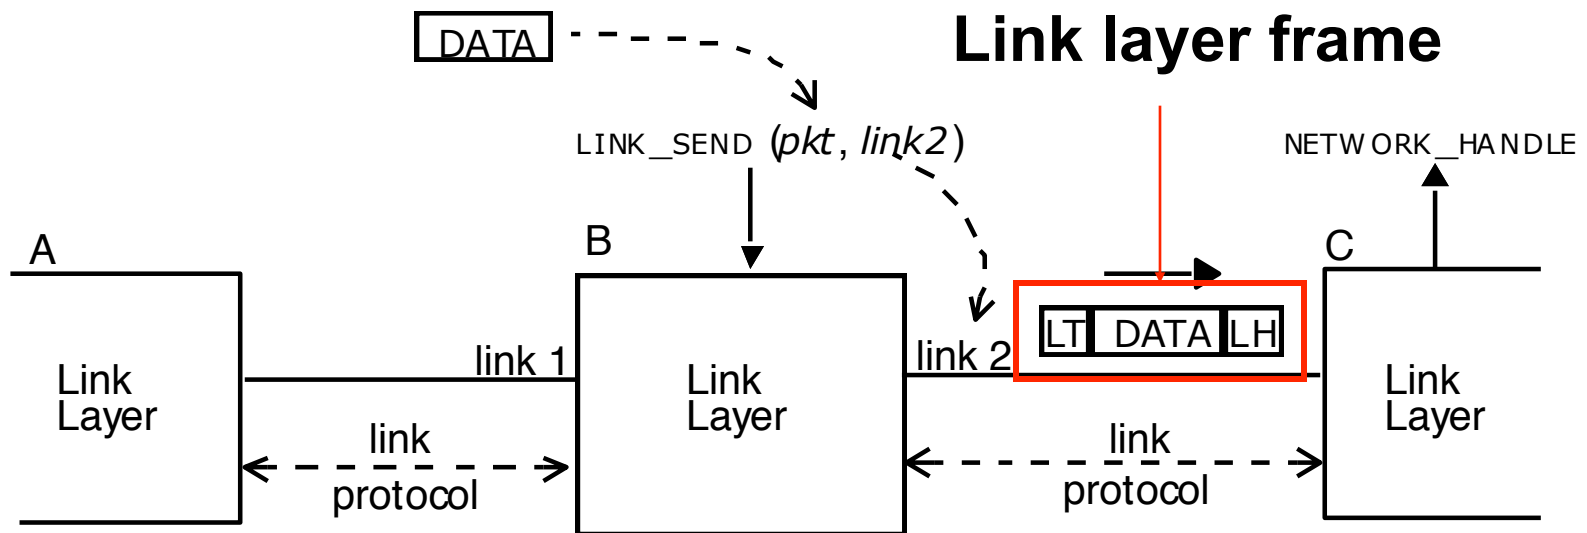
Example

- Node B has two links
- Forwarding A \leftrightarrow C requires SEND to name two links, and a buffer
 - Link identifier is *local* to device




Example

- Header, trailer add information to frame
 - E.g. which network protocol to be used;
source/destination link addresses;
checksum for error detection

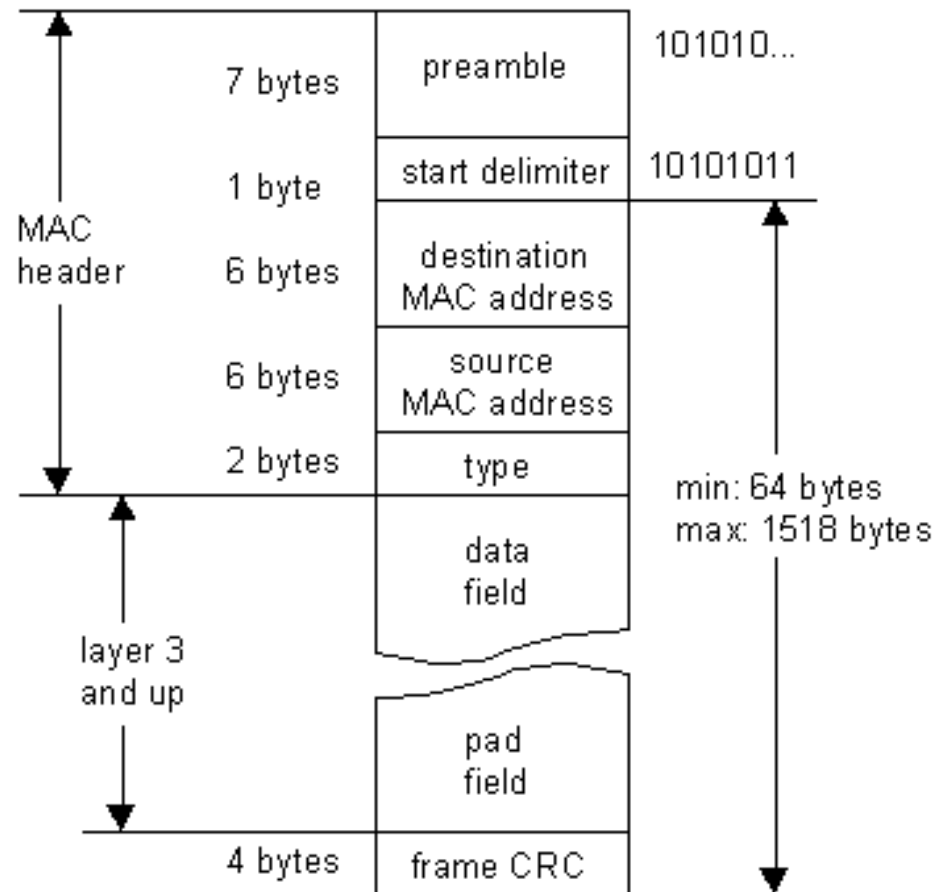


Receiving data

- Two approaches
 - Network layer polls link layer with a RECEIVE call to see if a frame has arrived
 - Link layer “up-calls” network layer when frames arrive – NETWORK_HANDLE



Example - Ethernet



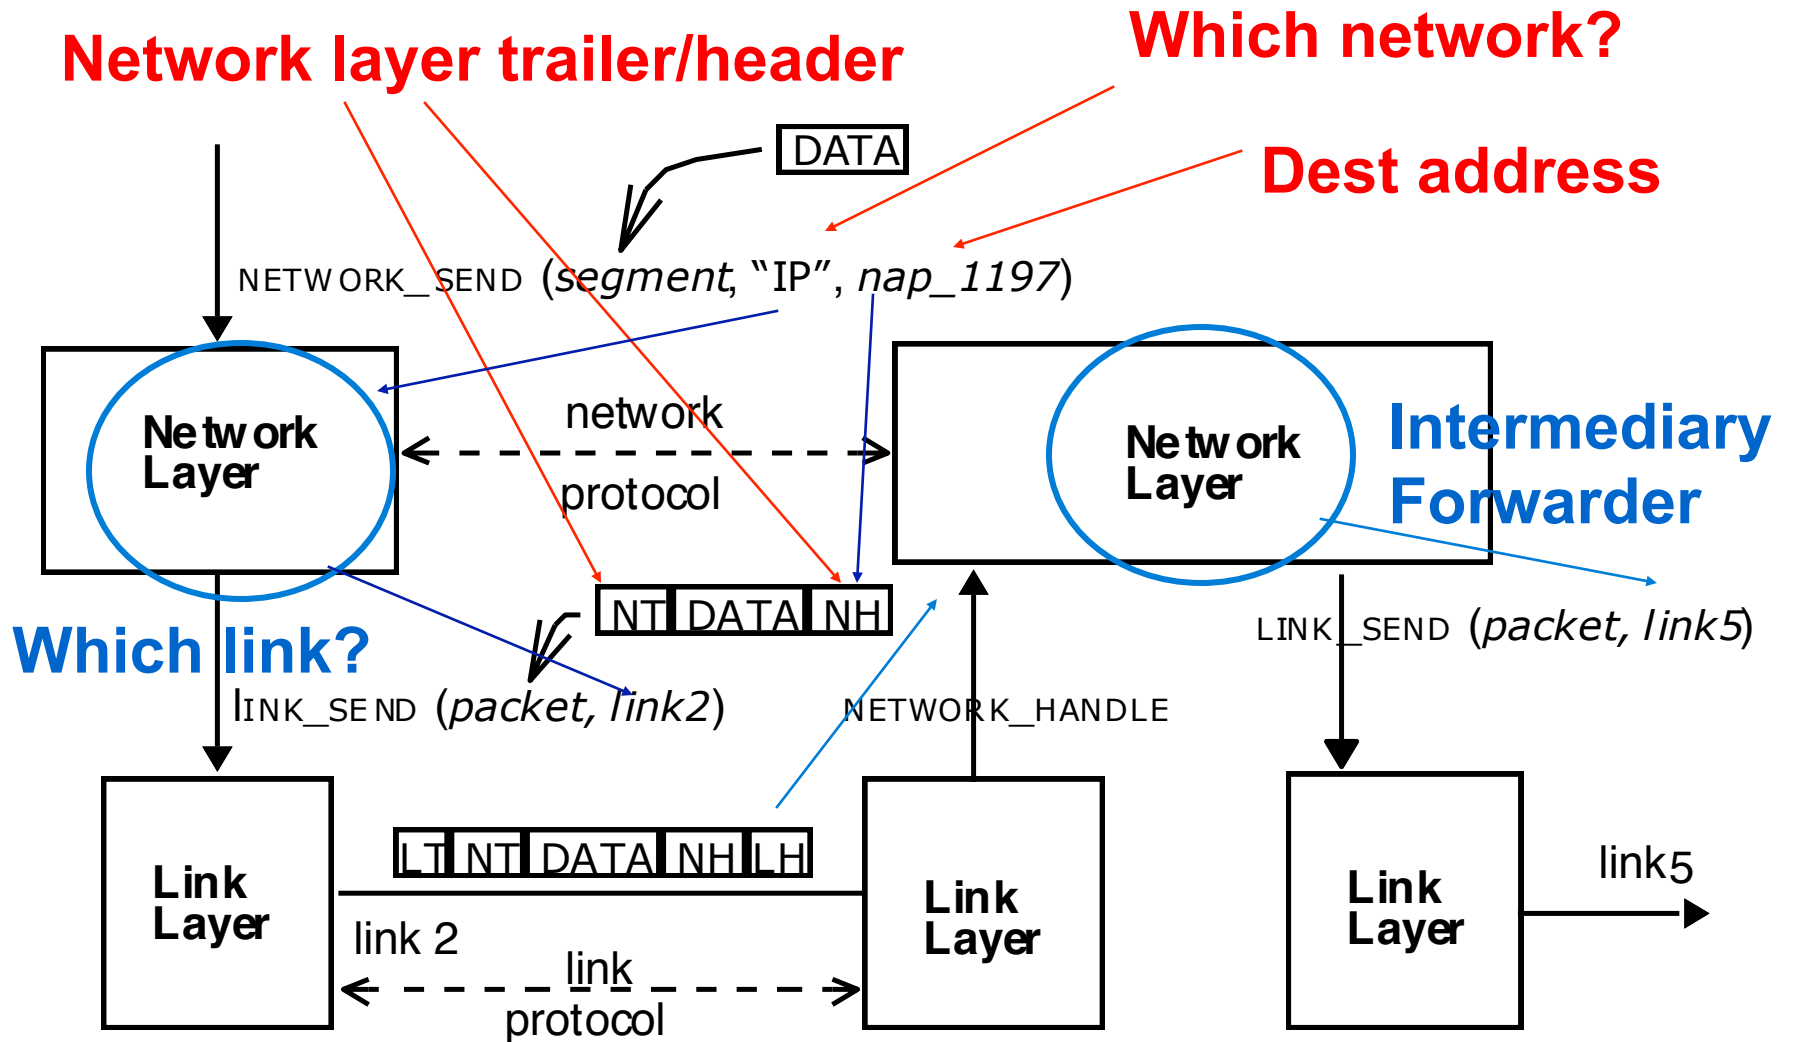
Network Layer

- A segment enters the network at a network attachment point (AP), and is received at another AP
 - Generally, need to traverse multiple links
- Network layer provides
 - Systematic naming of APs
 - Creating a packet
 - Determine which links to traverse and forwarding along this path

Network Layer

- Interface:
 - NETWORK_SEND (segment_buffer, network_identifier, destination)
- Segment buffer:
 - Payload
- Network_identifier:
 - Which network to send packet to
- Destination:
 - Name of destination network attachment point (network-layer address)

Network Layer



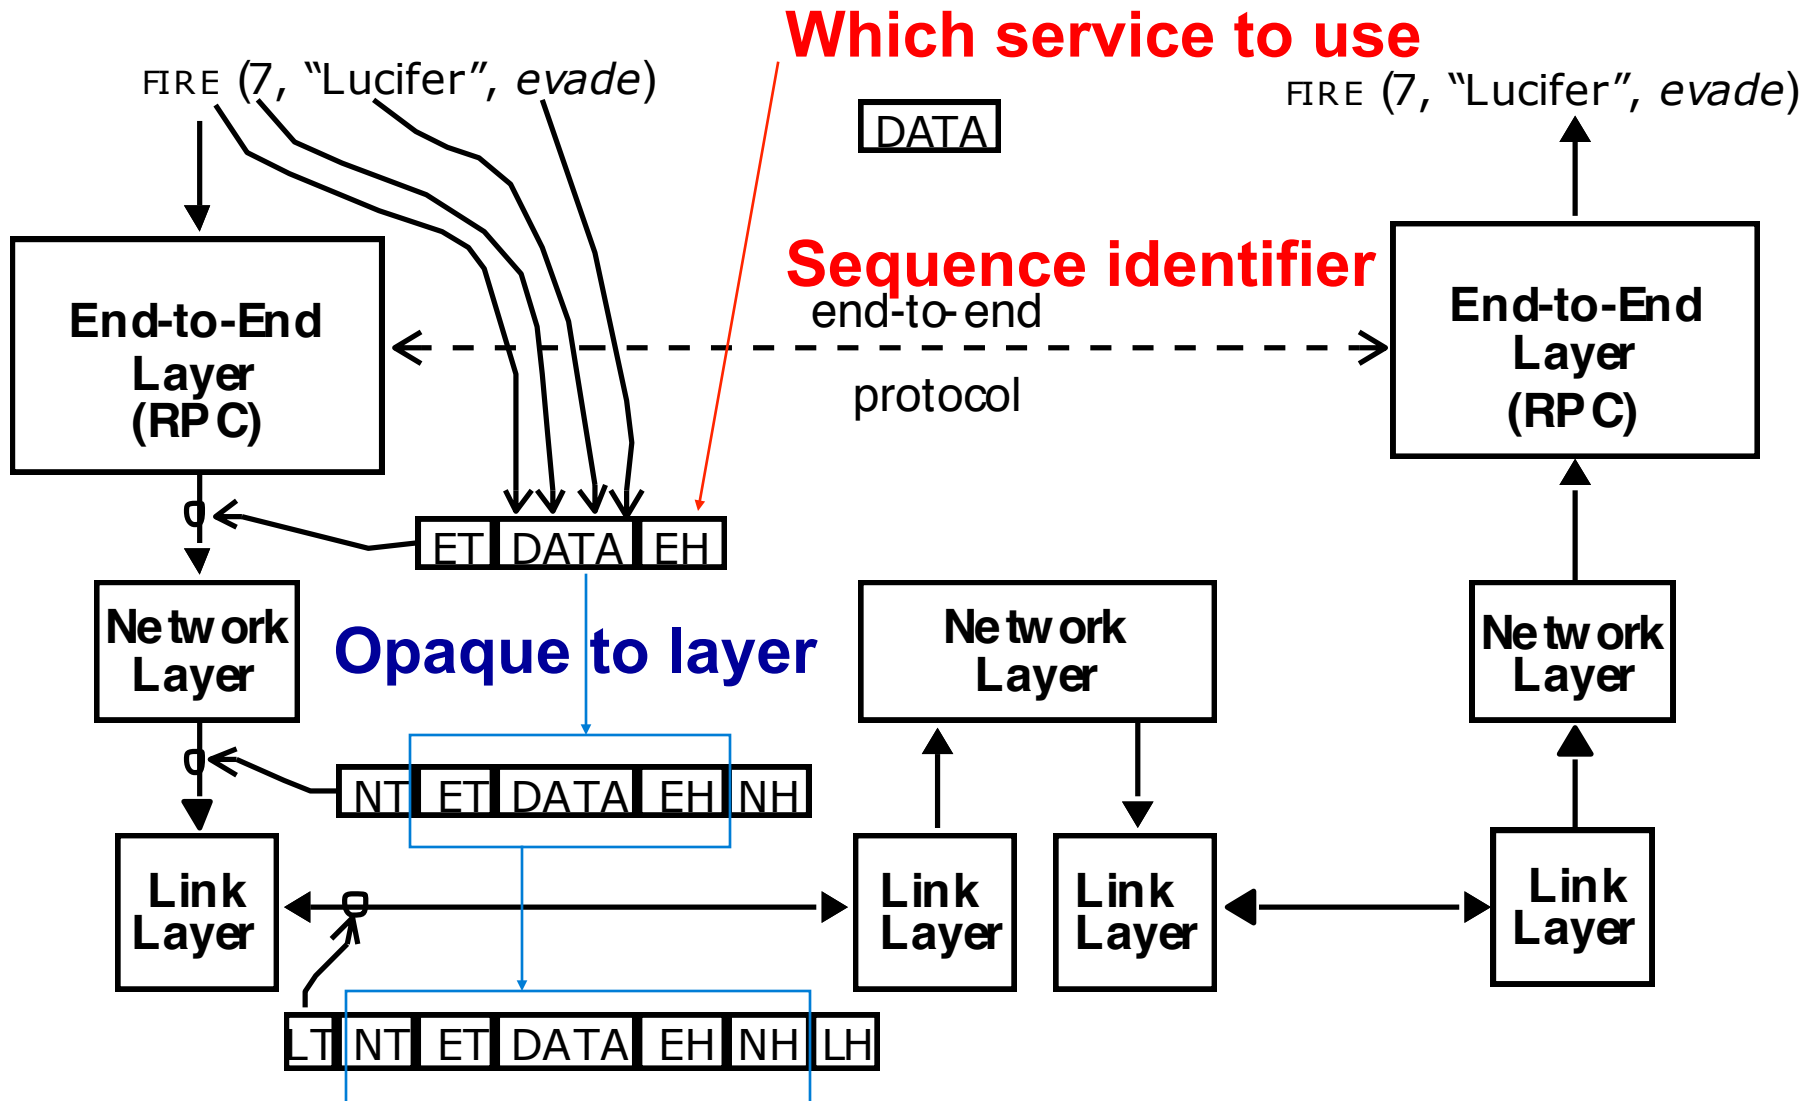
Example – IP packet header

0	4	8	16	19	31
Version	IHL	Type of Service	Total Length		
Identification			Flags	Fragment Offset	
Time To Live		Protocol	Header Checksum		
Source IP Address					
Destination IP Address					
Options					Padding

End-to-end Layer

- Network, link layers: best-effort contract
 - Packet loss, duplicates, out-of-order, errors
- End-End layer provides an easier to use interface for applications
- Example - RPC provides:
 - Presentation: Translates data formats, emulating semantics of procedure call
 - Session: Negotiating discovery and binding to locate and prepare to use service
 - Transport: Dividing streams/messages into segments, dealing with packet loss, duplicates, out-of-order

End-to-End Layer



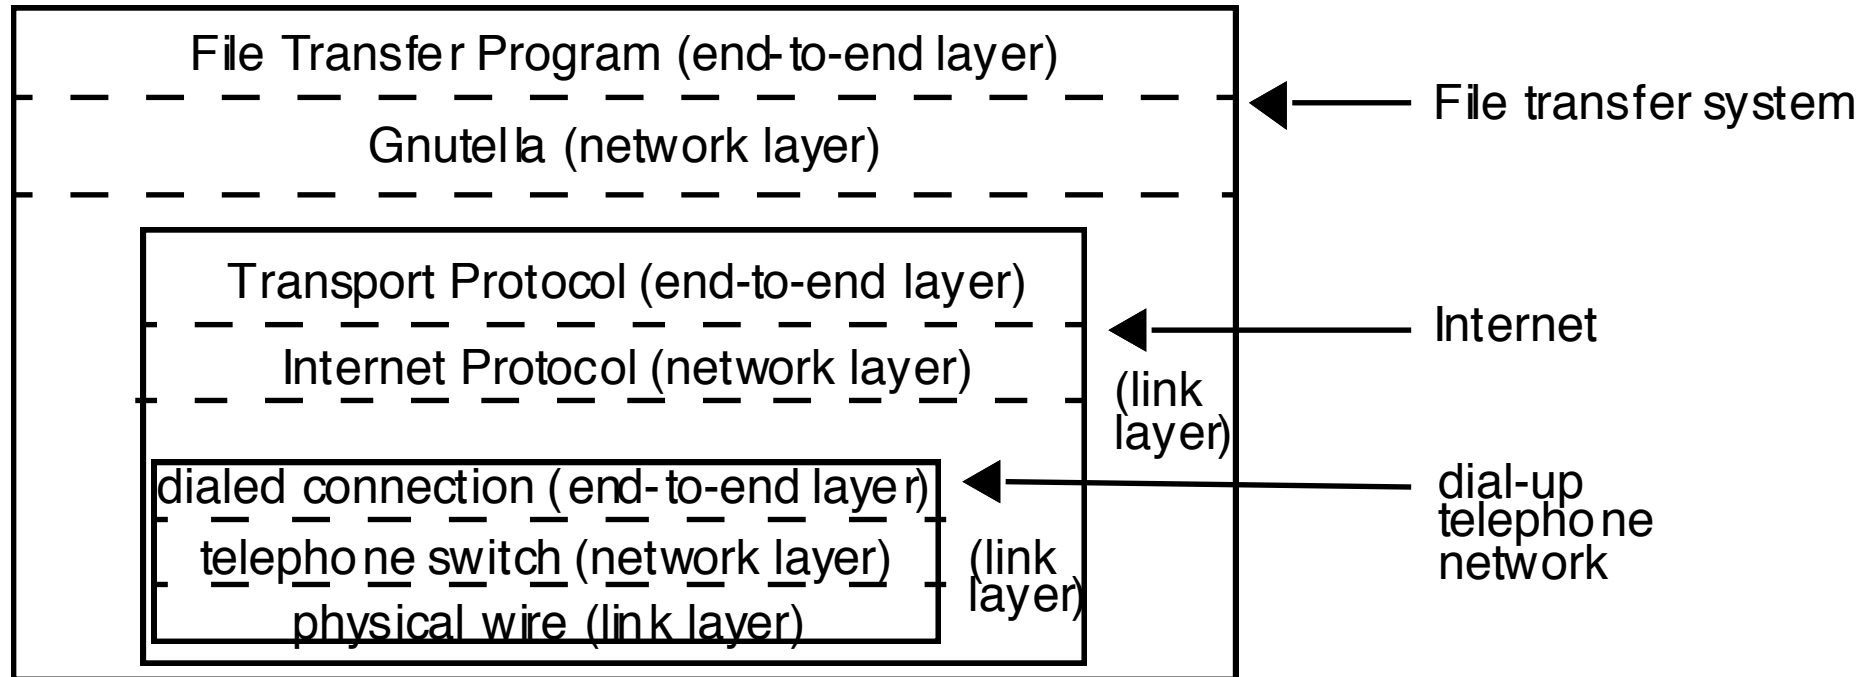
Example - Internet

- Various data links
 - E.g. dial-up, Ethernet, DSL, ...
- Network layer – Internet Protocol
 - IP version 4
 - IP version 6
- End-to-end
 - Transport:
 - TCP – connection-oriented
 - UDP – datagram-oriented
 - RTP – ‘real-time’ (voice, media streaming)
 - May be layered on top of UDP
 - Application
 - File transfer protocol (FTP); HTTP; SMTP, POP

Recursive composition

- Network layer rests on a link layer that itself is a layered network
 - Dial-up phone line
 - A 3-layer network of its own
 - Used as a link layer for Internet
 - Application-level overlays
 - E.g. Gnutella
 - Use end-to-end Internet layer for its links

Recursive composition



Mapping requirements to layers

- Different applications - different requirements
 - Voice:
 - Steady stream, data rate supporting audio quality
 - Low latency/jitter
 - Ok to lose small fraction of data (lost packets, data errors)
 - Reliable transfer of large files
 - High throughput
 - Guaranteed data integrity
 - Out-of-order, duplicate packets ok
 - Financial transaction/trading application
 - Low latency – time is money
 - At-most-once semantics, in-order delivery

Mapping requirements to layers

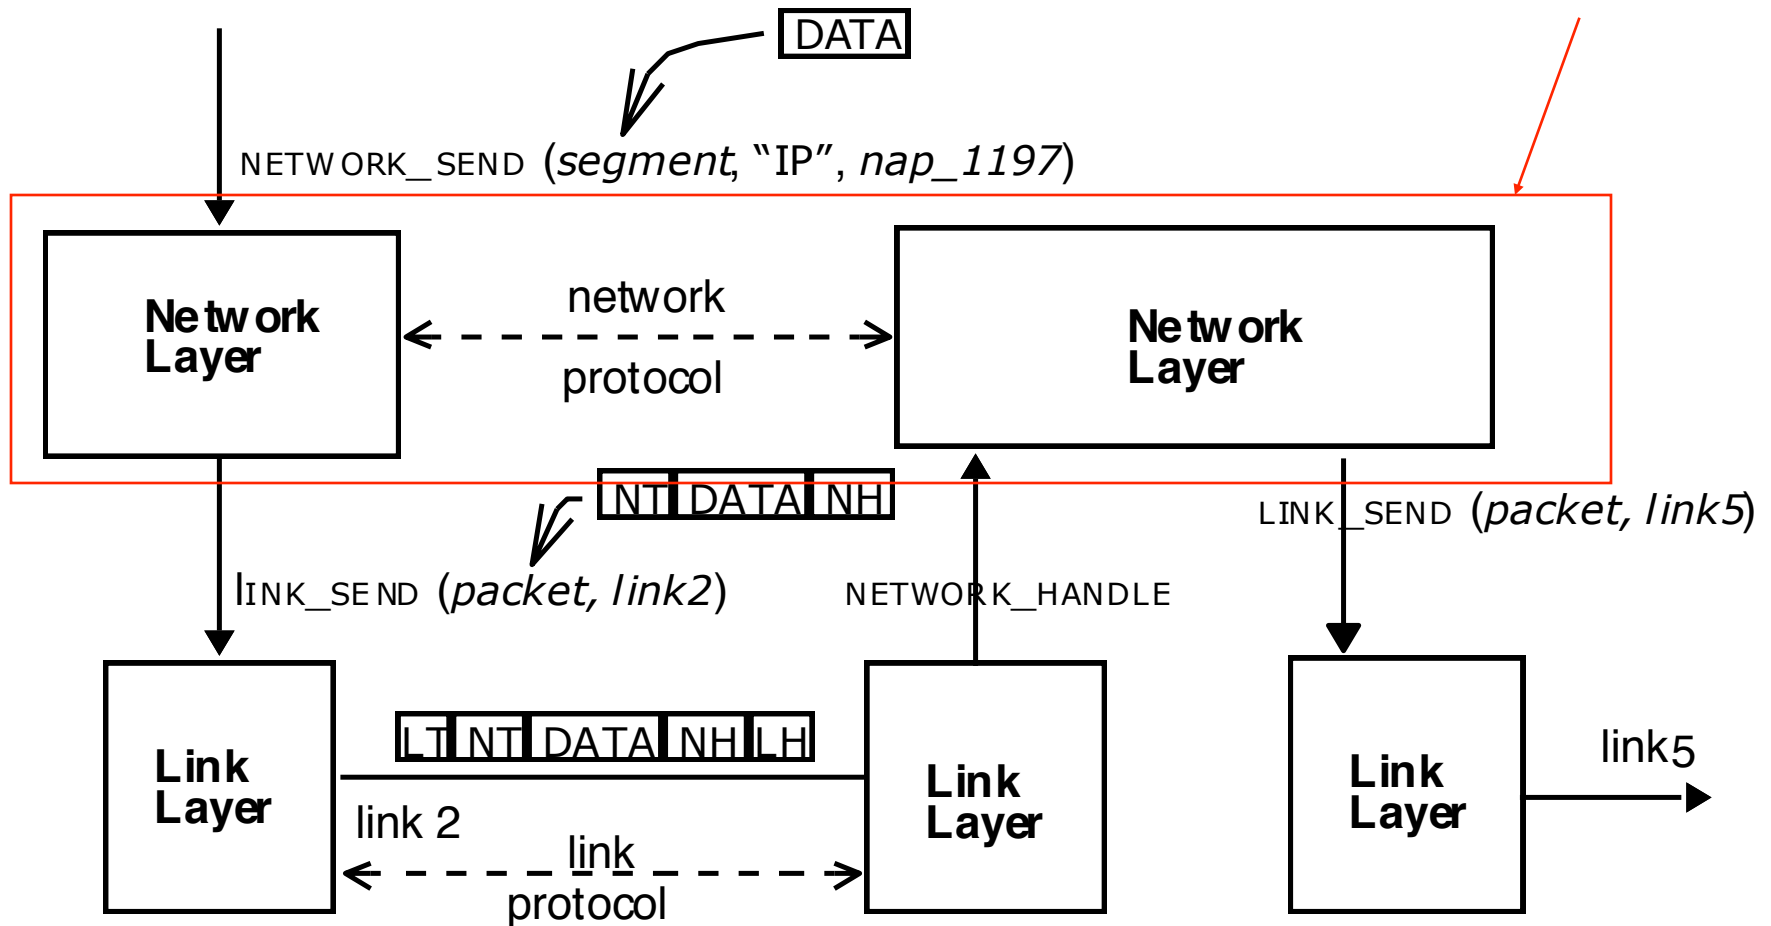
- Data link and network layers provide core functionality
 - Move data from A to B, possibly through intermediaries
- What else provided at these layers to support application requirements
 - Error detection/correction?
 - Dealing with duplicates?

End-to-end principle

- Design principle: “the application knows best”
 - Communication requirements are application specific
 - There are many applications that have similar requirements (e.g. RPC), but no single class covers all applications
 - Implementing support for features other than data movement at lower layers may not suit what is needed by application
 - Applications that do not use feature - inefficient
 - Still require applications for which feature is not quite what they need to re-implement

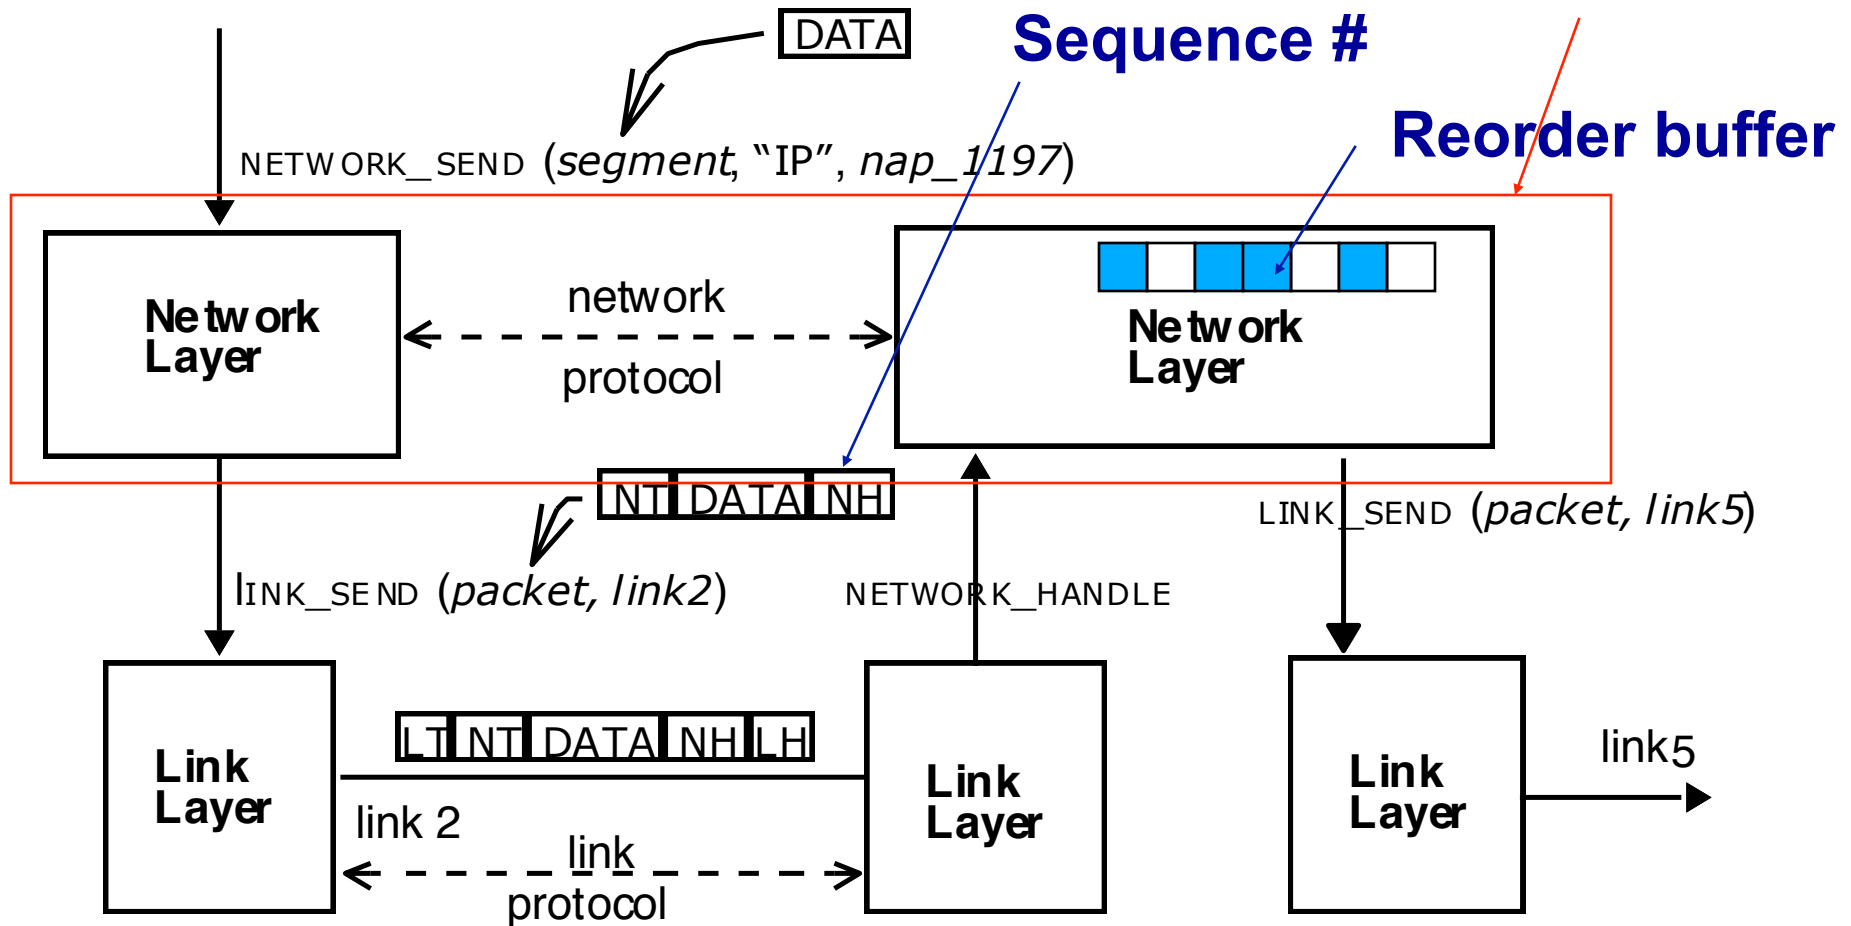
Example

In-order delivery?



Example

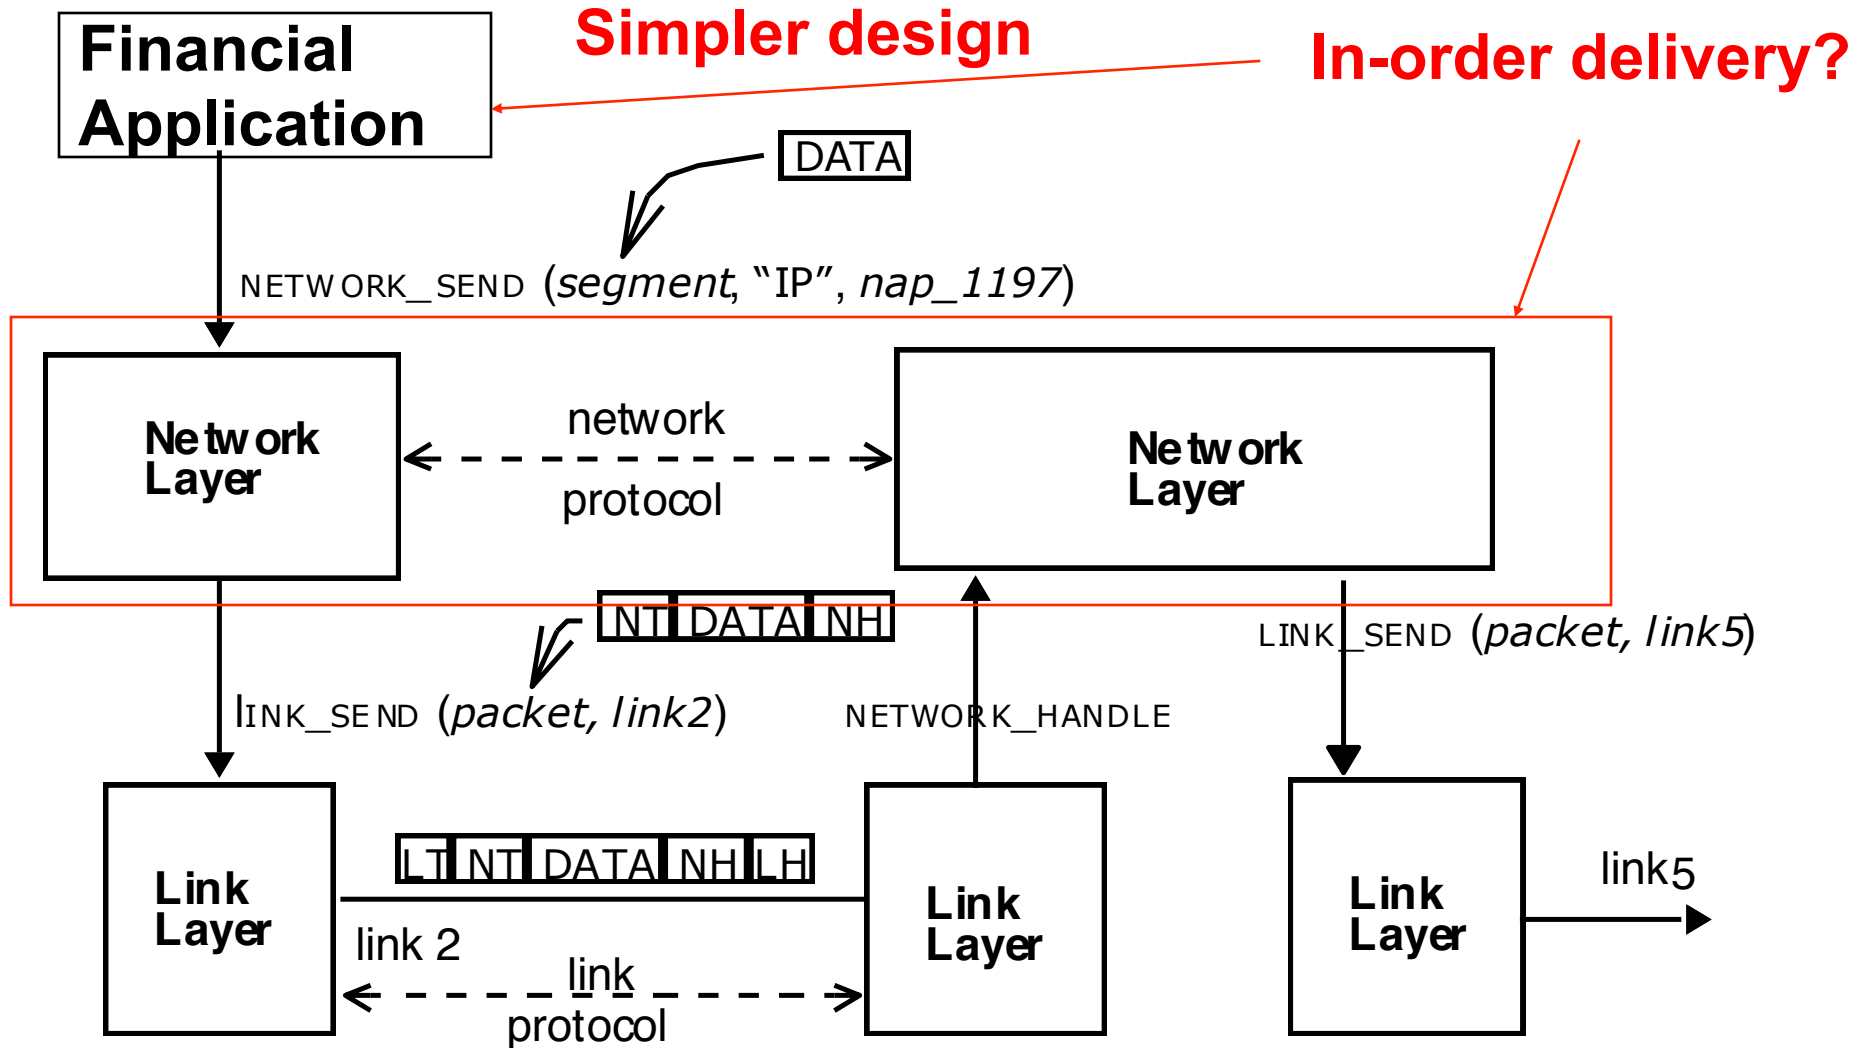
In-order delivery?



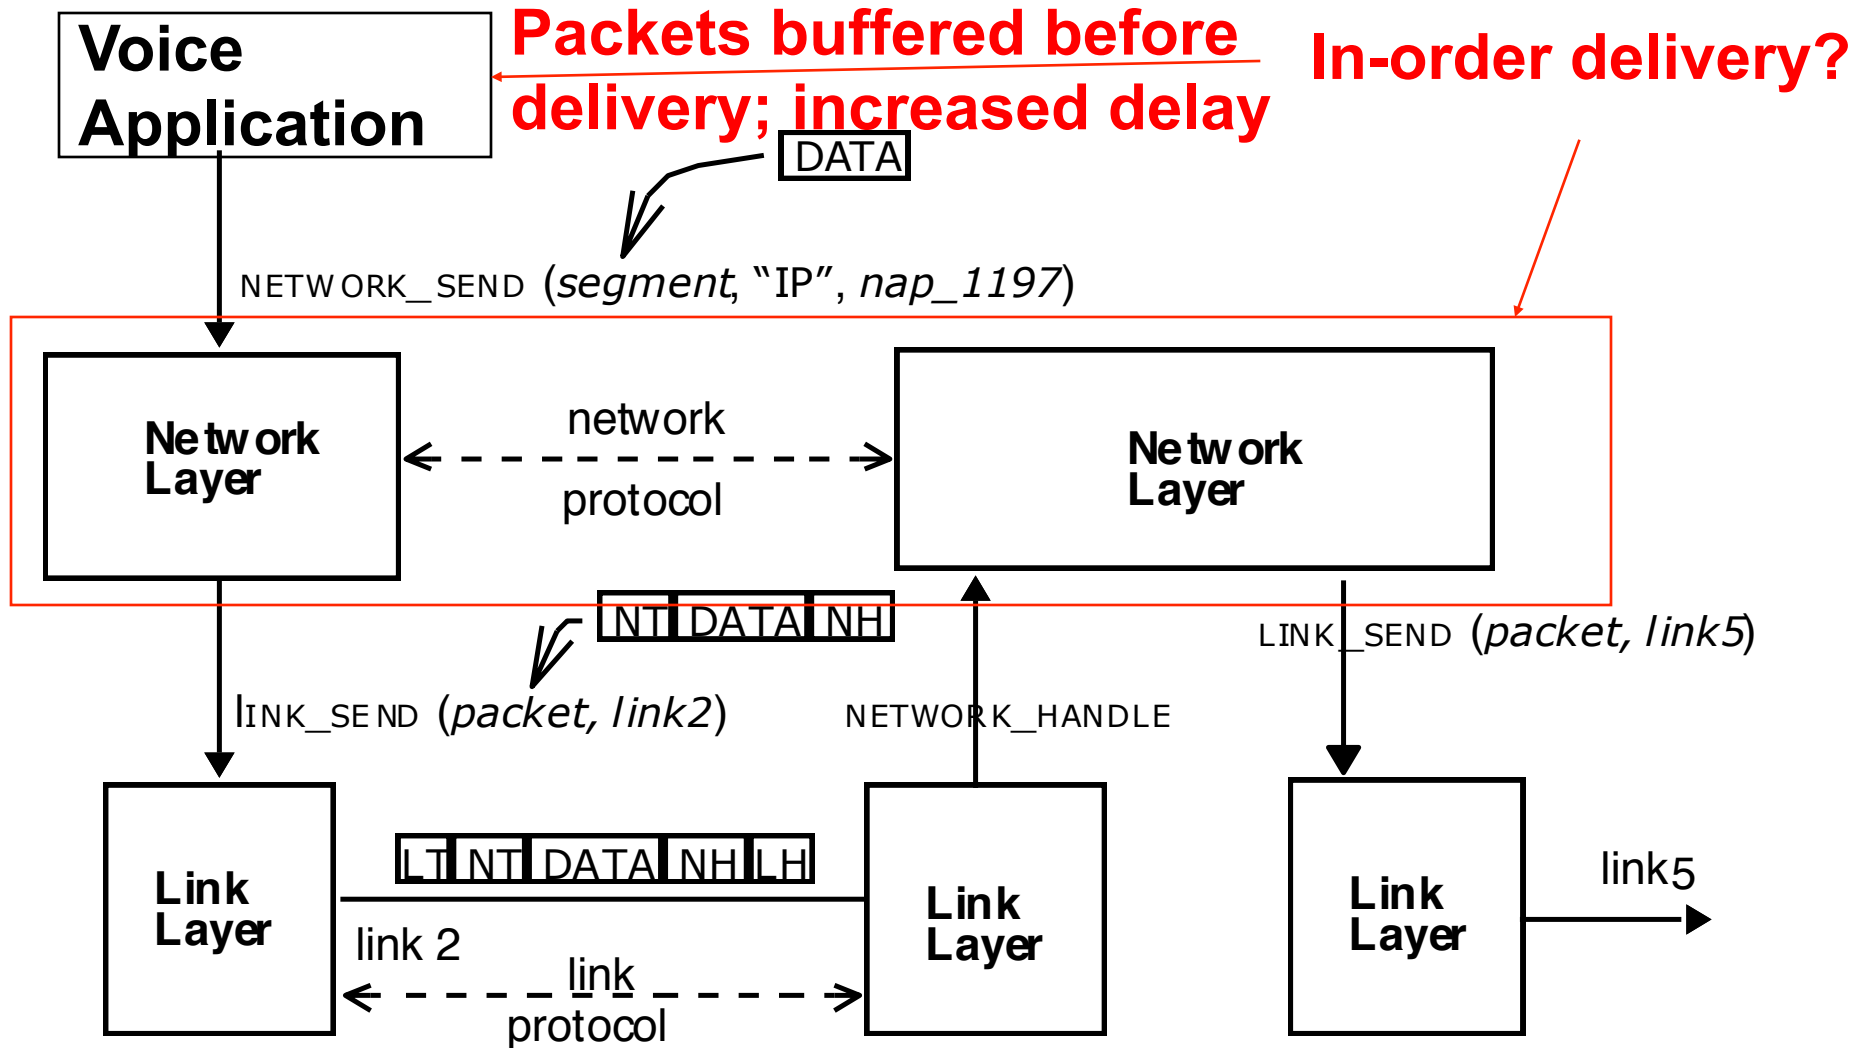
In-order delivery

- In-order delivery in network layer
- Advantage:
 - Design of modules in layer above can be simplified – assume in-order delivery
- Disadvantages:
 - Additional cost associated with every message to guarantee in-order delivery
 - Additional header information, e.g. sequence #
 - Additional code complexity, runtime
 - Additional buffering/delay for packets

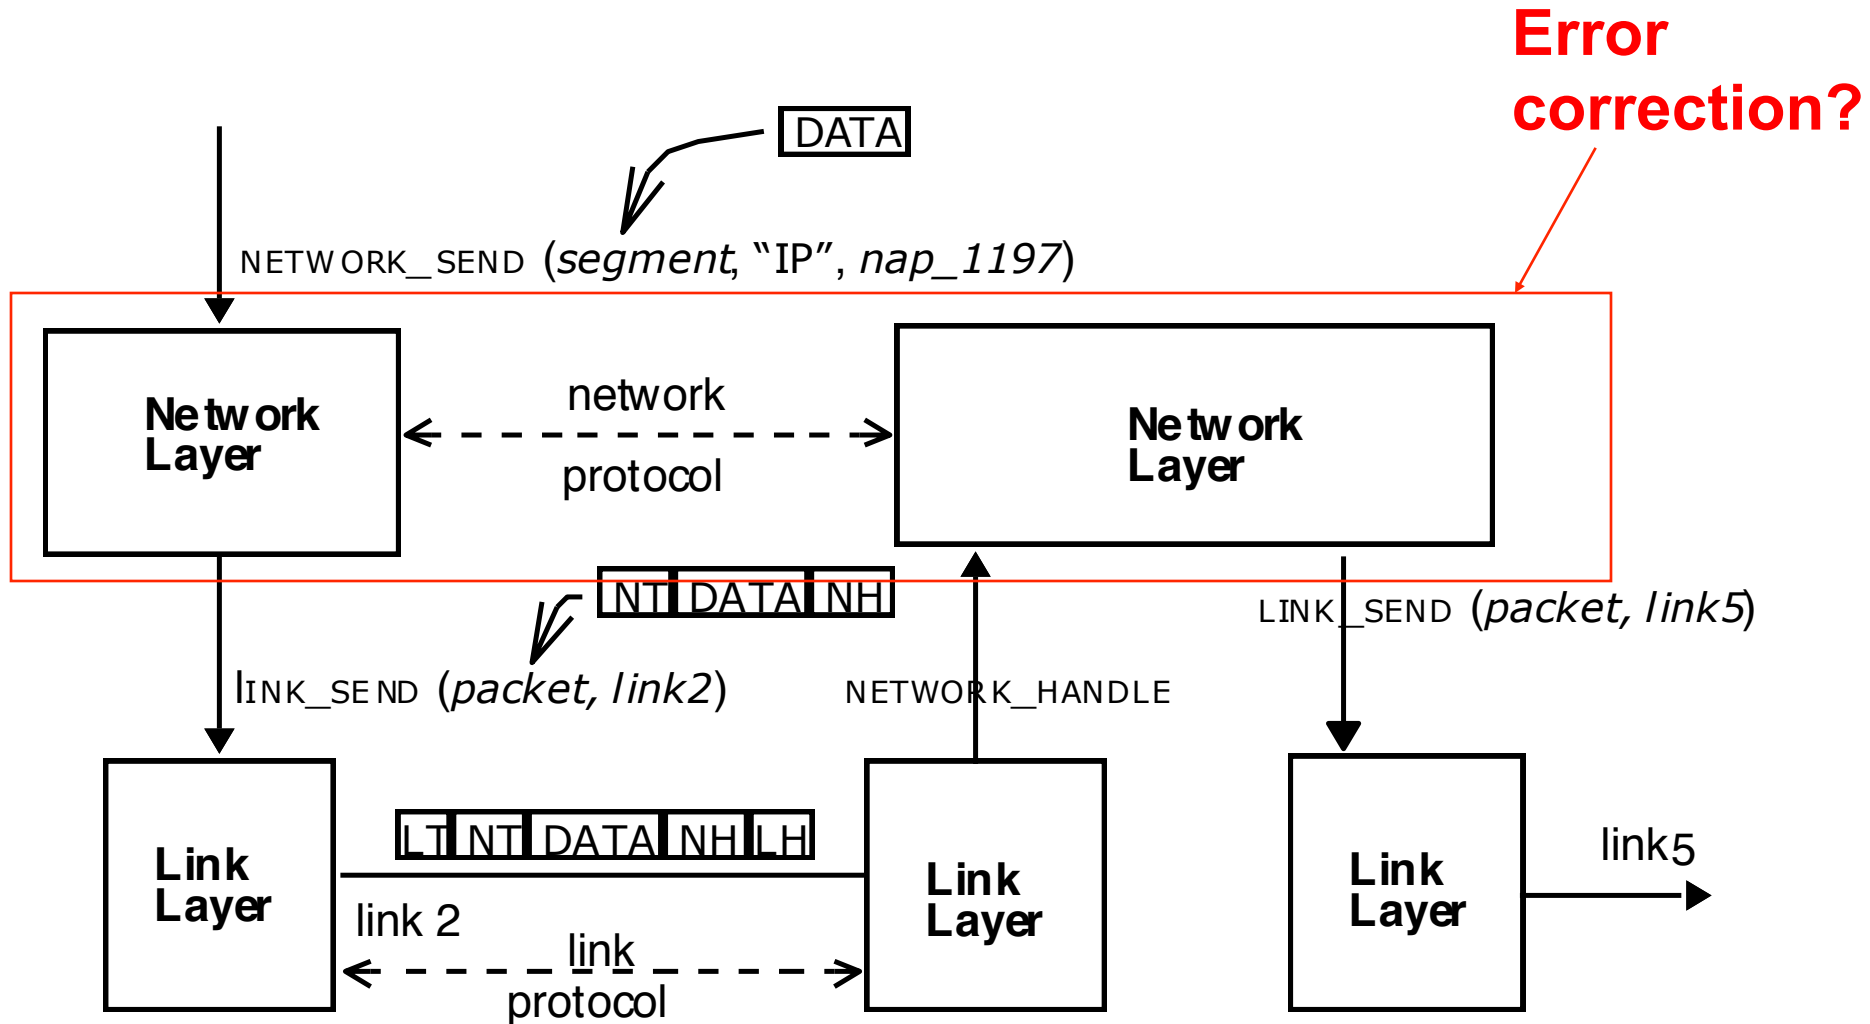
Example



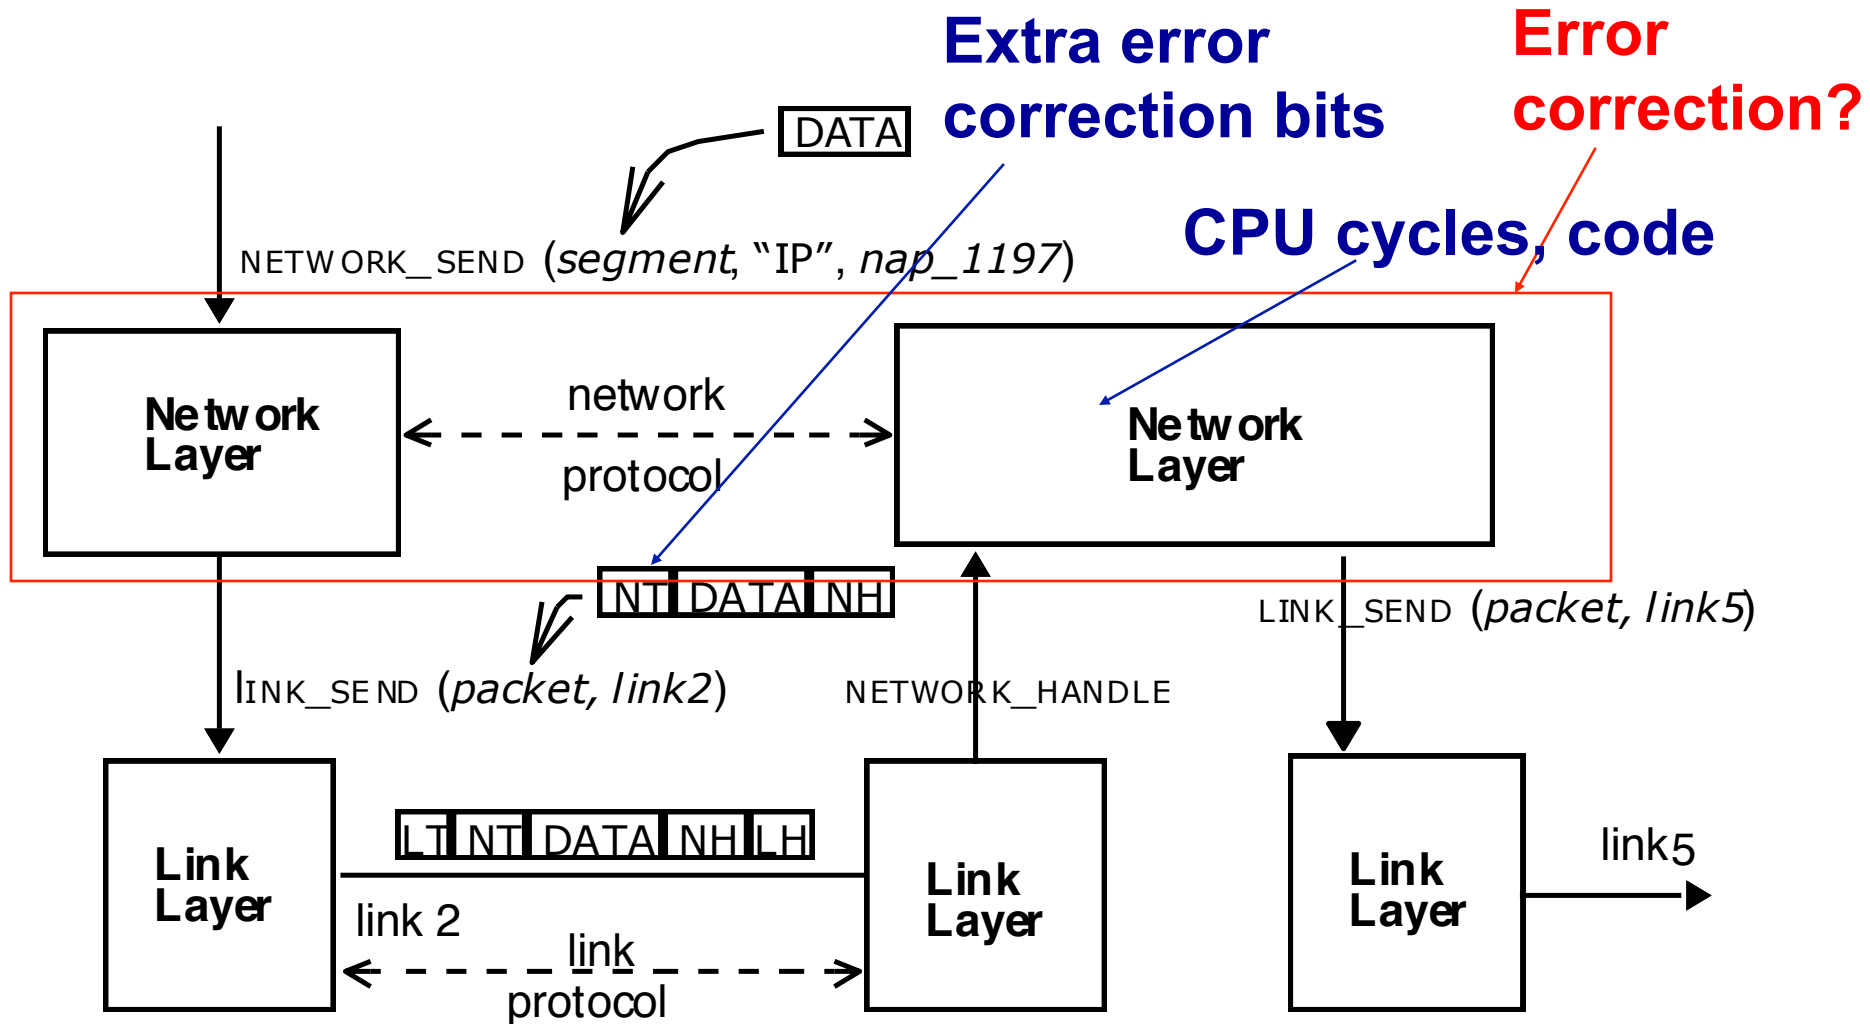
Example



Example



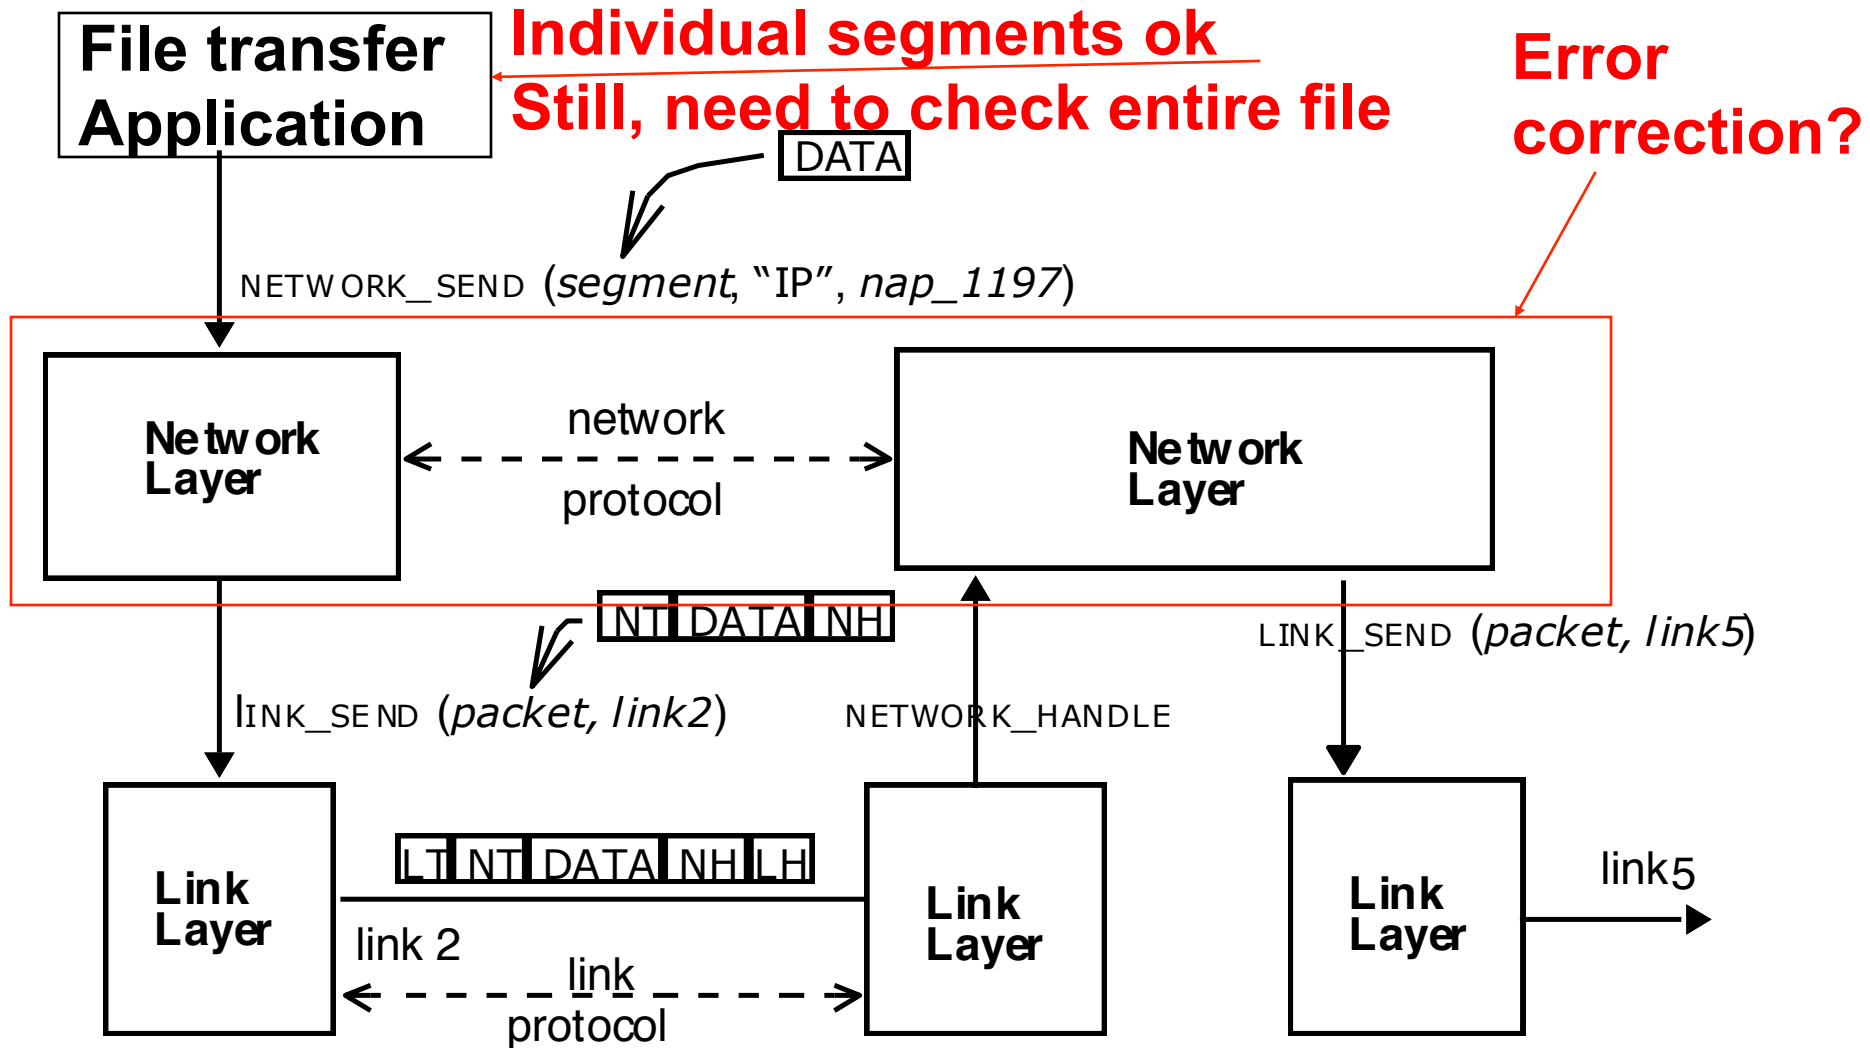
Example



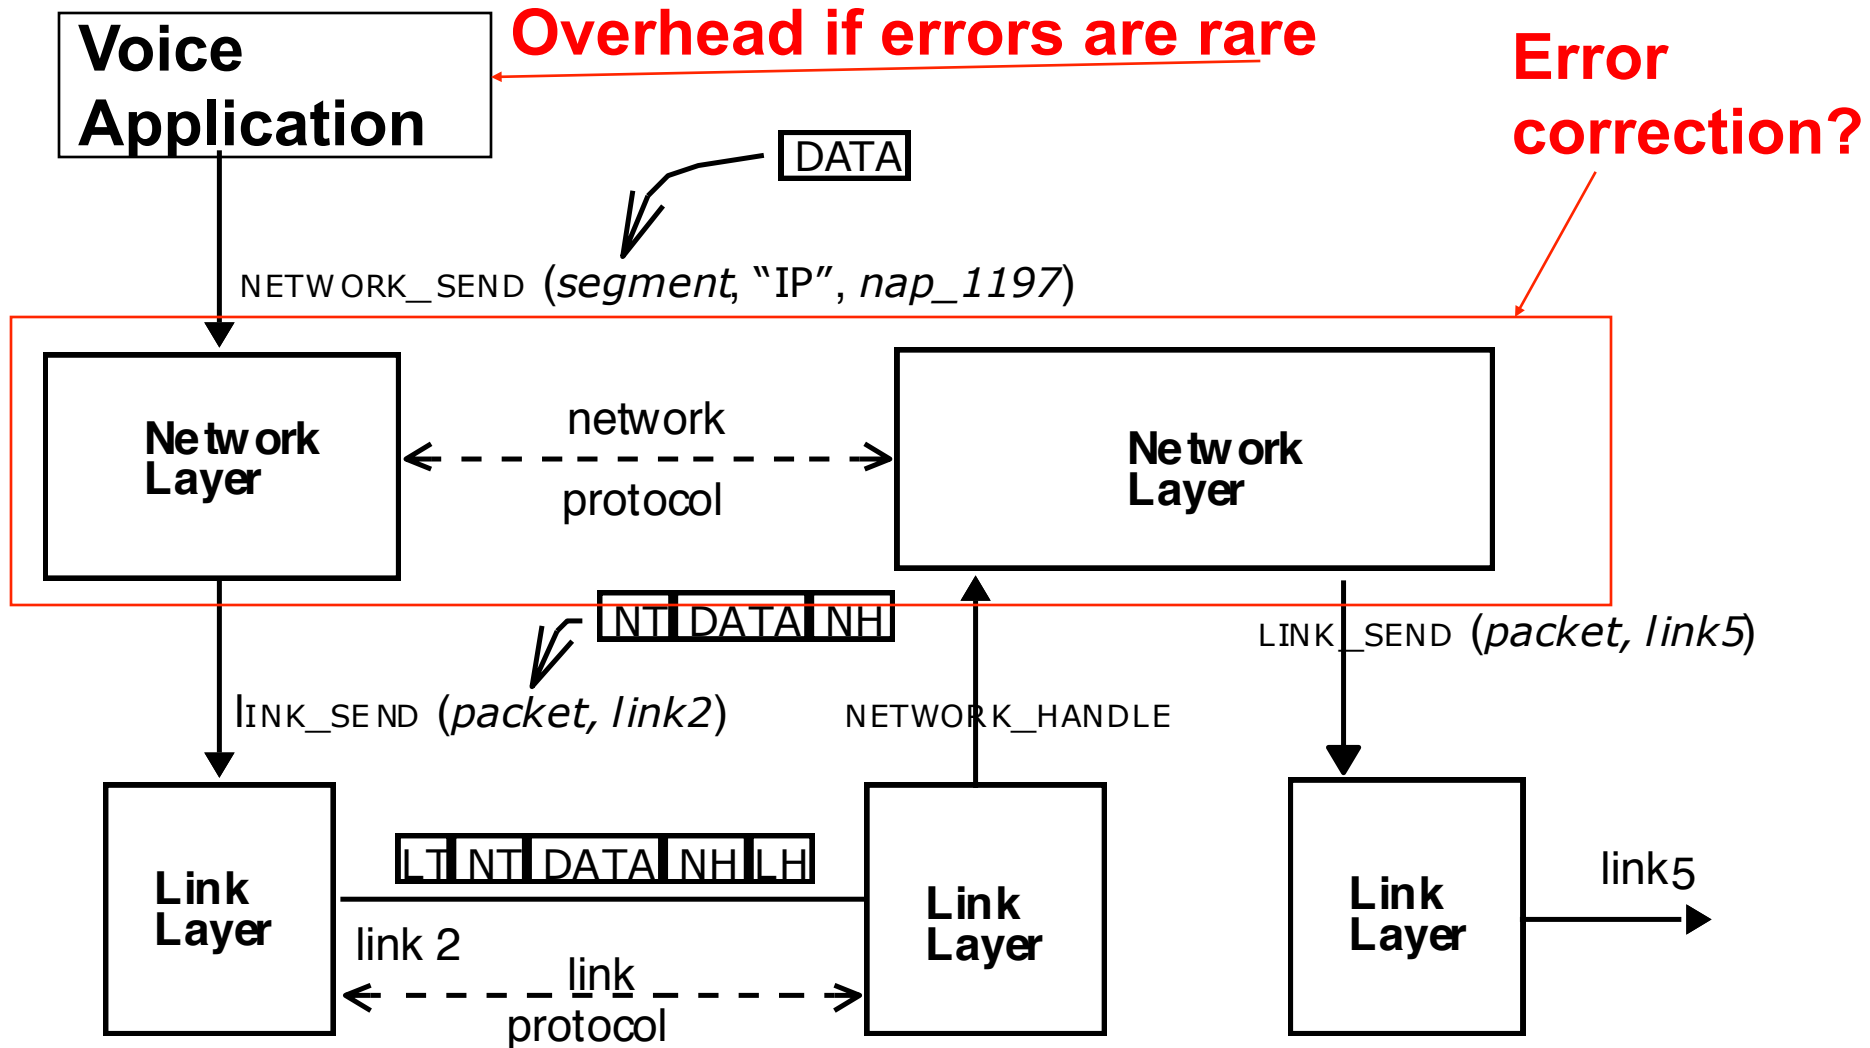
Error correction

- Error correction in network layer
- Advantage:
 - Design of modules in layer above can be simplified – assume if a packet arrives, it has no errors
- Disadvantages:
 - Additional cost associated with every message to guarantee in-order delivery
 - Additional trailer information, ECC
 - Additional code complexity, runtime
 - When correction code not sufficient, must still retransmit
 - Additional code complexity; delay if slows down other packets (e.g. in-order delivery)

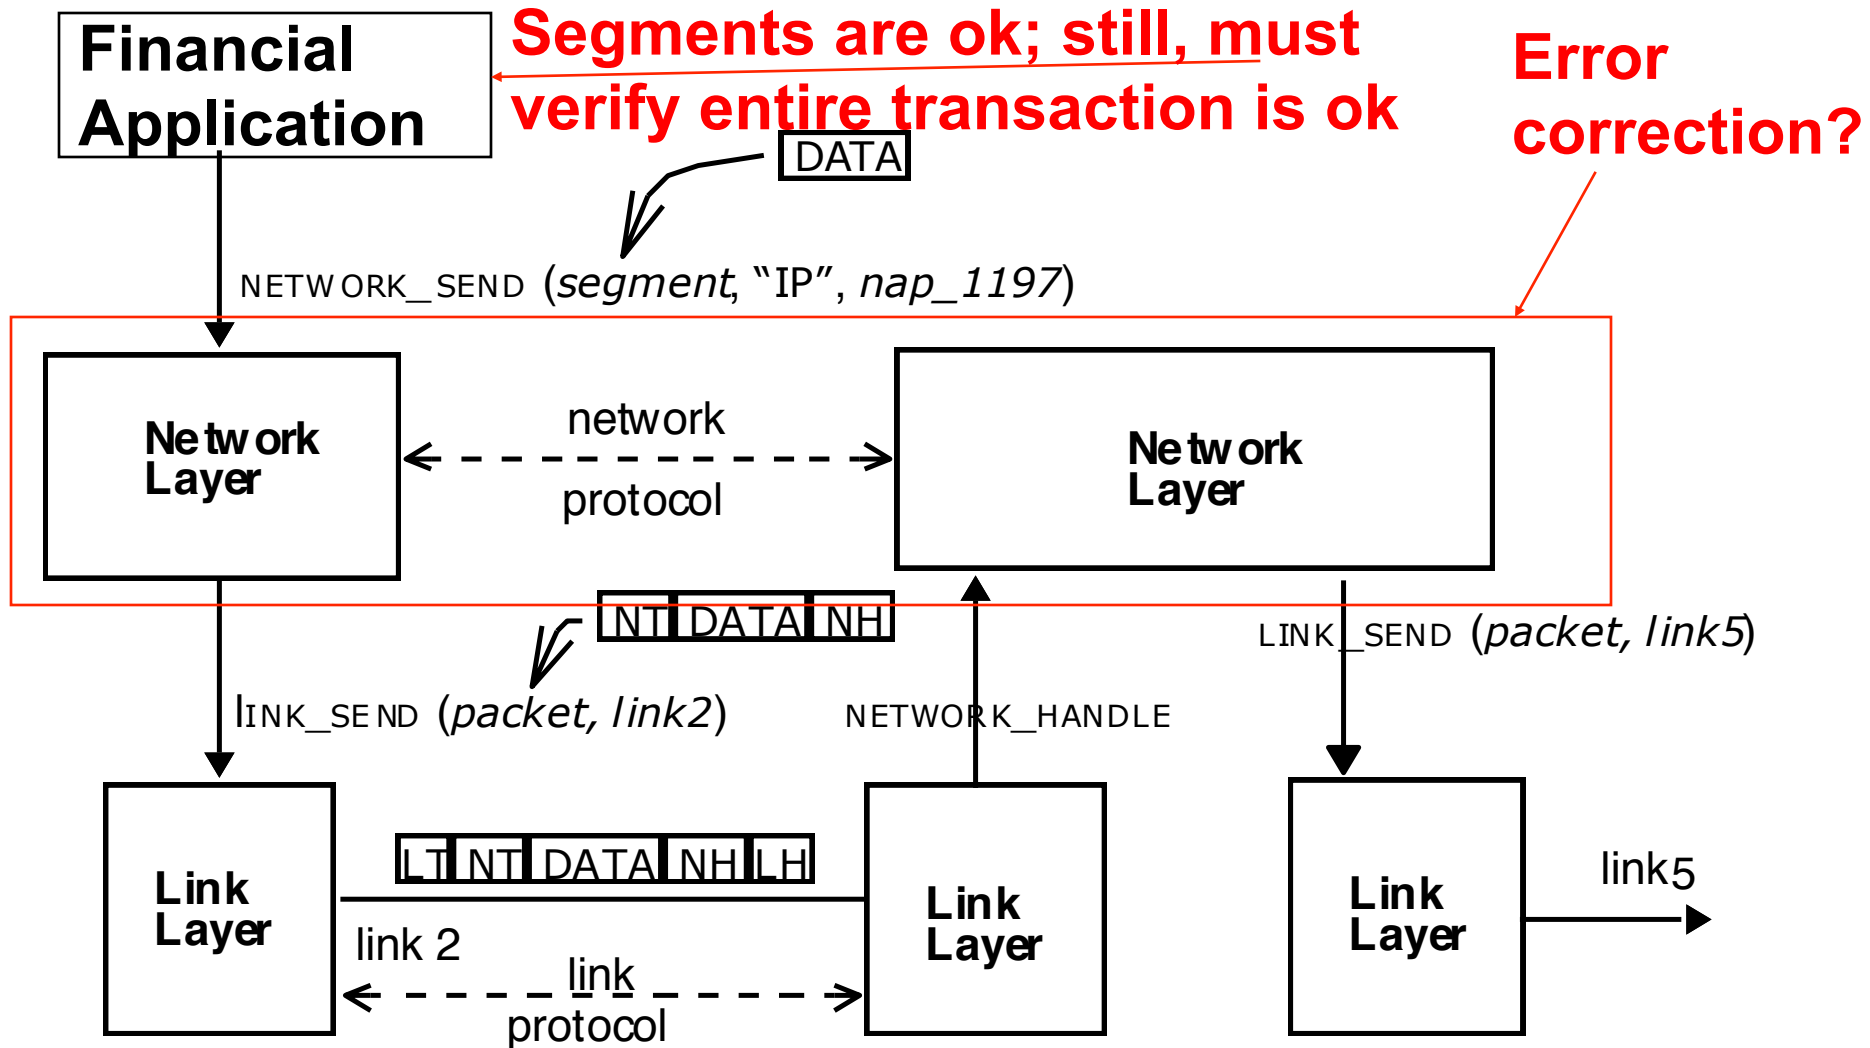
Example



Example



Example




File transfer example

- Reliable file transfer application
 - Even though packets are guaranteed to be error-free by network layer, many other errors can occur as the application interacts with other systems
 - Crashes during transmission while buffers still in memory
 - Faults in hardware – e.g. bad block, bit flip
 - Important to guarantee integrity of the file stored in disk
 - E.g. compute an “end-to-end” checksum of entire file on source, verify on destination after stored in disk
 - Correcting an error may require re-transmitting the whole file, regardless of error correction support at network layer

File transfer example

- Each individual source of error might be dealt with at its own layer to lower probability of its occurrence
 - Redundancy in disks (RAID)
 - Error correction in memory
 - Error detection, retry in network
- No matter how small their probabilities, errors can still happen, outside control of application
 - If application must guarantee a target probability that errors do not happen, it needs to control this aspect

File transfer application



What if target probability of error being detected is 99.9999999% ?

End-to-end principle

- Avoid implementing functionality that is not strictly necessary at a lower layer
- Note, still may be desired for some features may be applied repeatedly across layers
 - E.g. fault-tolerance:
 - Ethernet has data link layer checksum: error detection for frames
 - TCP has transport layer checksum: error detection for messages
 - Reliable file transfer application: whole file checksum
 - Why?

End-to-end principle

- Example: errors caused by different reasons
 - Data link – noisy transmission
 - Transport – errors in intermediate devices (e.g. bit-flip in router)
 - File – disk errors, crash
- How likely are these errors?
- Different error rates call for different checksum codes
- Dealing with more common errors at lower layer
 - Simpler, faster
 - Avoid potentially expensive upper-layer operations

End-to-end principle

- Application desires to provide its own level of fault tolerance
- Support at lower layers for fault tolerance may reduce likelihood of faults that application will perceive
 - Does not eliminate need for end-to-end error correction
 - Performance enhancement, not application requirement
 - Tradeoffs:
 - Re-sending a frame vs. re-sending a file
 - More bits for all frames vs. more retransmissions

End-to-end principle

- Applies to various aspects of systems design – not just networking
 - E.g. file backups
 - Reduced instruction set (RISC) vs. CISC
- Not an absolute decision technique
 - But an argument/principle that should guide the decision-making process when you are designing modules/layers in a complex system

Reading

- Section 7.3