

Principles of Computer System Design  
Midterm Exam  
Wednesday, Oct 8th, 2014

**NAME:**

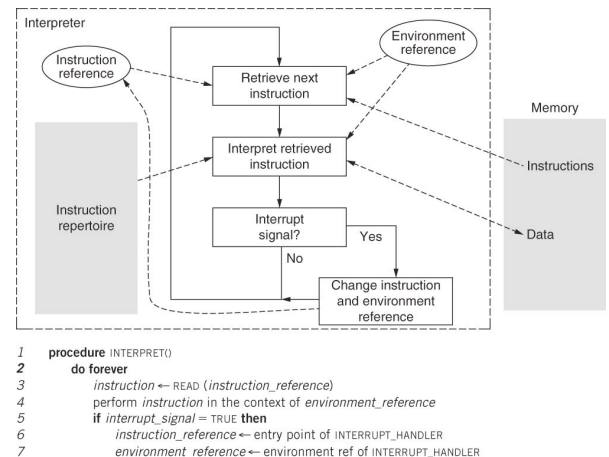
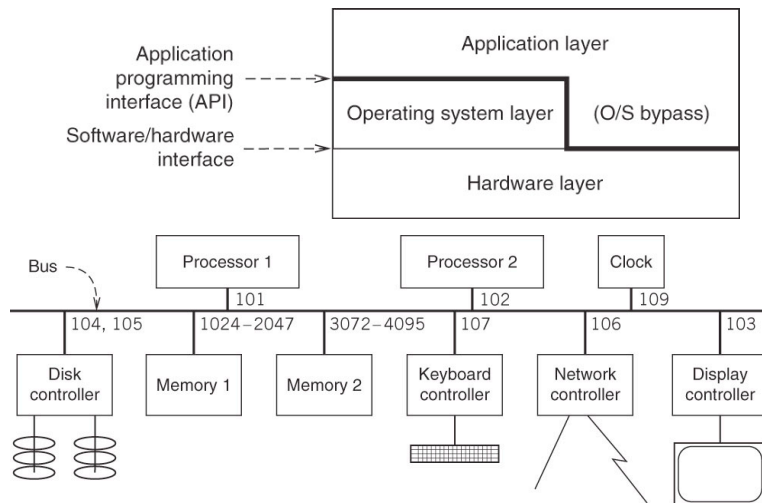
**UFID:**

Please read each question carefully, to avoid any confusion. This exam should have 10 pages printed double-sided (the last 2 pages are blank); before you begin, make sure your copy contains all pages. The exam is closed book, closed notes. Each question has its number of points identified in brackets.

GOOD LUCK!

<b>QUESTION</b>	<b>POINTS SCORED</b>
<b>1 [25]</b>	
<b>2 [25]</b>	
<b>3 [30]</b>	
<b>4 [20]</b>	
<b>TOTAL</b>	

## 1) [25] Naming and layers:



1.a) [10] Suppose each processor above has a single kernel/user bit, and a single memory domain register that enforces bounds (lower L, upper U) and permissions (Read, Write, eXec). Suppose processor 1's domain register has (L=1024, U=1535; perm=R, X). Suppose an application in P1 issues the following instructions; *which of these instructions result in crossing from the application to OS layer? Why?*  
*Notation:* instruction address: opcode, register, data address

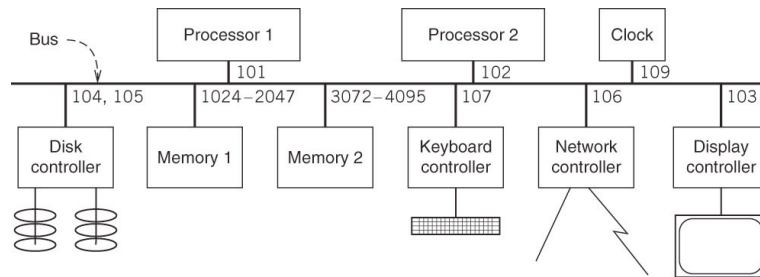
3072: LOAD R1, 1030

1024: STORE R2, 1400

1300: LOAD R3, 1200

1240: SVC

1.b) [15] Now assume that the processor supports page-based virtual memory. Suppose the page size is 1024 bytes. Given the page tables for processors P1 and P2 to be as follows (VA, PA: virtual and physical addresses, respectively).



**P1's page table**

VA	PA	Perm
1024	1024	RW
0	3072	RX
3072	0	R

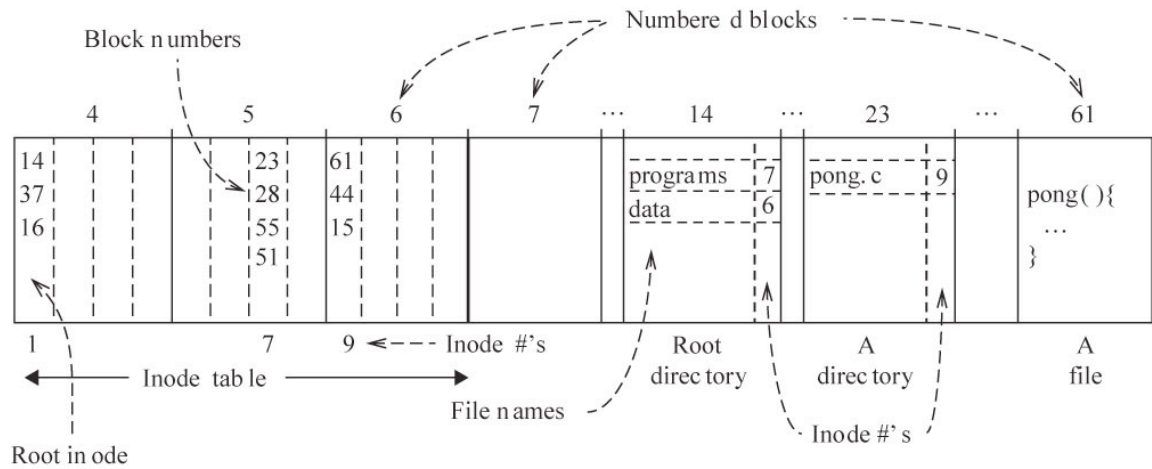
**P2's page table**

VA	PA	Perm
0	1024	RW
1024	3072	RX
2048	0	R

- i) Would it be possible for a thread in P1 to send a message to a thread in P2 using a shared buffer? If so, describe where the buffer and lock need to reside in physical memory
  
- ii) Would it be possible for threads in P1 and P2 to use *memory mapped I/O* to *read* data from the keyboard controller? What about to *send* data over the network interface?

## 2) [25] File systems, RPC, NFS:

a) [10] Consider the figure below discussed in class.



i) Suppose an application opens and reads `/programs/pong.c`. Name *all* block numbers read from disk, and the order in which they are read.

ii) Suppose you create a hard link `"/myhardlink"` that links to `/programs/pong.c`, and a soft link `"/mysoftlink"` that also links to `/programs/pong.c`. Which *i-nodes* and *blocks*, if any, would need to be modified to accommodate these two links? Explain.

b) [15] Alice is using both her desktop D and laptop L to work on a file served by an NFS server S. The NFS file system is mounted on both D and L in /home. Alice opens the file "/home/alice/file.txt" in both computers. Assume there is no caching.

i) What LOOKUP messages are sent to S when the system call to open() the file is issued by D and L? State your assumptions.

ii) Are the generation counts the same in the file handles received by D and L? Briefly explain.

iii) At a later point in time, Alice notices that two subsequent NFS reads issued by D, for the same file handle and offset, result in a successful read followed by a stale file handle error. How could that happen?

**3) [30]** Virtualization and enforced modularity

- i) [5] What must be the permissions associated with the memory domain where the bounded buffer is stored in order to enforce modularity?
  
  
  
  
  
  
  
  
  
  
- ii) [5] What privilege level is the processor on when the RSM instruction is issued to acquire the bounded buffer's lock? Why?
  
  
  
  
  
  
  
  
  
  
- iii) [10] Can a race condition occur in the bounded buffer send/receive with a *single* client and *single* service? Explain

Refer to the YIELD implementation:

```

1  shared structure processor_table[7] // each processor maintains the following information
2      integer topstack                // value of stack pointer
3      byte reference stack            // preallocated stack for this processor
4      integer thread_id               // identity of thread currently running on this processor
5  shared structure thread_table[7]   // each thread maintains the following information:
6      integer topstack                // value of the stack pointer
7      integer state                   // RUNNABLE, RUNNING, or FREE
8      boolean kill_or_continue       // terminate this thread? initialized to CONTINUE
9      byte reference stack            // stack for this thread

10 procedure YIELD ()
11     ACQUIRE (thread_table_lock)
12     ENTER_PROCESSOR_LAYER (GET_THREAD_ID(), CPUID) // See caption below!
13     RELEASE (thread_table_lock)
14     return
15
16 procedure SCHEDULER ()
17     while shutdown = FALSE do
18         ACQUIRE (thread_table_lock)
19         for i from 0 until 7 do
20             if thread_table[i].state = RUNNABLE then
21                 thread_table[i].state ← RUNNING
22                 processor_table[CPUID].thread_id ← i
23                 EXIT_PROCESSOR_LAYER (CPUID, i)
24                 if thread_table[i].kill_or_continue = KILL then
25                     thread_table[i].state ← FREE
26                     DEALLOCATE (thread_table[i].stack)
27                     thread_table[i].kill_or_continue = CONTINUE
28             RELEASE (thread_table_lock) // Go shut down this processor
29         return
30
31 procedure ENTER_PROCESSOR_LAYER (tid, processor)
32     thread_table[tid].state ← RUNNABLE
33     thread_table[tid].topstack ← SP // save state: store yielding's thread's SP
34     SP ← processor_table[processor].topstack // dispatch: load SP of processor thread
35     return
36
37 procedure EXIT_PROCESSOR_LAYER (processor, tid) // transfers control to after line 14
38     processor_table[processor].topstack ← SP // save state: store processor thread's SP
39     SP ← thread_table[tid].topstack // dispatch: load SP of thread
40     return

```

iv) [10] Suppose the ACQUIRE and RELEASE of lines 18 and 28 were moved to be between lines 20-21 (ACQUIRE) and 21-22 (RELEASE). Would this approach work? If so, explain why; if not, provide a concrete example of a race condition that might occur.

**4) [20]** Multiple choice and true/false questions

- i) [2] The file descriptor of an open file is stored in its inode  
[ True | False ]
- ii) [2] An open file's offset for reads and writes is stored in its inode  
[ True | False ]
- iii) [2] Symbolic links can be used both for directories and files  
[ True | False ]
- iv) [2] A name resolver that uses a table allows for synonyms  
[ True | False ]
- v) [2] The SVC instruction provides address of a gate as its argument  
[ True | False ]
- vi) [5] In contrast to a hierarchical file system, consider a "flat" file system which only has one context. *Circle all correct statements:*
  - i) A flat file system cannot support access permissions
  - ii) A flat file system cannot support symbolic links
  - iii) A flat file system cannot hold two files with the same name
  - iv) A flat file system cannot support 'hard' links
  - v) Lookups do not need to be recursive in a flat file system
- vii) [5] Application A in client C opens file F in an NFS mounted directory. While A holds F open, the server suffers a short power outage and quickly reboots. *Circle all correct statements:*
  - i) The application's file descriptor becomes invalid
  - ii) The file handle for file F becomes invalid
  - iii) The NFS client stub may re-send multiple RPC calls for the same idempotent NFS operation during the reboot period
  - iv) The NFS client must re-mount the file system to recover
  - v) The NFS server sends an RPC error message to the client upon reboot



**Scratch space**

**Scratch space**