

Principles of Computer System Design
Midterm Exam
Wednesday, Oct 8th, 2014

NAME: Solution

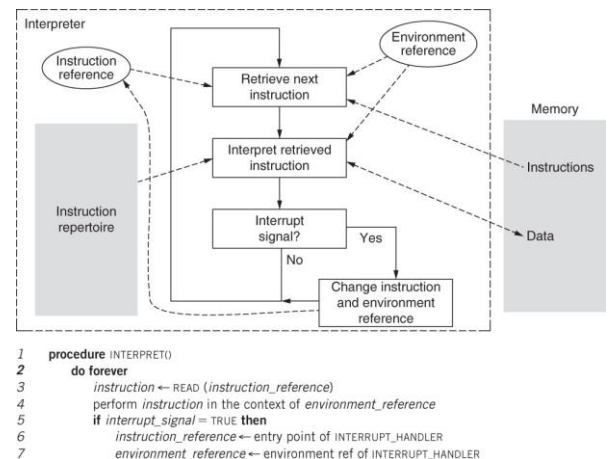
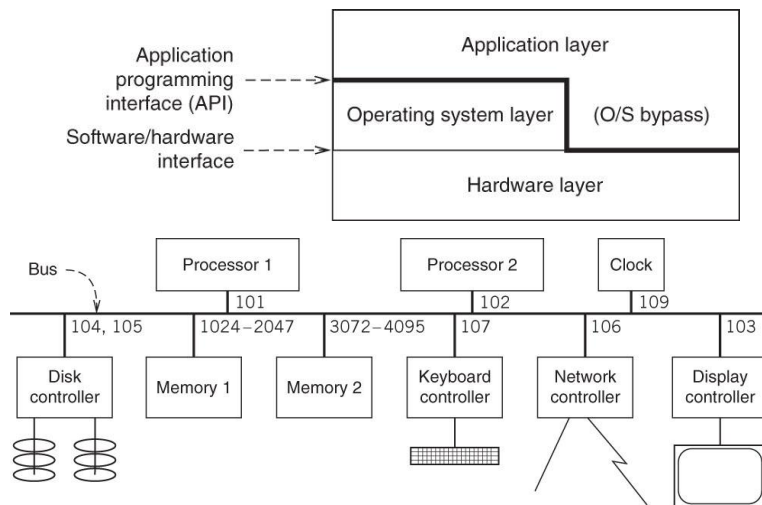
UFID:

Please read each question carefully, to avoid any confusion. This exam should have 10 pages printed double-sided (the last 2 pages are blank); before you begin, make sure your copy contains all pages. The exam is closed book, closed notes. Each question has its number of points identified in brackets.

GOOD LUCK!

QUESTION	POINTS SCORED
1 [25]	
2 [25]	
3 [30]	
4 [20]	
TOTAL	

1) [25] Naming and layers:



1.a) [10] Suppose each processor above has a single kernel/user bit, and a single memory domain register that enforces bounds (lower L, upper U) and permissions (Read, Write, eXec). Suppose processor 1's domain register has (L=1024, U=1535; perm=R, X). Suppose an application in P1 issues the following instructions; *which of these instructions result in crossing from the application to OS layer? Why?* Notation: instruction address: opcode, register, data address

3072: LOAD R1, 1030

Crosses – instruction address 3072 outside domain

1024: STORE R2, 1400

Crosses – permission for 1400 does not allow writes

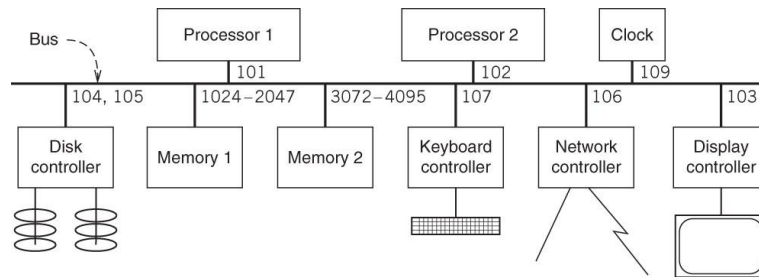
1300: LOAD R3, 1200

Does not cross; both instruction, data within bounds, permission ok

1240: SVC

Crosses – it's the behavior of this instruction

1.b) [15] Now assume that the processor supports page-based virtual memory. Suppose the page size is 1024 bytes. Given the page tables for processors P1 and P2 to be as follows (VA, PA: virtual and physical addresses, respectively).



P1's page table

VA	PA	Perm
1024	1024	RW
0	3072	RX
3072	0	R

P2's page table

VA	PA	Perm
0	1024	RW
1024	3072	RX
2048	0	R

- i) Would it be possible for a thread in P1 to send a message to a thread in P2 using a shared buffer? If so, describe where the buffer and lock need to reside in physical memory

Yes – page in PA 1024 shared read/write by both. Buffer and lock should be in this page

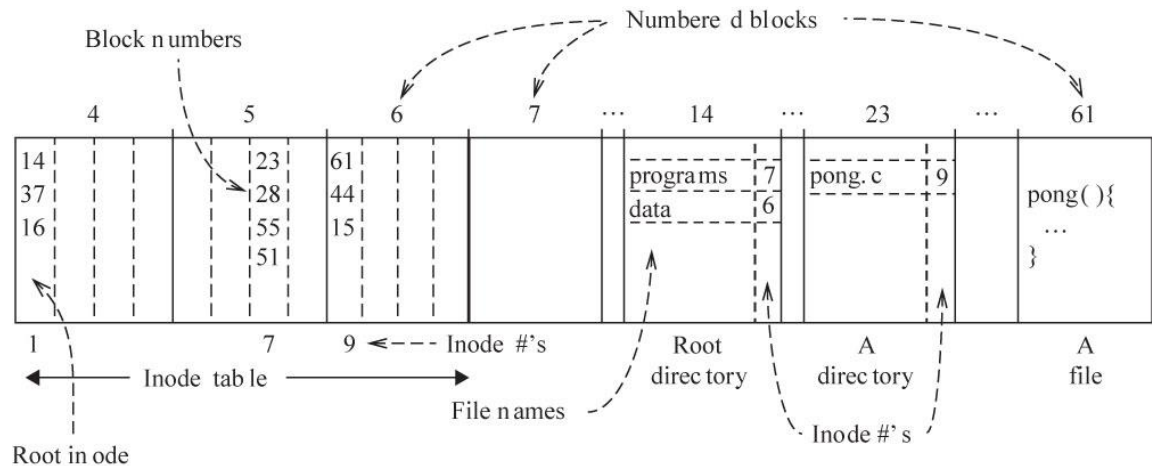
- ii) Would it be possible for threads in P1 and P2 to use *memory mapped I/O* to *read* data from the keyboard controller? What about to *send* data over the network interface?

Possible to read from keyboard controller (PA 0 mapped readable; keyboard is in address 107)

Not allowed to send to network (PA 0 mapped readable; network in address 106)

2) [25] File systems, RPC, NFS:

a) [10] Consider the figure below discussed in class.



i) Suppose an application opens and reads `/programs/pong.c`. Name *all* block numbers read from disk, and the order in which they are read.

4, 14, 5, 23, 6, 61 (44,15)

ii) Suppose you create a hard link `"/myhardlink"` that links to `/programs/pong.c`, and a soft link `"/mysoftlink"` that also links to `/programs/pong.c`. Which *i-nodes* and *blocks*, if any, would need to be modified to accommodate these two links? Explain.

myhardlink: change block 16 to add `"myhardlink | 9"`
increment reference count of inode #9

mysoftlink: create inode "n"; stores name `"/programs/pong.c"`
change block 16 to add `"mysoftlink | n"`

b) [15] Alice is using both her desktop D and laptop L to work on a file served by an NFS server S. The NFS file system is mounted on both D and L in /home. Alice opens the file "/home/alice/file.txt" in both computers. Assume there is no caching.

- i) What LOOKUP messages are sent to S when the system call to open() the file is issued by D and L? State your assumptions.

Two lookups: LOOKUP (root file handle, "alice") -> alicefh
 LOOKUP (alicefh, "file.txt") -> filefh

- ii) Are the generation counts the same in the file handles received by D and L? Briefly explain.

Yes; counts only differ if file unlinked and inode reused

- iii) At a later point in time, Alice notices that two subsequent NFS reads issued by D, for the same file handle and offset, result in a successful read followed by a stale file handle error. How could that happen?

Client L has unlinked file.txt in between the reads

3) [30] Virtualization and enforced modularity

- i) [5] What must be the permissions associated with the memory domain where the bounded buffer is stored in order to enforce modularity?

K, R, W – must be readable and writable by kernel only

- ii) [5] What privilege level is the processor on when the RSM instruction is issued to acquire the bounded buffer's lock? Why?

Kernel mode – the memory domain where the lock resides has K permission

- iii) [10] Can a race condition occur in the bounded buffer send/receive with a *single* client and *single* service? Explain

Yes (in the version that did not use locks) – e.g. writes to multiple-cell variables by sender (for instance, in) can be partially read by service

Refer to the YIELD implementation:

```

1  shared structure processor_table[7] // each processor maintains the following information
2      integer topstack                // value of stack pointer
3      byte reference stack            // preallocated stack for this processor
4      integer thread_id               // identity of thread currently running on this processor
5  shared structure thread_table[7] // each thread maintains the following information:
6      integer topstack                // value of the stack pointer
7      integer state                   // RUNNABLE, RUNNING, or FREE
8      boolean kill_or_continue       // terminate this thread? initialized to CONTINUE
9      byte reference stack            // stack for this thread

10 procedure YIELD ()
11     ACQUIRE (thread_table_lock)
12     ENTER_PROCESSOR_LAYER (GET_THREAD_ID(), CPUID) // See caption below!
13     RELEASE (thread_table_lock)
14     return
15
16 procedure SCHEDULER ()
17     while shutdown = FALSE do
18         ACQUIRE (thread_table_lock)
19         for i from 0 until 7 do
20             if thread_table[i].state = RUNNABLE then
21                 thread_table[i].state ← RUNNING
22                 processor_table[CPUID].thread_id ← i
23                 EXIT_PROCESSOR_LAYER (CPUID, i)
24                 if thread_table[i].kill_or_continue = KILL then
25                     thread_table[i].state ← FREE
26                     DEALLOCATE(thread_table[i].stack)
27                     thread_table[i].kill_or_continue = CONTINUE
28             RELEASE (thread_table_lock)
29         return // Go shut down this processor

30 procedure ENTER_PROCESSOR_LAYER (tid, processor)
31     thread_table[tid].state ← RUNNABLE
32     thread_table[tid].topstack ← SP // save state: store yielding's thread's SP
33     SP ← processor_table[processor].topstack // dispatch: load SP of processor thread
34     return

35 procedure EXIT_PROCESSOR_LAYER (processor, tid) // transfers control to after line 14
36     processor_table[processor].topstack ← SP // save state: store processor thread's SP
37     SP ← thread_table[tid].topstack // dispatch: load SP of thread
38     return

```

iv) [10] Suppose the ACQUIRE and RELEASE of lines 18 and 28 were moved to be between lines 20-21 (ACQUIRE) and 21-22 (RELEASE). Would this approach work? If so, explain why; if not, provide a concrete example of a race condition that might occur.

It does not work; for instance, two processors reach line 20 and decide that the same thread "i" is runnable; then, same thread will run in two processors

4) [20] Multiple choice and true/false questions

- i) [2] The file descriptor of an open file is stored in its inode
[True | False]
- ii) [2] An open file's offset for reads and writes is stored in its inode
[True | False]
- iii) [2] Symbolic links can be used both for directories and files
[True | False]
- iv) [2] A name resolver that uses a table allows for synonyms
[True | False]
- v) [2] The SVC instruction provides address of a gate as its argument
[True | False]
- vi) [5] In contrast to a hierarchical file system, consider a "flat" file system which only has one context. *Circle all correct statements:*
 - i) A flat file system cannot support access permissions
 - ii) A flat file system cannot support symbolic links
 - iii) A flat file system cannot hold two files with the same name
 - iv) A flat file system cannot support 'hard' links
 - v) Lookups do not need to be recursive in a flat file system
- vii) [5] Application A in client C opens file F in an NFS mounted directory. While A holds F open, the server suffers a short power outage and quickly reboots. *Circle all correct statements:*
 - i) The application's file descriptor becomes invalid
 - ii) The file handle for file F becomes invalid
 - iii) The NFS client stub may re-send multiple RPC calls for the same idempotent NFS operation during the reboot period
 - iv) The NFS client must re-mount the file system to recover
 - v) The NFS server sends an RPC error message to the client upon reboot

Scratch space

Scratch space