

EEL 5764 Computer Architecture

Sandip Ray

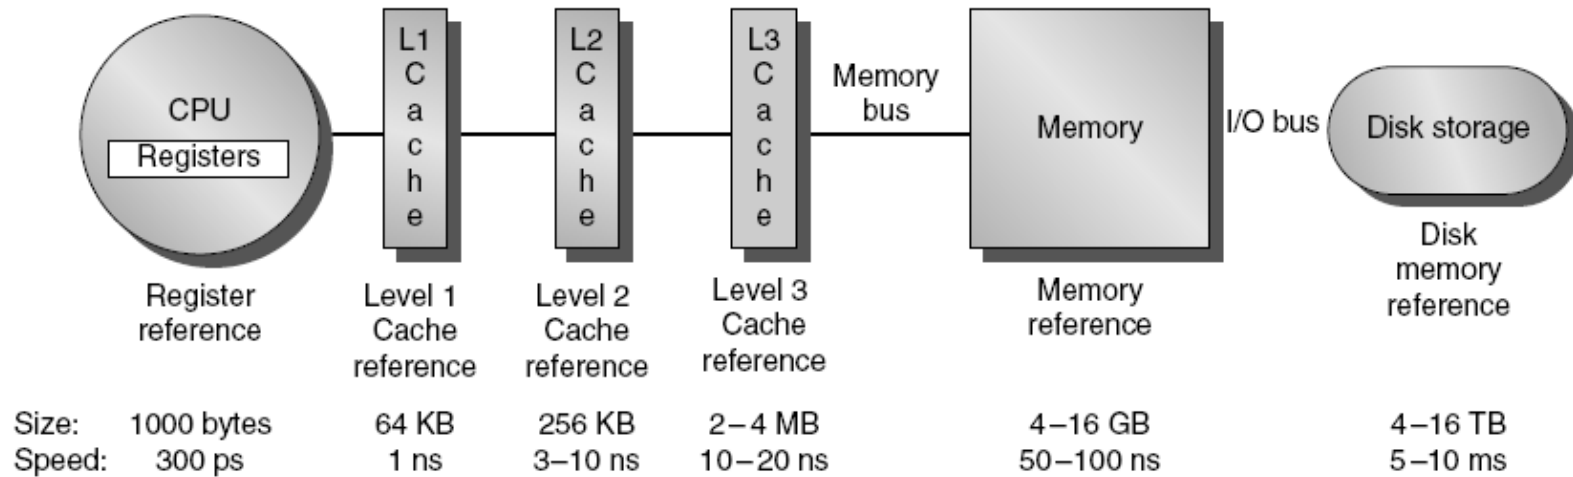
Department of Electrical and Computer Engineering
University of Florida

Lecture 6: Memory Hierarchy and Cache


Announcements

- **The time of the two in-class exams will be changed (sorry)**
 - Same day, time in the evening (likely 8:30-10:00pm)
 - Please keep yourselves available for the entire evening those days to allow for the exam
 - Please let me know (in writing, with documentation) by **midnight Friday September 6** if you can't take the exams on the requisite day.

Memory Hierarchy



(a) Memory hierarchy for server




(b) Memory hierarchy for a personal mobile device

Principle of Locality – Review

- Programs often access a small proportion of their address space at any time
- Temporal locality
 - Items accessed recently are likely to be accessed again soon
 - e.g., instructions in a loop, induction variables
- Spatial locality
 - Items near those accessed recently are likely to be accessed soon
 - E.g., sequential instruction access, array data


Cache Basics

- When a word is found in cache -> **cache hit**.
- When a word is not found in the cache, a **miss** occurs:
 - Fetch word from lower level in hierarchy, requiring a higher latency reference
 - Lower level may be another cache or the main memory
 - Also fetch the other words contained within the **block**
 - Takes advantage of spatial locality
 - Place block into cache in any location within its **set**, determined by address




Cache Basics

- Inclusive cache



- Exclusive cache




Memory Hierarchy Questions


- **Block placement** - where Can a block be placed in the upper level?
- **Block identification** – how to find a block in the upper level?
- **Block replacement** - which block should be replaced on a miss?
- **Write strategy** – how to handle writes?

Direct Mapped Cache

- Only one choice (**Block address**) MOD (**#Blocks in cache**)



- *#Blocks is a power of 2*
- *Use low-order address bits*



Tags and Valid Bits

- Multiple memory blocks mapped to a single cache location.
- How do we know which particular memory block is stored in a cache location?
 - Store **tag** as well as the data
 - Tag is the high-order bits of a block address
- What if there is no data in a location?
 - **Valid** bit: 1 = present, 0 = not present
 - Initially 0

Cache Example

- 8-blocks, 1 word/block, direct mapped
- Initial state

Index	V	Tag	Data
000	0		
001	0		
010	0		
011	0		
100	0		
101	0		
110	0		
111	0		

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
22	10 110	Miss	110

Index	V	Tag	Data
000	0		
001	0		
010	0		
011	0		
100	0		
101	0		
110	1	10	Mem[10110]
111	0		

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
26	11 010	Miss	010

Index	V	Tag	Data
000	0		
001	0		
010	1	11	Mem[11010]
011	0		
100	0		
101	0		
110	1	10	Mem[10110]
111	0		

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
16	10 000	Miss	000
3	00 011	Miss	011
16	10 000	Hit	000

Index	V	Tag	Data
000	1	10	Mem[10000]
001	0		
010	1	11	Mem[11010]
011	1	00	Mem[00011]
100	0		
101	0		
110	1	10	Mem[10110]
111	0		

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
18	10 010	Miss	010

Index	V	Tag	Data
000	1	10	Mem[10000]
001	0		
010	1	10	Mem[10010]
011	1	00	Mem[00011]
100	0		
101	0		
110	1	10	Mem[10110]
111	0		

Cache Misses

- On cache *hit*, CPU proceeds normally
- On cache *miss*
 - Stall the CPU
 - Fetch a block from next level of mem. hierarchy
 - Instruction cache miss
 - Restart instruction fetch
 - Data cache miss
 - Complete data access
- Note that speculative and multithreaded processors may execute other instructions during a miss
 - Reduces performance impact of misses