

EEL 5764 Computer Architecture

Sandip Ray

Department of Electrical and Computer Engineering
University of Florida

Lecture 14-15:

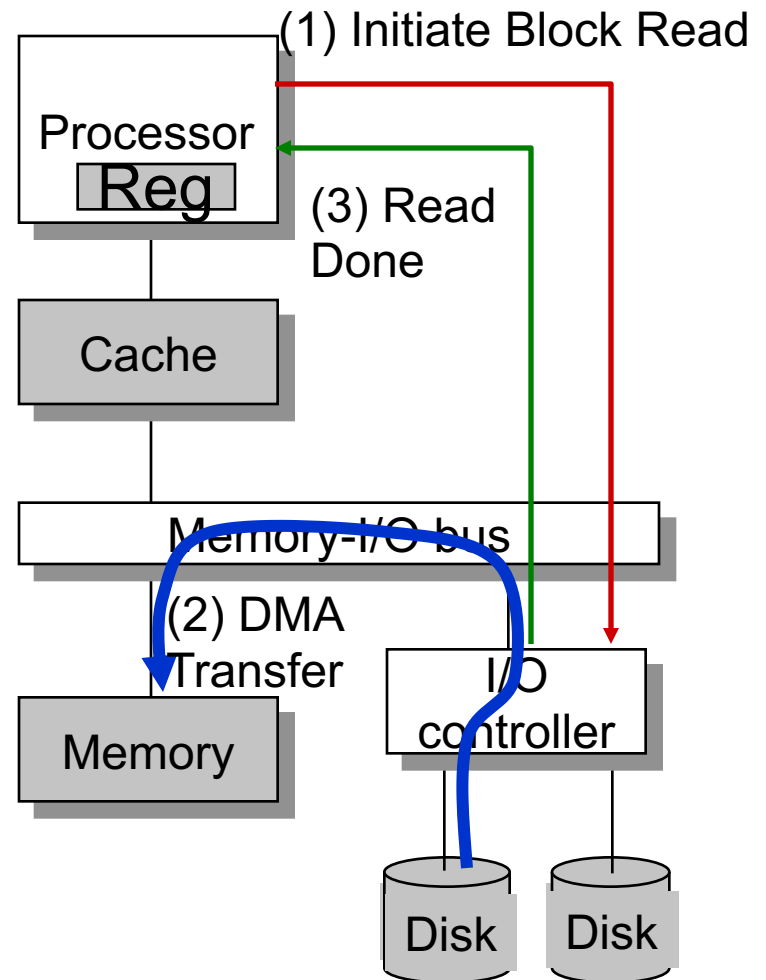
- Virtual Memories
- Instruction Set Architecture

Announcement

- Please use Canvas and preferably the discussion section for your course-related questions
- For any message you send through Canvas please copy the **entire teaching staff (instructor and TA)**
- For HW-related questions and clarifications, please contact TA first

Servicing Page Faults

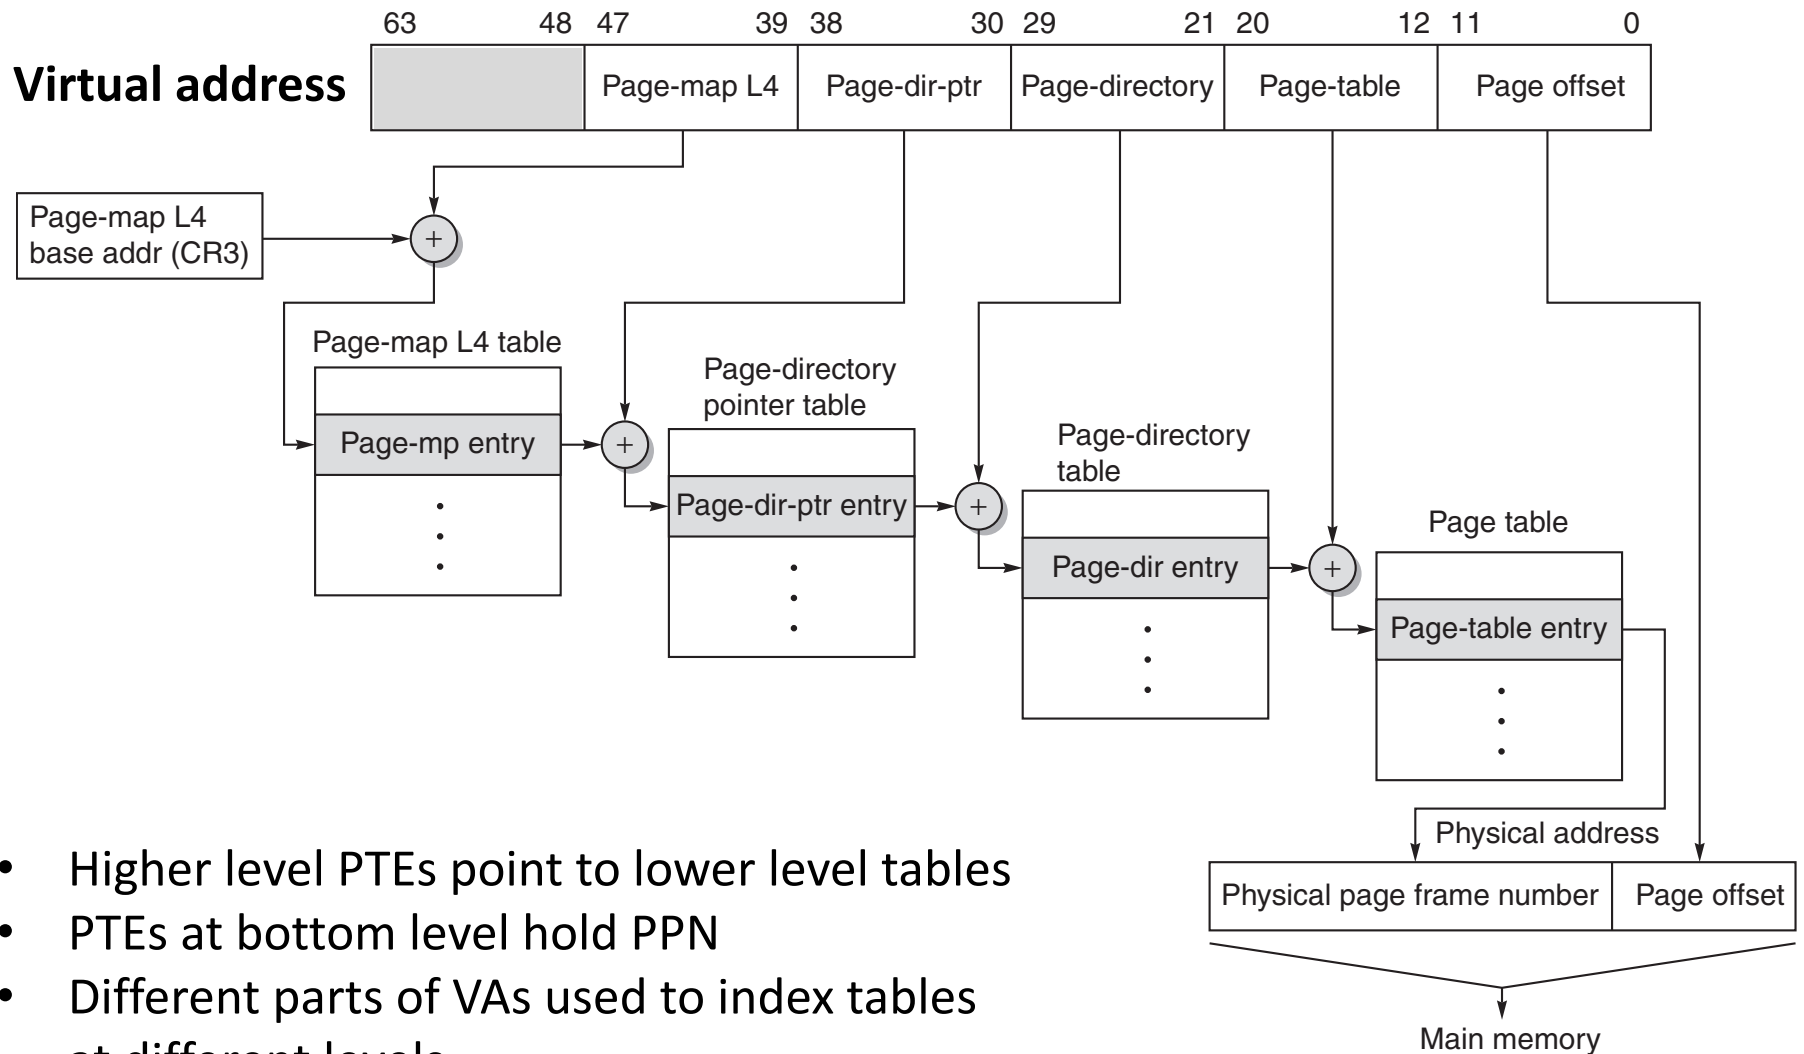
- Processor signals controller
 - Read block of length P starting at disk address X and store starting at memory address Y
- Read occurs
 - Direct Memory Access (DMA)
 - Under control of I/O controller
- I/O controller signals completion
 - Interrupt processor
 - OS resumes suspended process



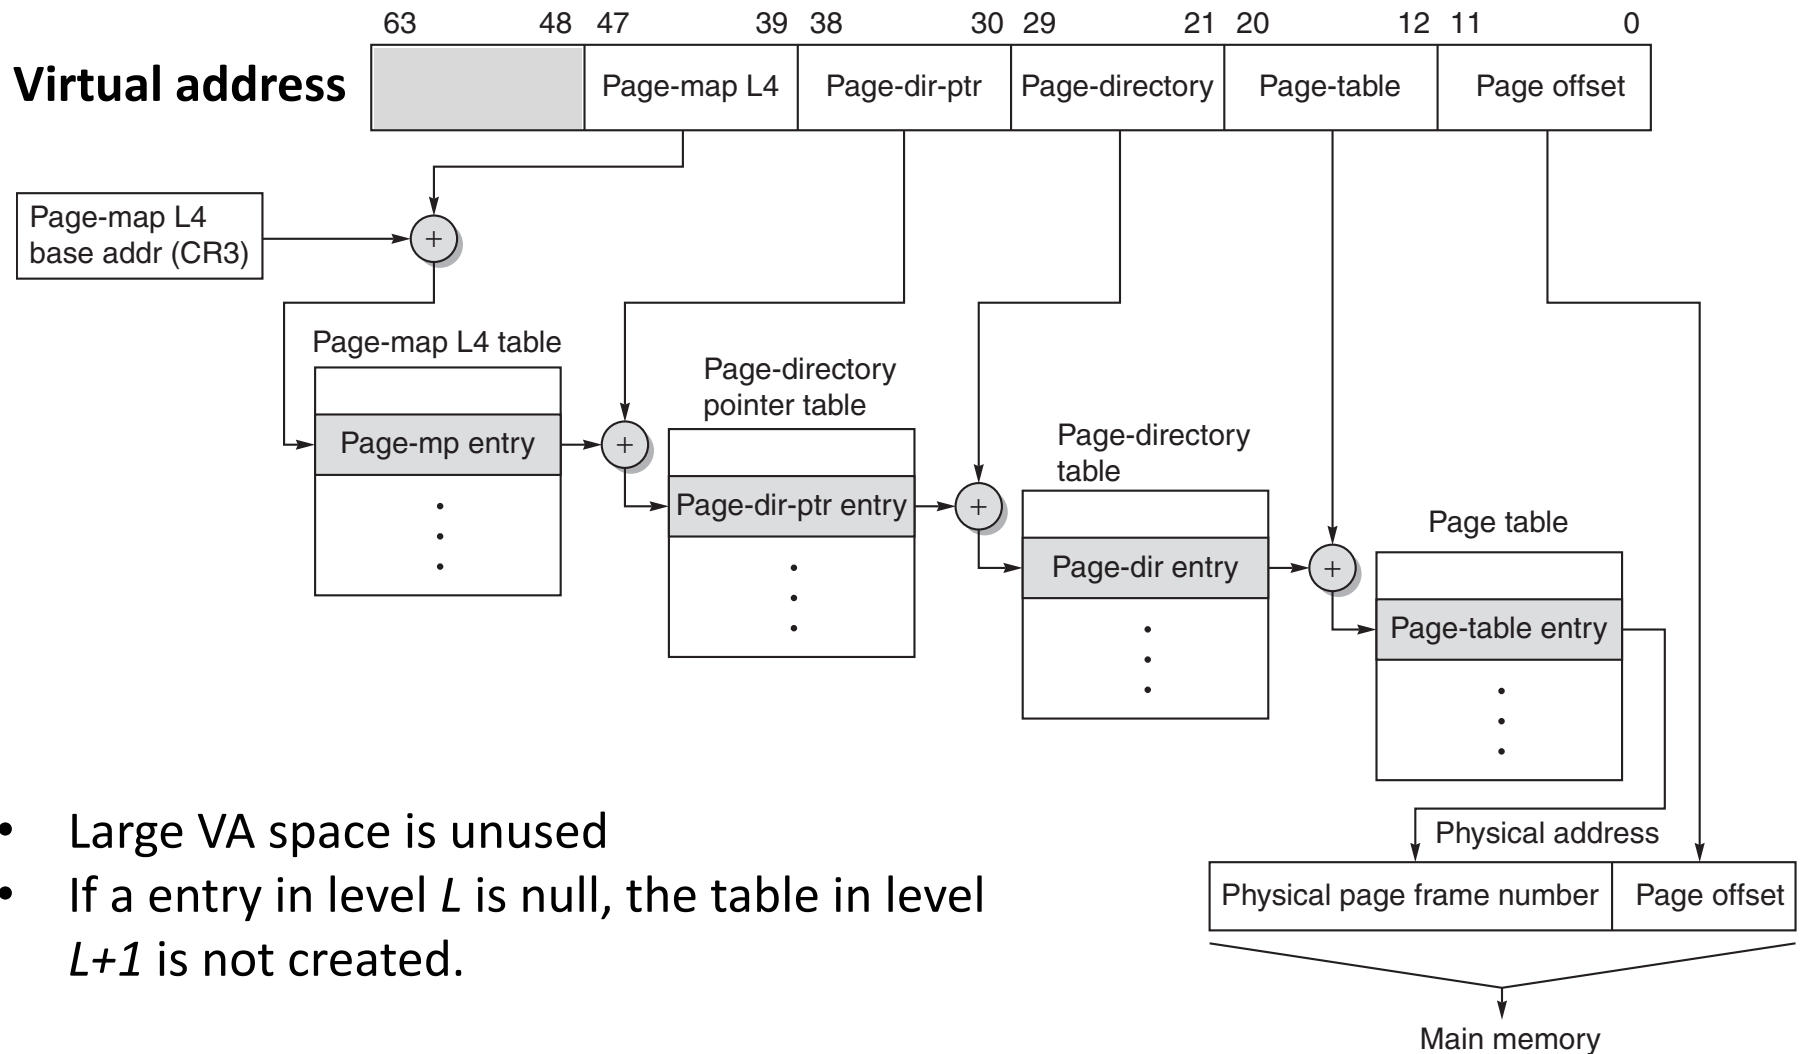
Page Replacement

- When there are no available free pages to handle a fault we must find a page to replace.
- This is determined by the *page replacement algorithm*.
- The goal of the replacement algorithm is to reduce the fault rate by selecting the best victim page to remove.
- LRU is the most popular replacement policy.
 - Each page is associated with a *use/reference* bit.
 - It is set whenever a page is accessed.
 - OS periodically clears these bits.

Multi-Level Page Table



Multi-Level Page Table



- Large VA space is unused
- If a entry in level L is null, the table in level $L+1$ is not created.

Selecting Page Size

- The larger the page size, the smaller page table
 - A page table can occupy a lot of space
- Larger page size -> larger L1 cache
- Larger page size -> less TLB misses, faster xlation
- More efficient to xfer larger pages from 2nd storage
- Large page size
 - Wasted storage
 - Wasted IO bandwidth
 - Slower process start up

Virtual Memory – Protection

- Virtual memory and multiprogramming
 - Multiple processes sharing processor and physical memory
- Protection via virtual memory
 - Keeps processes in their own memory space
- Role of architecture:
 - Provide user mode and supervisor mode
 - Protect certain aspects of CPU state: PC, register, etc
 - Provide mechanisms for switching between user mode and supervisor mode
 - Provide mechanisms to limit memory accesses
 - Provide TLB to translate addresses

Page Table Entries

- Address – physical page number
- Valid/present bit
- Modified/dirty bit
- Reference bit
 - For LRU
- Protection bits – Access right field defining allowable accesses
 - E.g., read only, read-write, execute only
 - Typically support multiple protection modes (kernel vs user)

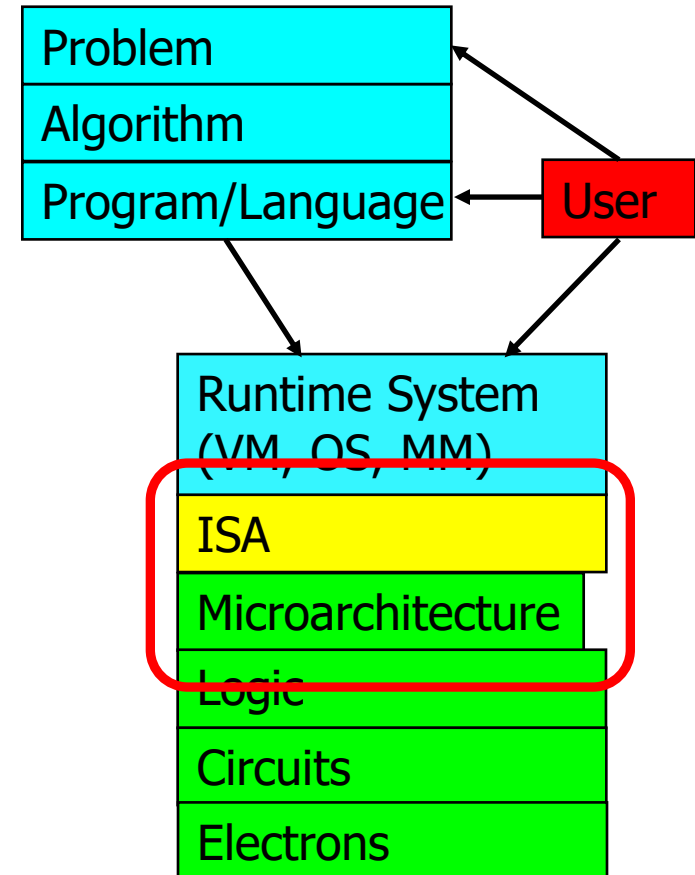
Summary

- OS virtualizes memory and IO devices
 - Each process has an illusion of private CPU and memory
- Virtual memory
 - Arbitrarily large memory, isolation/protection, inter-process communication
 - Reduce page table size
 - Translation buffers – cache for page table
 - Manage TLB misses and page faults

Instruction Set Architecture

Abstractions

- Abstraction helps us deal with complexity
 - Hide lower-level detail
- Instruction set architecture (ISA)
 - The hardware/software interface
 - Defines storage, operations, etc
- Implementation
 - The details underlying the interface
 - An ISA can have multiple implementations



Levels of Program Code

- High-level language
 - Level of abstraction closer to problem domain
 - Provides for productivity and portability
- Assembly language
 - Textual representation of instructions
- Hardware representation
 - Machine code - Binary digits (bits)
 - Encoded instructions and data
- Compiler
 - Translate HL prgm to assembly
- Assembler
 - Translate assembly to machine code

High-level
language
program
(in C)

```
swap(int v[], int k)
{int temp;
  temp = v[k];
  v[k] = v[k+1];
  v[k+1] = temp;
}
```

Compiler

Assembly
language
program
(for MIPS)

```
swap:
    muli $2, $5, 4
    add  $2, $4, $2
    lw   $15, 0($2)
    lw   $16, 4($2)
    sw   $16, 0($2)
    sw   $15, 4($2)
    jr   $31
```


Assembler

Binary machine
language
program
(for MIPS)

```
000000001010000100000000000011000
000000000000110000001100000100001
100011000110001000000000000000000
100011001111001000000000000000100
101011001111001000000000000000000
101011001111001000000000000000000
101011000110001000000000000000100
00000011111000000000000000001000
```

Program Execution Model

- A computer is just a FSM
 - States stored in registers, memory, PC, etc
 - States changed by instruction execution
- An instruction is executed in
 - **Fetch** an instruction into CPU from memory
 - **Decode** it to generate control signals
 - **Execute** it (add, mult, etc)
 - **Write back** output to reg or memory
- *Programs and data* coexist in memory



What Makes a Good ISA?

- **Programmability**

- Who does assembly programming these days?

- **Performance/Implementability**

- Easy to design high-performance implementations?

- Easy to design low-power implementations?

- Easy to design low-cost implementations?

- **Compatibility**

- Easy to maintain as languages, programs evolve

- x86 (IA32) generations: 8086, 286, 386, 486, Pentium, Pentium-II, Pentium-III, Pentium4, Core2, Core i7, ...

Performance

- **Execution time = $IC * CPI * \text{cycle time}$**
- **IC:** instructions executed to finish program
 - Determined by program, compiler, ISA
- **CPI:** number of cycles needed for each instruction
 - Determined by program, compiler, ISA, u-architecture
- **Cycle time:** inverse of clock frequency
 - Determined by u-architecture & technology
- **Ideally optimize all three**
 - Their optimizations often against each other
 - Compiler plays a significant role

Instruction Granularity

- **CISC** (Complex Instruction Set Computing) **ISAs**
 - Big heavyweight instructions (lots of work per instruction)
 - + Low “insns/program”
 - Higher “cycles/insn” and “seconds/cycle”
 - We have the technology to get around this problem
- **RISC** (Reduced Instruction Set Computer) **ISAs**
 - Minimalist approach to an ISA: simple insns only
 - + Low “cycles/insn” and “seconds/cycle”
 - Higher “insn/program”, but hopefully not as much
 - Rely on compiler optimizations

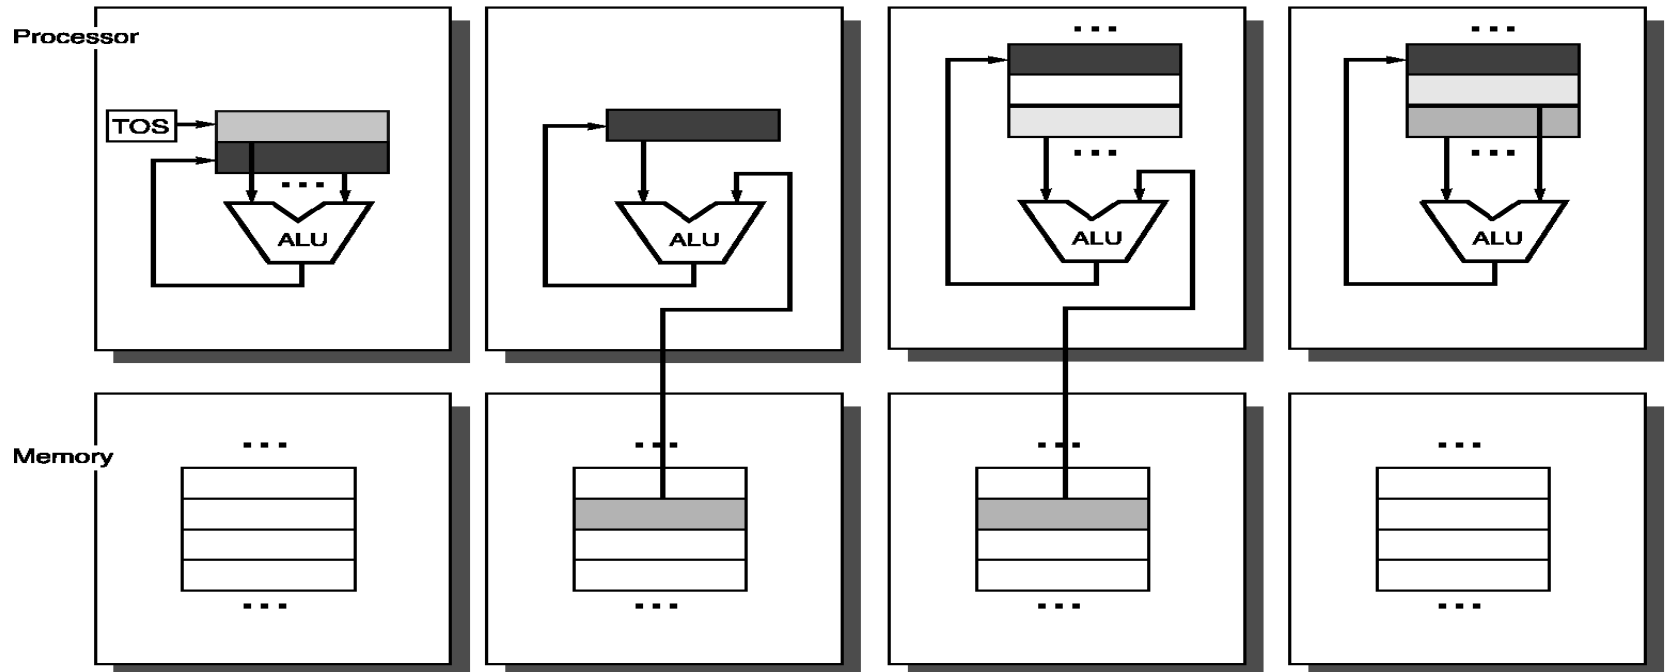
Classifying Architectures

- One important classification scheme is by the type of addressing modes supported.
- **Stack architecture**: Operands implicitly on top of a stack.
- **Accumulator** architecture: One operand is implicitly an accumulator (a special register).
- **General-purpose register** arch.: Operands may be any of a large (typically 10s-100s) # of registers.
 - Register-memory architectures: One op may be memory.
 - Load-store architectures: All ops are registers, except in special load and store instructions.

Illustrating Architecture Types

Assembly for $C := A + B$:

Stack	Accumulator	Register (register-memory)	Register (load-store)
Push A	Load A	Load R1, A	Load R1, A
Push B	Add B	Add R1, B	Load R2, B
Add	Store C	Store C, R1	Add R3, R1, R2
Pop C			Store C, R3



Number of Registers

- Registers have advantages
 - faster than memory, good for compiler optimization, hold variables
- Have as many as possible?
 - **No**
- One reason registers are faster: there are
 - **fewer of them** – small is fast (hardware truism)
- Another: they are **directly addressed**
 - More registers, means more bits per register in instruction
 - Thus, fewer registers per instruction or larger instructions
- More registers means **more saving/restoring**
 - Across function calls, traps, and context switches
- Trend toward more registers:
 - 8 (x86) → 16 (x86-64), 16 (ARM v7) → 32 (ARM v8)

Number of Operands

- A further classification is by the **max. number of operands**, and **# that can be memory**: e.g.,
 - 2-operand (e.g. $a += b$)
 - src/dest(reg), src(reg)
 - src/dest(reg), src(mem) IBM 360, x86, 68k
 - src/dest(mem), src(mem) VAX
 - 3-operand (e.g. $a = b + c$)
 - dest(reg), src1(reg), src2(reg) MIPS, PPC, SPARC, &c.
 - dest(reg), src1(reg), src2(mem)
 - dest(mem), src1(mem), src2(mem) VAX


Memory Addressing

- Byte Addressing
 - Each byte has a unique address
- Addressing units
 - Half-word: 16-bit (or 2 bytes)
 - Word: 32-bit (or 4 bytes)
 - Double word : 64-bit (or 8 bytes)
 - Quad word: 128-bit (or 16 bytes)
- Two issues
 - **Alignment** specifies whether there are any boundaries for word addressing
 - **Byte order** (Big Endian vs. Little Endian)
 - specifies how multiple bytes within a word are mapped to memory addresses

Memory Addressing

- Alignment


→ Must half word, words, double words begin mod 2, mod 4, mod 8 boundaries



Aligned if $A \bmod s = 0$


Memory Addressing

- Alignment
 - Or there no alignment restrictions



Memory Addressing

- Non-aligned memory references may cause multiple memory accesses



- Consider a system in which memory reads return 4 bytes and a reference to a word spans a 4-byte boundary: two memory accesses are required
- Complicates memory and cache controller design
- Assemblers typically force alignment for efficiency