

EEL 5764 Computer Architecture

Fall 2019

Lecture 22-23: ILP and Pipeline Scheduling

Sandip Ray

Department of Electrical and Computer Engineering
University of Florida

Announcements

- **Reminder: Mid-term 1 on October 22 8:30pm**
 - Room **NEB 0202**
- EDGE Students must be registered and have appointment for the test
 - Have Proctor U account (if you don't have already)
 - Schedule appointment with Proctor U to take the exam
 - Exam title is "EEL5764- Computer Architecture – Exam 1 Fall 2019"
- No Makeup test other than for documented emergencies recognized by University of Florida
- You are allowed
 - 1 US-letter sized crib sheet of hand-written notes (no typing).
You can use both sides
 - Scientific calculators
 - **Nothing else is allowed during test**

Dependences

- Parallelism with **basic block** is limited
 - Typical size of basic block = 3-6 instructions
 - Must optimize across branches
- Increase ILP – Loop-Level Parallelism
 - Unroll loop statically or dynamically
 - Use SIMD (vector processors and GPUs)
- Challenges: Data dependency
 - Instruction j is data dependent on instruction i if
 - Instruction i produces a result that may be used by instruction j**
 - Instruction j is data dependent on instruction k and instruction k is data dependent on instruction i**
- Dependent instructions cannot be executed simultaneously

Data Dependence

- Dependencies are a property of programs
- Pipeline organization determines if dependence is detected and if it causes a stall
- **Data dependence** conveys:
 - Possibility of a hazard
 - Order in which results must be calculated
 - Upper bound on exploitable instruction level parallelism
- Dependencies that flow through memory locations are difficult to detect
 - SD R1, 45(R5)
 - LD R2, 90(R7)

Data Hazards

- A **hazard** exists between two dependent instructions that are close in a program such that overlapping their executions would change the order of accesses to the operand involved in the dependence.
- **Data Hazards**
 - Read after write (RAW)
 - Write after write (WAW)
 - Write after read (WAR)

Control Dependences

- **Control Dependence** determines ordering of instruction i with respect to a branch instruction
 - Instruction control dependent on a branch cannot be moved before the branch so that its execution is no longer controlled by the branch
 - An instruction not control dependent on a branch cannot be moved after the branch so that its execution is controlled by the branch

Control Dependence – Examples

- Example 1:

```
DADDU    R1,R2,R3
BEQZ     R4,L
DSUBU    R1,R1,R6
L:  ...
OR       R7,R1,R8
```

- OR instruction dependent on DADDU and DSUBU
→ Cannot move DSUBU before the branch

- Example 2:

```
DADDU    R1,R2,R3
BEQZ     R12,skip
DSUBU    R4,R5,R6
DADDU    R5,R4,R9
skip:
OR       R7,R8,R9
```

- Reorder instruction if it does not change program semantics
- Assume R4 isn't used after skip
→ Possible to move DSUBU before the branch

Pipeline Scheduling

- Separate dependent instruction from the source instruction by the **pipeline latency** of the source instruction
 - Pipeline latency – cycles to separate two dependent instructions to avoid hazards
- By inserting NOP (bubbles) – pipeline stalls
- By inserting independent instructions -
Scheduling
 - Statically
 - Dynamically
- Scheduling relies on
 - ILP in programs
 - latencies of functional units

Static Scheduling

- Transform and re-arrange the code to reduce pipeline latency

- **Example:**

```
for (i=999; i>=0; i=i-1)
    x[i] = x[i] + s;
```

```
Loop:  L.D      F0, 0(R1)
        ADD.D   F4, F0, F2
        S.D     F4, 0(R1)
        DADDUI  R1, R1, #-8
        BNE     R1, R2, Loop
```

Assume branch delay = 1 cycle

Instruction producing result	Instruction using result	Latency in clock cycles
FP ALU op	Another FP ALU op	3
FP ALU op	Store double	2
Load double	FP ALU op	1
Load double	Store double	0

Static Scheduling

Loop: L.D F0, 0(R1)
stall
ADD.D F4, F0, F2
stall
stall
S.D F4, 0(R1)
DADDUI R1, R1, #-8
stall
BNE R1, R2, Loop

Instruction producing result	Instruction using result	Latency in clock cycles
FP ALU op	Another FP ALU op	3
FP ALU op	Store double	2
Load double	FP ALU op	1
Load double	Store double	0

Static Scheduling

Scheduled code:

```
Loop:  L.D      F0, 0(R1)
        DADDUI  R1, R1, #-8
        ADD.D   F4, F0, F2
        stall
        stall
        S.D     F4, 8(R1)
        BNE     R1, R2, Loop
```

- ILP in basic blocks is limited
- Loop unrolling to increase ILP

Instruction producing result	Instruction using result	Latency in clock cycles
FP ALU op	Another FP ALU op	3
FP ALU op	Store double	2
Load double	FP ALU op	1
Load double	Store double	0

Loop Unrolling

- Unroll by a factor of 4 (assume # elements is divisible by 4)
 - Eliminate unnecessary instructions

```
Loop:   L.D      F0, 0(R1)
        ADD.D    F4, F0, F2
        S.D      F4, 0(R1)    ;drop DADDUI & BNE
        L.D      F6, -8(R1)
        ADD.D    F8, F6, F2
        S.D      F8, -8(R1)   ;drop DADDUI & BNE
        L.D      F10, -16(R1)
        ADD.D    F12, F10, F2
        S.D      F12, -16(R1) ;drop DADDUI & BNE
        L.D      F14, -24(R1)
        ADD.D    F16, F14, F2
        S.D      F16, -24(R1)
        DADDUI   R1, R1, #-32
        BNE      R1, R2, Loop
```

Run in 27 cycles
without scheduling

Number of live registers
increase vs. original loop
→ Avoid dependence
→ Increase ILP

Loop Unrolling/Pipeline Scheduling

- Pipeline schedule the unrolled loop:

Loop:

L.D	F0,0(R1)
L.D	F6,-8(R1)
L.D	F10,-16(R1)
L.D	F14,-24(R1)
ADD.D	F4,F0,F2
ADD.D	F8,F6,F2
ADD.D	F12,F10,F2
ADD.D	F16,F14,F2
S.D	F4,0(R1)
S.D	F8,-8(R1)
DADDUI	R1, R1, #-32
S.D	F12, -16(R1)
S.D	F16, -24(R1)
BNE	R1, R2, Loop

⇒ Run in 14 cycles after scheduling

⇒ 3.5 cycles per iteration

Dealing with Unknown Loop Bound

- Number of iterations = n ,
- Goal: make k copies of the loop body
- Solution: generate pair of loops:
 - First executes $n \bmod k$ times of original loop body
 - Second executes n / k times of unrolled loop body by k times.

Limitations of Loop unrolling

- Loop iteration dependence
- Code size increase
- Register shortfall

Why Dynamic Scheduling

- In-Order pipeline
 - Issue, execution, and completion of instructions follow program order
- Pipeline stalls if there is data dependence
- But, independent instructions are stalled too.

DIV.D	F0, F2, F4
ADD.D	F10 F0, F8
SUB.D	F12, F8, F14

- Static scheduling may not identify some dependences

Dynamic Scheduling

- Hardware rearrange order of instructions
 - To reduce stalls while maintaining data flow and exception
- Advantages:
 - Compiler doesn't need to have knowledge of microarchitecture
 - Handles cases where dependencies are unknown at compile time
 - Tolerate unpredictable delays in M access, FP operations, etc.
- Disadvantage:
 - Substantial increase in hardware complexity
 - Complicates exceptions

Dynamic Scheduling

- ID stage is divided into
 - **Issue** – decode instruction, and check for S-hazards
 - **Read Operands** – wait until no D-hazards, then read operands
- Instructions pass **Issue** in-order
- They are stalled or bypass each other in **Read Operands**
 - Enter **Execute** out-of-order

Dynamic Scheduling

- Dynamic scheduling implies:
 - In-order issue
 - Out-of-order execution
 - Out-of-order completion
- Creates the possibility for WAR and WAW hazards

DIV.D	F0, F2, F4
ADD.D	F6, F0, F8
SUB.D	F8, F10, F14
MUL.D	F6, F10, F8

- WAR/WAW hazards can be avoided by **register renaming**

Dynamic Scheduling

- Tomasulo's Approach
 - Tracks when operands are available
 - To eliminate RAW hazards
 - Introduces **register renaming** in hardware
 - To eliminate WAW and WAR hazards

Register Renaming

- Example:

DIV.D F0,F2,F4

ADD.D **F6**,F0,F8


S.D F6,0(R1)

SUB.D F8,F10,F14

MUL.D **F6**,F10,F8

antidependence via F8

Output dependence via F6



Register Renaming

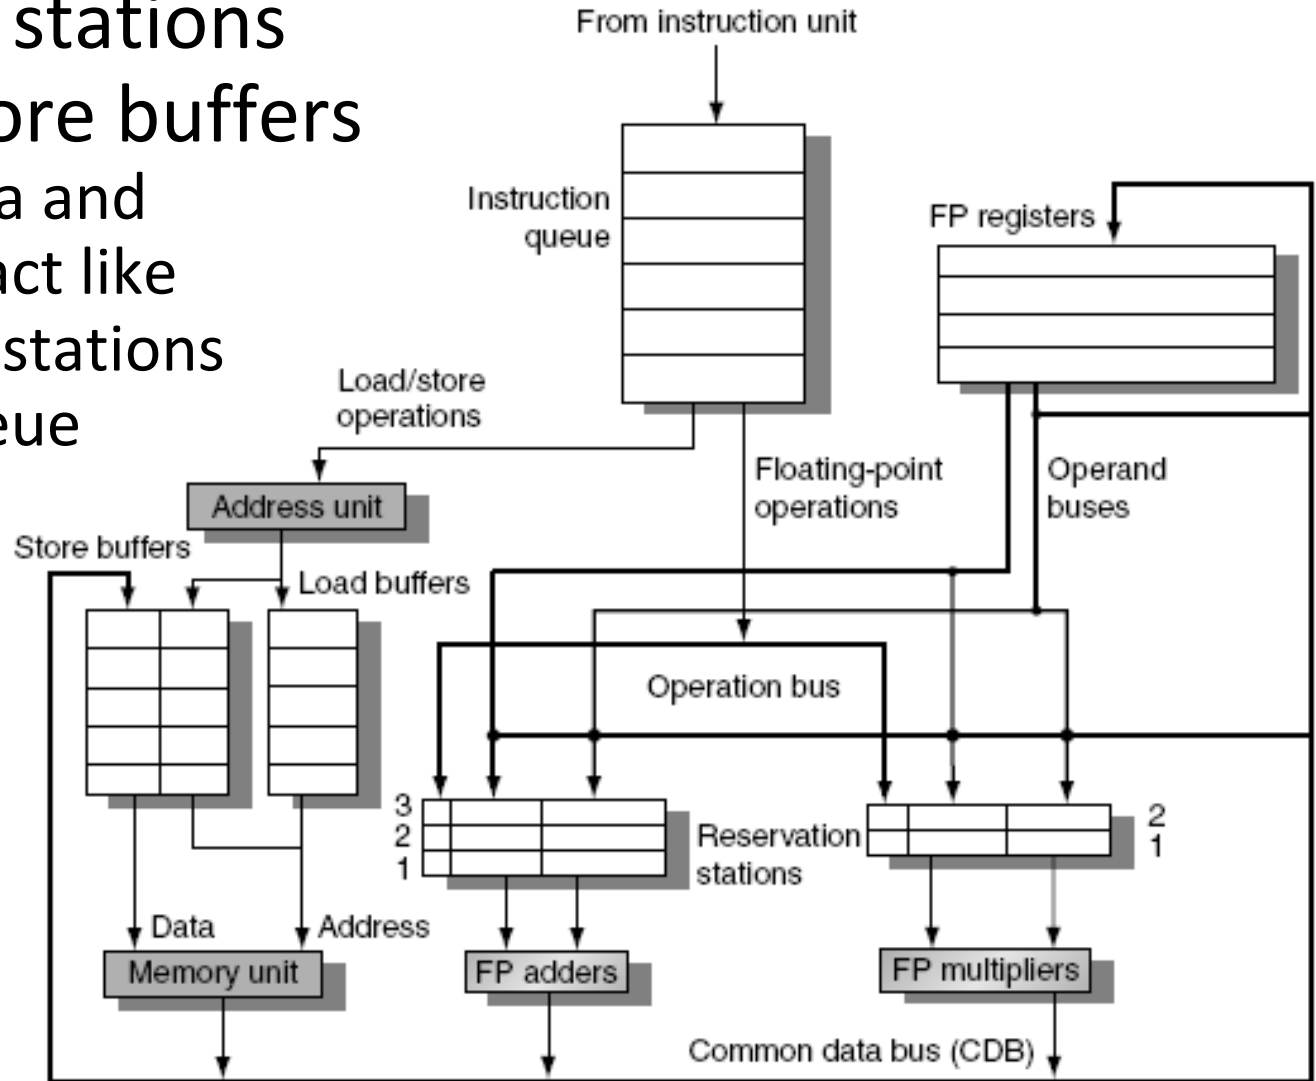
- Example: assume two additional registers **S** and **T**

DIV.D	F0,F2,F4	
ADD.D	S , F0, F8	; F6 -> S
S.D	S , 0(R1)	; F6 -> S
SUB.D	T , F10, F14	; F8 -> T
MUL.D	F6, F10, T	; F8 -> T

- Now only RAW hazards remain, which must be strictly ordered

Tomasulo's Algorithm – Structure

- Reservation stations
- Load and store buffers
 - Contain data and addresses, act like reservation stations
- Instruction queue
- CDB

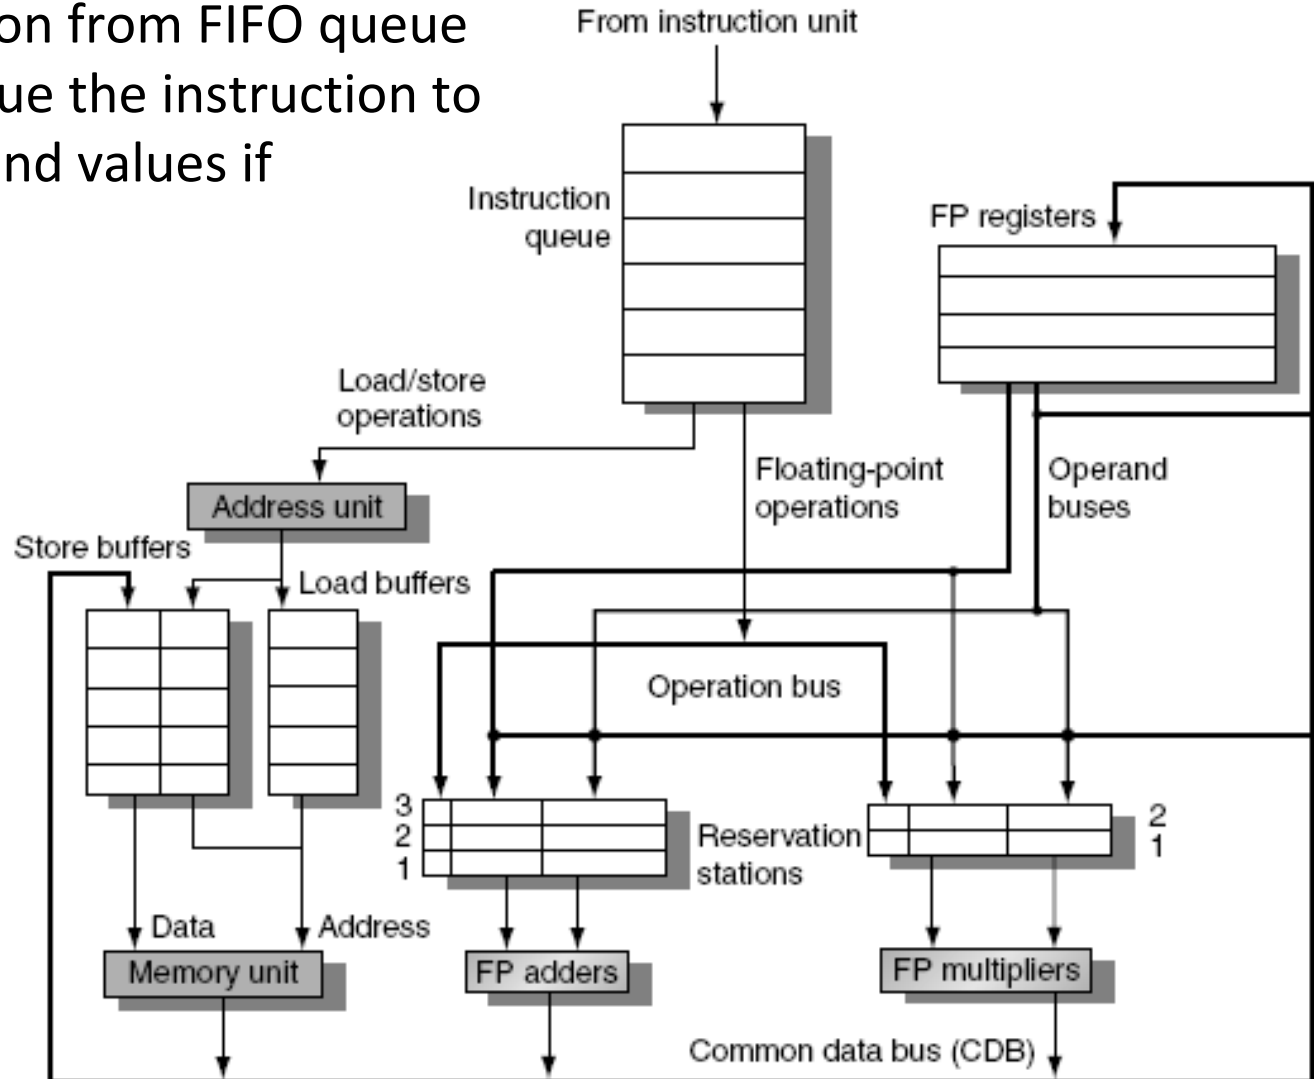


Register Renaming By Reservation Stations

- Reservation stations (RS) buffer
 - The instruction opcode
 - Buffered operand values (when available)
 - Reservation station number of instruction providing the operand values
- RS fetches and buffers an operand as soon as it becomes available (not necessarily involving register file)
- Pending instructions designate the RS to which they will send their output
- Result values broadcast on a result bus, common data bus (CDB)
- Only the last output updates the register file
- As instructions are issued, the register specifiers are renamed with the reservation station
- May be more reservation stations than registers

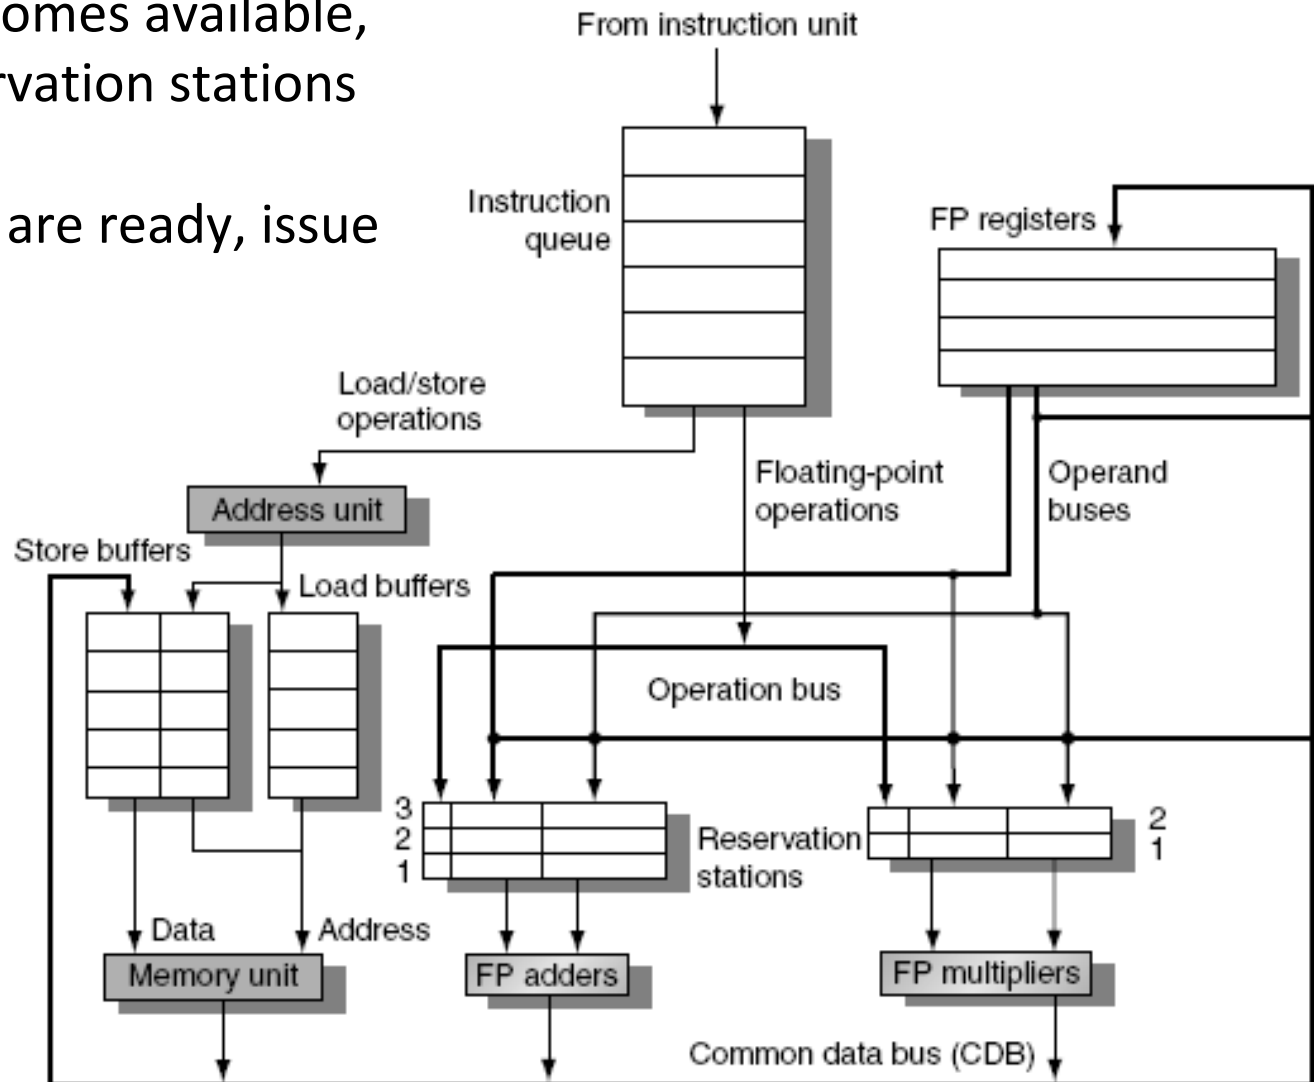
Tomasulo's Algorithm – Issue

- Get next instruction from FIFO queue
 - If available RS, issue the instruction to the RS with operand values if available
- If operand values not available, stall the instruction
- Renaming in this step



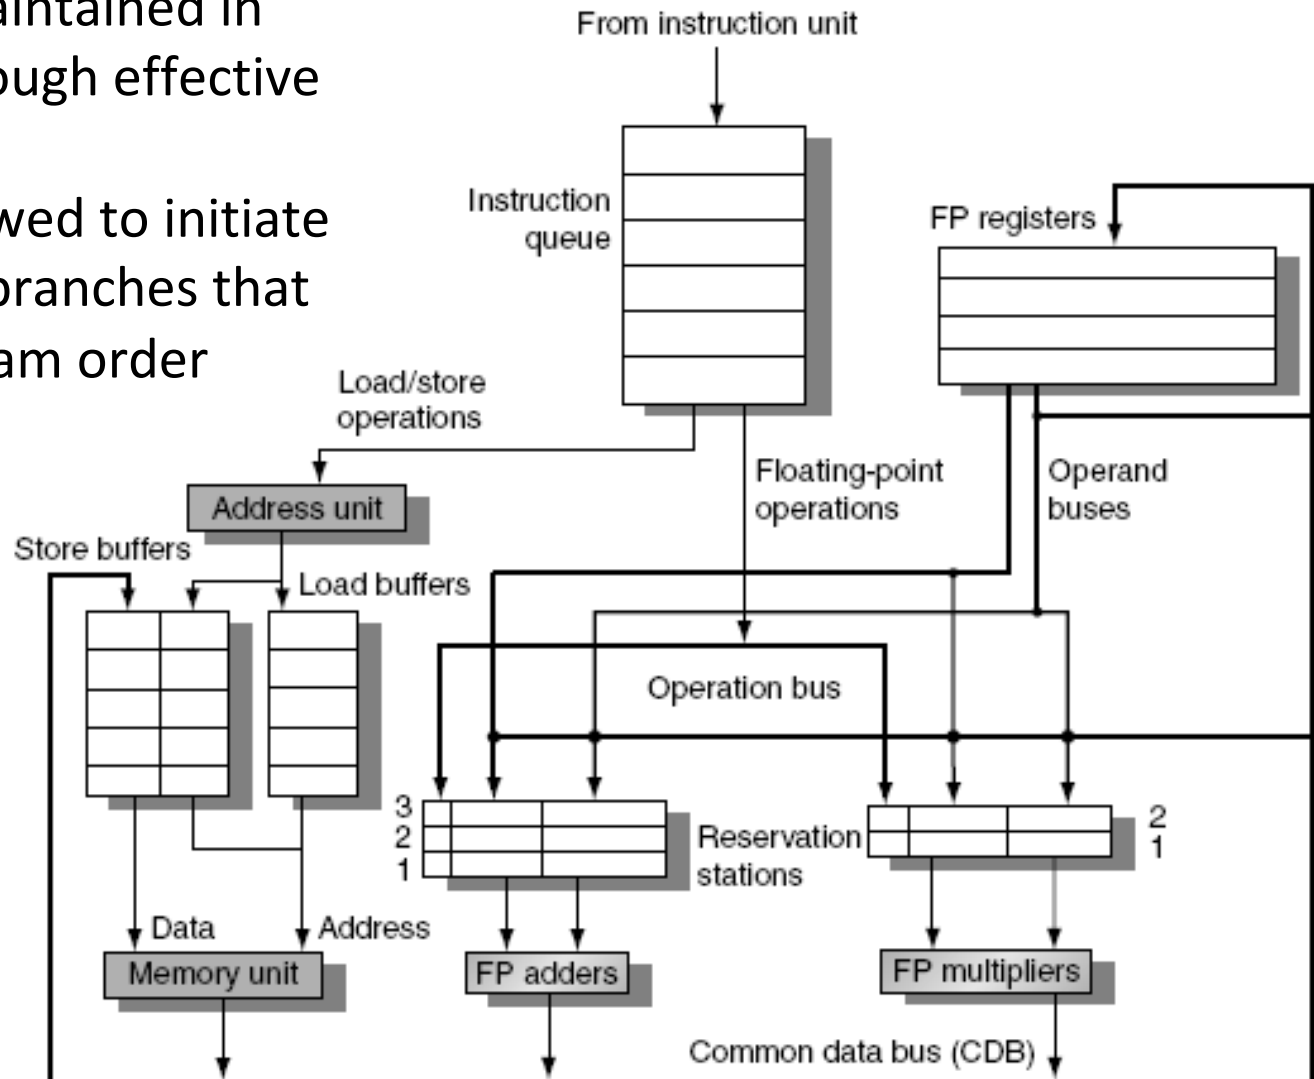
Tomasulo's Algorithm – Execute

- When operand becomes available, store it in any reservation stations waiting for it
- When all operands are ready, issue the instruction



Tomasulo's Algorithm – Execute

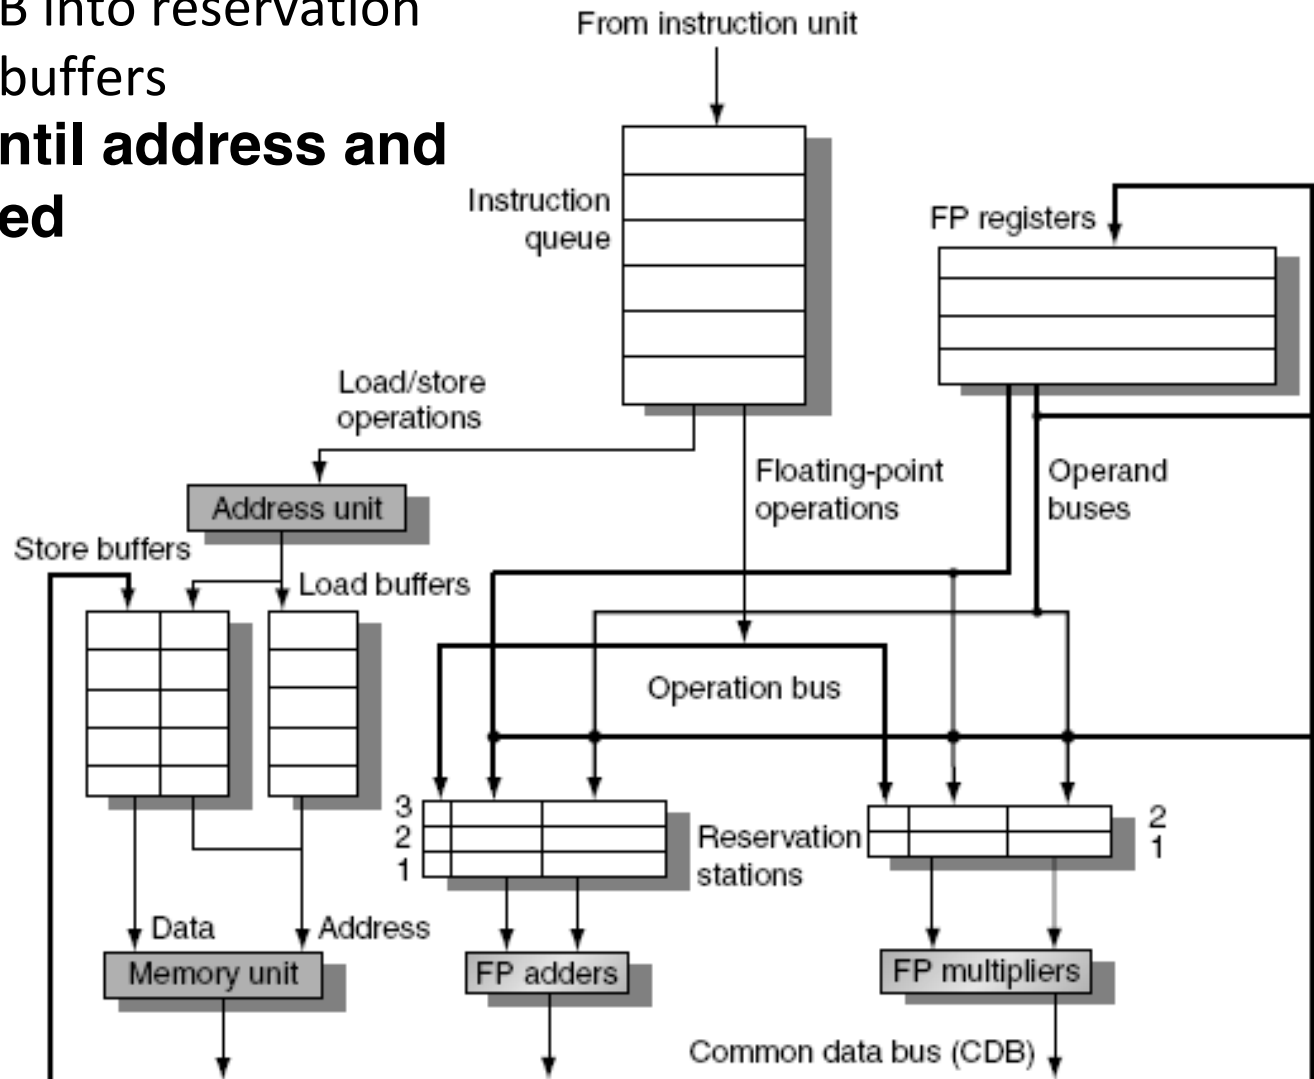
- Loads and store maintained in program order through effective address
- No instruction allowed to initiate execution until all branches that proceed it in program order have completed



Tomasulo's Algorithm – Write Result

- Write result on CDB into reservation stations and store buffers

Stores must wait until address and value are received



Tomasulo's Algorithm – Some Notes

- Each RS, register, store buffer is attached with a tag
 - Tags refer to FUs or buffers that produce operands
- FUs or memory put results on CDB
- RS, registers, store buffers monitor the CDB,
 - Get the data where its tag matches tag on CDB
- Forwarding through broadcast on CDB

Example

Instruction		Instruction status		
		Issue	Execute	Write Result
L.D	F6,32(R2)	✓	✓	✓
L.D	F2,44(R3)	✓	✓	
MUL.D	F0,F2,F4	✓		
SUB.D	F8,F2,F6	✓		
DIV.D	F10,F0,F6	✓		
ADD.D	F6,F8,F2	✓		

Reservation stations							
Name	Busy	Op	Vj	Vk	Qj	Qk	A
Load1	No						
Load2	Yes	Load					44 + Regs[R3]
Add1	Yes	SUB		Mem[32 + Regs[R2]]	Load2		
Add2	Yes	ADD			Add1	Load2	
Add3	No						
Mult1	Yes	MUL		Regs[F4]	Load2		
Mult2	Yes	DIV		Mem[32 + Regs[R2]]	Mult1		

Register status									
Field	F0	F2	F4	F6	F8	F10	F12	...	F30
Qi	Mult1	Load2		Add2	Add1	Mult2			

Reservation Station Fields

- Op – inst opcode
- Qj, Qk – RS that produce the operands
- Vj, Vk – operand values
- A – hold memory address
- Qi – RS whose output should be stored into this register

Reservation stations								
Name	Busy	Op	Vj	Vk		Qj	Qk	A
Load1	No							
Load2	Yes	Load						44 + Regs[R3]
Add1	Yes	SUB		Mem[32 + Regs[R2]]		Load2		
Add2	Yes	ADD				Add1	Load2	
Add3	No							
Mult1	Yes	MUL		Regs[F4]		Load2		
Mult2	Yes	DIV		Mem[32 + Regs[R2]]		Mult1		

Register status									
Field	F0	F2	F4	F6	F8	F10	F12	...	F30
Qi	Mult1	Load2		Add2	Add1	Mult2			

A Loop Example

```
Loop: L.D      F0, 0(R1)
      MUL.D    F4, F0, F2
      S.D      F3, 0(R1)
      DADDIU   R1, R1, -8
      BNE      R1, R2, Loop
```

Fx: FP registers

Rx: integer registers

Assume predict-taken for loop unrolling

		Instruction status			
Instruction		From iteration	Issue	Execute	Write result
L.D	F0,0(R1)	1	✓	✓	
MUL.D	F4,F0,F2	1	✓		
S.D	F4,0(R1)	1	✓		
L.D	F0,0(R1)	2	✓	✓	
MUL.D	F4,F0,F2	2	✓		
S.D	F4,0(R1)	2	✓		

Reservation stations							
Name	Busy	Op	Vj	Vk	Qj	Qk	A
Load1	Yes	Load					Regs[R1] + 0
Load2	Yes	Load					Regs[R1] – 8
Add1	No						
Add2	No						
Add3	No						
Mult1	Yes	MUL		Regs[F2]	Load1		
Mult2	Yes	MUL		Regs[F2]	Load2		
Store1	Yes	Store	Regs[R1]			Mult1	
Store2	Yes	Store	Regs[R1] – 8			Mult2	

Register status									
Field	F0	F2	F4	F6	F8	F10	F12	...	F30
Qi	Load2		Mult2						

Ordering of Load and Store

- If load and store accesses different memory, they can be re-ordered
- Otherwise, reordering causes WAW/RAW hazards

L.D	F1, XXX	S.D	F2, XXX
S.D	F2, XXX	S.D	F2, YYY

- Consider **load** is next to issue
 - Compute effective address for **load**
 - Check the **A** field in all **store buffers** for match
 - Hold **load** if a match is found
- Store is handled similarly
 - Check both **load** and **store** buffers

Tomasulo's Algorithm – Summary

- Rename registers to RS
 - To eliminates WAW and WAR hazards
- Buffer RS # for operands currently unavailable
 - To eliminate RAW hazards
- Tolerate unpredictable delays of mem hierarchy
- Achieve high performance with compiler optimization
- HW complexity increases substantially
 - Increase area and power